

Curious OCaml

Lukasz Stafiniak

Contents

Curious OCaml	4
Chapter 1: Logic	4
1.1 In the Beginning there was Logos	4
1.2 Rules for Logical Connectives	4
1.3 Logos was Programmed in OCaml	6
1.4 Exercises	8
Chapter 2: Algebra	10
2.1 A Glimpse at Type Inference	10
2.2 Algebraic Data Types	11
2.3 Syntactic Conventions	13
2.4 Pattern Matching	13
2.5 Interpreting Algebraic Data Types as Polynomials	15
2.6 Differentiating Algebraic Data Types	18
2.7 Exercises	20
Chapter 3: Computation	23
3.1 Function Composition	23
3.2 Evaluation Rules (Reduction Semantics)	25
3.3 Symbolic Derivation Example	28
3.4 Tail Calls and Tail Recursion	29
3.5 First Encounter of Continuation Passing Style	31
3.6 Exercises	32
Chapter 4: Functions	33
4.1 Review: Computation by Hand	33
4.2 Language and Rules of the Untyped Lambda-Calculus	35
4.3 Booleans	36
4.4 If-then-else and Pairs	37
4.5 Pair-Encoded Natural Numbers	38
4.6 Church Numerals	38
4.7 Recursion: Fixpoint Combinators	41
4.8 Encoding Lists and Trees	43
4.9 Looping Recursion	45
4.10 Exercises	46
Chapter 5: Polymorphism and Abstract Data Types	48

5.1 Type Inference	48
5.2 Parametric Types	49
5.3 Type Inference, Formally	50
5.4 Algebraic Specification	53
5.5 Homomorphisms	55
5.6 Example: Maps	55
5.7 Modules and Interfaces (Signatures): Syntax	56
5.8 Implementing Maps: Association Lists	56
5.9 Implementing Maps: Binary Search Trees	58
5.10 Implementing Maps: Red-Black Trees	59
Exercises	61
Chapter 6: Folding and Backtracking	63
6.1 Basic Generic List Operations	63
6.2 Making Fold Tail-Recursive	65
6.3 Map and Fold for Trees and Other Structures	66
6.4 Point-Free Programming	69
6.5 Reductions and More Higher-Order Functions	70
6.6 Grouping and Map-Reduce	71
6.7 Higher-Order Functions for the Option Type	73
6.8 The Countdown Problem Puzzle	73
6.9 The Honey Islands Puzzle	76
6.10 Constraint-Based Puzzles	80
6.11 Exercises	81
Chapter 7: Laziness	82
7.1 Evaluation Strategies and Parameter Passing	82
7.2 Call-by-name: Streams	83
7.3 Lazy Values	84
7.4 Power Series and Differential Equations	86
7.5 Arbitrary Precision Computation	89
7.6 Circular Data Structures: Double-Linked Lists	90
7.7 Input-Output Streams	91
7.8 Pipes	91
7.9 Example: Pretty-Printing	93
7.10 Exercises	97
Chapter 8: Monads	99
8.1 List Comprehensions	99
8.2 Generalized Comprehensions: Binding Operators	101
8.3 Monads	102
8.4 Monad Laws	104
8.5 Monoid Laws and Monad-Plus	104
8.6 Backtracking: Computation with Choice	105
8.7 Monad Flavors	106
8.8 Interlude: The Module System	107
8.9 The Two Metaphors	108
8.10 Monad Classes and Instances	109
8.11 Monad Instances	111

8.12 Monad Transformers	118
8.13 Probabilistic Programming	121
8.14 Lightweight Cooperative Threads	128
8.15 Exercises	132
Chapter 9: Compiler	134
Compilation, Runtime, Optimization, and Parsing	134
9.1 OCaml Compilers	134
9.2 Imperative Features in OCaml	139
9.3 OCaml Garbage Collection	140
9.4 Stack Frames and Closures	142
9.5 Profiling and Optimization	143
9.6 Parsing: ocamllex and Menhir	148
9.7 Example: Phrase Search	162
9.8 Exercises	168
Chapter 10: Functional Reactive Programming	170
10.1 Zippers	170
10.2 Example: Context Rewriting	172
10.3 Adaptive Programming (Incremental Computing)	174
10.4 Functional Reactive Programming	176
10.5 Reactivity by Stream Processing	179
10.6 Reactivity by Incremental Computing	183
10.7 Direct Control	185
10.8 Graphical User Interfaces	188
10.9 Exercises	191
Chapter 11: The Expression Problem	192
11.1 The Expression Problem: Definition	193
11.2 Functional Programming Non-Solution: Ordinary Algebraic Datatypes	194
11.3 Lightweight FP Non-Solution: Extensible Variant Types . .	196
11.4 Object-Oriented Programming: Subtyping	198
11.5 Direct Object-Oriented Non-Solution	199
11.6 OOP Non-Solution: The Visitor Pattern	203
11.7 Polymorphic Variants	205
11.8 Polymorphic Variants with Recursive Modules	208
11.9 Parser Combinators	211
11.10 Parser Combinators: Implementation	212
11.11 Parser Combinators: Tying the Recursive Knot	214
11.12 Parser Combinators: Dynamic Code Loading	215
11.13 Parser Combinators: Toy Example	215
11.14 Exercises	216

Curious OCaml

Chapter 1: Logic

From logic rules to programming constructs

1.1 In the Beginning there was Logos

What logical connectives do you know?

\top	\perp	\wedge	\vee	\rightarrow
truth	falsehood	$a \wedge b$	$a \vee b$	$a \rightarrow b$
“trivial”	“impossible”	conjunction a and b	disjunction a or b	implication a gives b
	shouldn’t get	got both	got at least one	given a , we get b

How can we define them? Think in terms of *derivation trees*:

$$\frac{\begin{array}{c} \text{a premise} & \text{another premise} & \text{this we have by default} \\ \hline \text{some fact} & & \text{another fact} \end{array}}{\text{final conclusion}}$$

We define connectives by providing rules for using them. For example, a rule $\frac{a \ b}{c}$ matches parts of the tree that have two premises, represented by variables a and b , and have any conclusion, represented by variable c .

Design principle: Try to use only the connective you define in its definition.

1.2 Rules for Logical Connectives

Introduction rules say how to *produce* a connective.

Elimination rules say how to *use* it.

Text in parentheses is comments. Letters are variables that can stand for anything.

Connective	Introduction Rules	Elimination Rules
\top	$\overline{\top}$	doesn’t have
\perp	doesn’t have	$\frac{\perp}{\perp}$ (i.e., anything)
\wedge	$\frac{a \ b}{a \wedge b}$	$\frac{\frac{a \wedge b}{a}}{a}$ (take first) $\frac{\frac{a \wedge b}{b}}{b}$ (take second)
\vee	$\frac{a}{a \vee b}$ (put first) $\frac{b}{a \vee b}$ (put second)	$\frac{\frac{[a]^x \quad [b]^y}{a \vee b} \quad \frac{\vdots_c \quad \vdots_c}{c}}{c}$ using x, y

Connective	Introduction Rules	Elimination Rules
\rightarrow	$[a]^x$ $\frac{\vdots^b}{a \rightarrow b}$ using x	$\frac{a \rightarrow b}{b}$

Notation for Hypothetical Derivations The notation $\frac{[a]^x}{\vdots^b}$ (sometimes written as a tree) matches any subtree that derives b and can use a as an assumption (marked with label x), even though a might not otherwise be warranted.

For example, we can derive “sunny \rightarrow happy” by showing that *assuming* it’s sunny, we can derive happiness:

$$\frac{\frac{\frac{\overline{\text{sunny}}^x}{\overline{\text{go outdoor}}}}{\overline{\text{playing}}}}{\text{happy}} \text{ using } x$$

$$\text{sunny} \rightarrow \text{happy}$$

Such assumptions can only be used in the matched subtree! But they can be used several times. For example, if someone’s mood is more difficult to influence:

$$\frac{\frac{\frac{\overline{\text{sunny}}^x}{\overline{\text{go outdoor}}}}{\text{playing}} \quad \frac{\overline{\text{sunny}}^x}{\text{nice view}} \quad \frac{\overline{\text{sunny}}^x}{\overline{\text{go outdoor}}}}{\text{happy}} \text{ using } x$$

$$\text{sunny} \rightarrow \text{happy}$$

Reasoning by Cases The elimination rule for disjunction represents **reasoning by cases**.

How can we use the fact that it is sunny \vee cloudy (but not rainy)?

$$\frac{\text{sunny} \vee \text{cloudy} \quad \text{forecast}}{\text{no-umbrella}} \quad \frac{\frac{\overline{\text{sunny}}^x}{\text{no-umbrella}} \quad \frac{\overline{\text{cloudy}}^y}{\text{no-umbrella}}}{\text{no-umbrella}} \text{ using } x, y$$

We know that it will be sunny or cloudy (by watching the weather forecast). If it will be sunny, we won’t need an umbrella. If it will be cloudy, we won’t need an umbrella. Therefore, we won’t need an umbrella.

Reasoning by Induction We need one more kind of rule to do serious math: **reasoning by induction** (somewhat similar to reasoning by cases). Example rule for induction on natural numbers:

$$\frac{p(0) \quad [p(x)]^x}{p(n)} \frac{\vdots^{p(x+1)}}{\text{by induction, using } x}$$

We get property p for any natural number n , provided we can: 1. Establish $p(0)$ (the base case) 2. Show that assuming $p(x)$ holds, we can derive $p(x + 1)$ (the inductive step)

Here x is a unique variable—we cannot substitute a particular number for it because we write “using x ” on the side.

1.3 Logos was Programmed in OCaml

There is a deep correspondence between logic and programming, known as the **Curry-Howard correspondence** (or “propositions as types”). The following table shows how logical connectives correspond to programming constructs:

Logic	Type	Expression
\top	<code>unit</code>	<code>()</code>
\perp	<code>'a</code>	<code>raise</code>
\wedge	<code>*</code>	<code>(,)</code>
\vee	<code> </code>	<code>match</code>
\rightarrow	<code>-></code>	<code>fun</code>
induction	—	<code>rec</code>

Typing rules for OCaml constructs:

- **Unit (truth):** $\frac{}{\text{()}: \text{unit}}$
- **Exception (falsehood):** $\frac{\text{oops!}}{\text{raise } \text{exn}: \text{c}}$ — can produce any type
- **Pair (conjunction):**
 - Introduction: $\frac{s:a \quad t:b}{(s,t): a * b}$
 - Elimination: $\frac{p:a * b}{\text{fst } p:a} \text{ and } \frac{p:a * b}{\text{snd } p:b}$
- **Variant (disjunction):**
 - Introduction: $\frac{s:a}{\text{A}(s): \text{A or } a \mid \text{B or } b}$
 - Elimination (match): given t of variant type and branches for each case, produce result c
- **Function (implication):**
 - Introduction: $\frac{[x:a] \quad e:b}{\text{fun } x \rightarrow e: a \rightarrow b}$
 - Elimination (application): $\frac{f:a \rightarrow b \quad t:a}{f \ t:b}$
- **Recursion (induction):** $\frac{[x:a]}{\text{rec } x = e: a}$

1.3.1 Definitions Writing out expressions and types repetitively is tedious: we need definitions.

Type definitions are written: `type ty = some type`.

- Writing `A(s) : A of a | B of b` in the table was cheating. Usually we have to define the type and then use it. For example, using `int` for `a` and `string` for `b`:

```
type int_string_choice = A of int | B of string
```

This allows us to write `A(s) : int_string_choice`.

- Without the type definition, it is difficult to know what other variants there are when one *infers* (i.e., “guesses”, computes) the type!
- In OCaml we can write ``A(s) : [`A of a | `B of b]`. With “```” variants (polymorphic variants), OCaml does guess what other variants there are. These types are interesting, but we will not focus on them in this book.
- Tuple elements don’t need labels because we always know at which position a tuple element stands. But having labels makes code more clear, so we can define a *record type*:

```
type int_string_record = {a: int; b: string}
```

and create its values: `{a = 7; b = "Mary"}`.

- We access the *fields* of records using the dot notation: `{a=7; b="Mary"}.b = "Mary"`.

1.3.2 Expression Definitions The recursive expression `rec x = e` in the table was cheating: `rec` (usually called `fix` in theory) cannot appear alone in OCaml! It must be part of a definition.

Definitions for expressions are introduced by rules a bit more complex:

$$\frac{e_1 : a \quad \frac{[x:a]}{e_2 : b}}{\text{let } x = e_1 \text{ in } e_2 : b}$$

(Note that this rule is the same as introducing and eliminating \rightarrow .)

For recursive definitions:

$$\frac{\frac{[x:a]}{e_1 : a} \quad \frac{[x:a]}{e_2 : b}}{\text{let rec } x = e_1 \text{ in } e_2 : b}$$

We will cover what is missing in the above rules when we discuss **polymorphism**.

1.3.3 Scoping Rules

- **Type definitions** we have seen above are *global*: they need to be at the top-level (not nested in expressions), and they extend from the point they occur till the end of the source file or interactive session.
- **let-in definitions** for expressions: `let x = e1 in e2` are *local*—*x* is only visible in *e₂*. But **let definitions** without **in** are global: placing `let x = e1` at the top-level makes *x* visible from after *e₁* till the end of the source file or interactive session.
- In the interactive session (toplevel/REPL), we mark the end of a top-level “sentence” with `;;`—this is unnecessary in source files.

1.3.4 Operators

Operators like `+`, `*`, `<`, `=` are names of functions. Just like other names, you can use operator names for your own functions:

```
let (+:) a b = String.concat "" [a; b] (* Special way of defining *)
"Alpha" +: "Beta" (* but normal way of using operators *)
```

Operators in OCaml are **not overloaded**. This means that every type needs its own set of operators: `-`, `+`, `*`, `/` work for integers `-.`, `+. .`, `/. .` work for floating point numbers

Exception: Comparisons `<`, `=`, etc. work for all values other than functions.

1.4 Exercises

Exercises from *Think OCaml: How to Think Like a Computer Scientist* by Nicholas Monje and Allen Downey.

1. Assume that we execute the following assignment statements:

```
let width = 17
let height = 12.0
let delimiter = '.'
```

For each of the following expressions, write the value of the expression and the type (of the value of the expression), or the resulting type error.

1. `width/2`
 2. `width/.2.0`
 3. `height/3`
 4. `1 + 2 * 5`
 5. `delimiter * 5`
2. Practice using the OCaml interpreter as a calculator:
 1. The volume of a sphere with radius *r* is $\frac{4}{3}\pi r^3$. What is the volume of a sphere with radius 5? (*Hint*: 392.6 is wrong!)

- Suppose the cover price of a book is \$24.95, but bookstores get a 40% discount. Shipping costs \$3 for the first copy and 75 cents for each additional copy. What is the total wholesale cost for 60 copies?
- If I leave my house at 6:52 am and run 1 mile at an easy pace (8:15 per mile), then 3 miles at tempo (7:12 per mile) and 1 mile at easy pace again, what time do I get home for breakfast?
- You've probably heard of the Fibonacci numbers before, but in case you haven't, they're defined by the following recursive relationship:

$$\begin{cases} f(0) = 0 \\ f(1) = 1 \\ f(n + 1) = f(n) + f(n - 1) \quad \text{for } n = 2, 3, \dots \end{cases}$$

Write a recursive function to calculate these numbers.

- A palindrome is a word that is spelled the same backward and forward, like "noon" and "redivider". Recursively, a word is a palindrome if the first and last letters are the same and the middle is a palindrome.

The following are functions that take a string argument and return the first, last, and middle letters:

```
let first_char word = word.[0]
let last_char word =
  let len = String.length word - 1 in
  word.[len]
let middle word =
  let len = String.length word - 2 in
  String.sub word 1 len
```

- Enter these functions into the toplevel and test them out. What happens if you call `middle` with a string with two letters? One letter? What about the empty string ""?
- Write a function called `is_palindrome` that takes a string argument and returns `true` if it is a palindrome and `false` otherwise.
- The greatest common divisor (GCD) of a and b is the largest number that divides both of them with no remainder.

One way to find the GCD of two numbers is Euclid's algorithm, which is based on the observation that if r is the remainder when a is divided by b , then $\gcd(a, b) = \gcd(b, r)$. As a base case, we can consider $\gcd(a, 0) = a$.

Write a function called `gcd` that takes parameters `a` and `b` and returns their greatest common divisor.

If you need help, see http://en.wikipedia.org/wiki/Euclidean_algorithm.

Chapter 2: Algebra

Algebraic Data Types and some curious analogies

2.1 A Glimpse at Type Inference

For a refresher, let us apply the type inference rules introduced in Chapter 1 to some simple examples. We will start with the identity function `fun x -> x`. In the derivations below, [?] means “type unknown yet.”

We begin with an incomplete derivation:

$$\frac{[?]}{\mathbf{fun} \ x \ -\!\!-\> \ x : [?]}$$

Using the \rightarrow introduction rule, we need to derive the body `x` assuming `x` has some type `a`:

$$\frac{\overline{x : a}^x}{\mathbf{fun} \ x \ -\!\!-\> \ x : [?] \rightarrow [?]}$$

The premise $\overline{x : a}^x$ matches the pattern for hypothetical derivations since $e = x$. Since the body `x` has type `a` (from our assumption), and the parameter `x` also has type `a`, we conclude:

$$\frac{\overline{x : a}^x}{\mathbf{fun} \ x \ -\!\!-\> \ x : a \rightarrow a}$$

Because `a` is arbitrary (we made no assumptions constraining it), OCaml introduces a *type variable* '`a`' to represent it:

```
# fun x -> x;;
- : 'a -> 'a = <fun>
```

A More Complex Example Let us try `fun x -> x+1`, which is the same as `fun x -> ((+) x) 1` (try it in OCaml!). We will use the notation `[?α]` to mean “type unknown yet, but the same as in other places marked `[?α]`.”

Starting the derivation and applying \rightarrow introduction:

$$\frac{[?]}{\mathbf{fun} \ x \ -\!\!-\> \ ((+) \ x) \ 1 : [?] \rightarrow [?]} \quad \frac{[?]}{\mathbf{fun} \ x \ -\!\!-\> \ ((+) \ x) \ 1 : [?] \rightarrow [?α]}$$

Applying \rightarrow elimination (function application) to `((+) x) 1`:

$$\frac{\frac{[?] \quad [?]}{\text{(+) } x : [? \beta] \rightarrow [? \alpha]} \quad \frac{[?]}{1 : [? \beta]}}{\text{fun } x \rightarrow ((+) x) \ 1 : [?] \rightarrow [? \alpha]}$$

We know that $1 : \text{int}$, so $[? \beta] = \text{int}$:

$$\frac{\frac{[?] \quad \text{(constant)} \quad 1:\text{int}}{\text{(+) } x : \text{int} \rightarrow [? \alpha]} \quad \frac{[?]}{1 : [? \alpha]}}{\text{fun } x \rightarrow ((+) x) \ 1 : [?] \rightarrow [? \alpha]}$$

Applying function application again to $(+) x$:

$$\frac{\frac{\frac{[?] \quad [?] \quad \text{(constant)} \quad 1:\text{int}}{\text{(+) } x : [? \gamma] \rightarrow \text{int} \rightarrow [? \alpha]} \quad \frac{[?]}{x : [? \gamma]}}{\text{(+) } x : \text{int} \rightarrow [? \alpha]} \quad \frac{[?]}{1 : [? \alpha]}}{\text{fun } x \rightarrow ((+) x) \ 1 : [?] \rightarrow [? \alpha]}$$

Since $(+) : \text{int} \rightarrow \text{int} \rightarrow \text{int}$, we have $[? \gamma] = \text{int}$ and $[? \alpha] = \text{int}$:

$$\frac{\frac{\frac{(\text{constant}) \quad \bar{x}:\text{int}^x \quad (\text{constant})}{(\text{+) } \text{int} \rightarrow \text{int} \rightarrow \text{int}} \quad \frac{(\text{constant})}{(\text{+) } x : \text{int} \rightarrow \text{int}}} {\text{(+) } x : \text{int} \rightarrow \text{int}} \quad \frac{1:\text{int}}{1 : \text{int}}} {\text{fun } x \rightarrow ((+) x) \ 1 : \text{int} \rightarrow \text{int}}$$

2.1.1 Curried Form When there are several arrows “on the same depth” in a function type, it means that the function returns a function. For example, $(+) : \text{int} \rightarrow \text{int} \rightarrow \text{int}$ is just a shorthand for $(+) : \text{int} \rightarrow (\text{int} \rightarrow \text{int})$. This is very different from:

```
fun f -> (f 1) + 1 : (int → int) → int
```

In the first case, $(+)$ is a function that takes an integer and returns a function from integers to integers. In the second case, we have a function that takes a function as an argument.

For addition, instead of $\text{fun } x \rightarrow x + 1$ we can write $((+) 1)$. What expanded form does $((+) 1)$ correspond to exactly (computationally)? It corresponds to $\text{fun } y \rightarrow 1 + y$.

We will become more familiar with functions returning functions when we study the *lambda calculus* in a later chapter.

2.2 Algebraic Data Types

In Chapter 1, we learned about the `unit` type and variant types like:

```
type int_string_choice = A of int | B of string
```

We also covered tuple types, record types, and type definitions. Let us now explore these concepts more deeply.

Variants Without Arguments Variants do not have to carry arguments. Instead of writing `A of unit`, we can simply use `A`. This is more convenient and idiomatic:

```
type color = Red | Green | Blue
```

A subtle point about OCaml: In OCaml, variants take multiple arguments rather than taking tuples as arguments. This means `A of int * string` is different from `A of (int * string)`. The first takes two separate arguments, while the second takes a single tuple argument. This distinction is usually not important unless you encounter situations where it matters.

Recursive Type Definitions Type definitions can be recursive! This allows us to define data structures of arbitrary size:

```
type int_list = Empty | Cons of int * int_list
```

Let us see what values inhabit `int_list`: - `Empty` represents the empty list - `Cons (5, Empty)` is a list containing just 5 - `Cons (5, Cons (7, Cons (13, Empty)))` is a list containing 5, 7, and 13

The built-in type `bool` can be viewed as if it were defined as `type bool = true | false`. Similarly, `int` can be thought of as a very large variant: `type int = 0 | -1 | 1 | -2 | 2 | ...`

Parametric Type Definitions Type definitions can be *parametric* with respect to the types of their components. This allows us to define generic data structures that work with any element type. For example, a list of elements of arbitrary type:

```
type 'elem list = Empty | Cons of 'elem * 'elem list
```

Several conventions and syntax rules apply to parametric types:

- Type variables must start with '`'`, but since OCaml will not remember the names we give, it is customary to use the names OCaml uses: `'a`, `'b`, `'c`, `'d`, etc.
- The OCaml syntax places the type parameter before the type name, mimicking English word order. A silly example:

```
type 'white_color dog = Dog of 'white_color
```

- With multiple parameters, OCaml uses parentheses:

```
type ('a, 'b) choice = Left of 'a | Right of 'b
```

Compare this to F# syntax: `type choice<'a,'b> = Left of 'a | Right of 'b`

```
And Haskell syntax: data Choice a b = Left a | Right b
```

2.3 Syntactic Conventions

Constructor Naming Names of variants, called *constructors*, must start with a capital letter. If we wanted to define our own booleans, we would write:

```
type my_bool = True | False
```

Only constructors and module names can start with capital letters in OCaml. *Modules* are organizational units (like “shelves”) containing related values. For example, the `List` module provides operations on lists, including `List.map` and `List.filter`.

Accessing Record Fields We can use dot notation to access record fields: `record.field`. For example, if we have `let person = {name="Alice"; age=30}`, we can write `person.name` to get "Alice".

Function Definition Shortcuts Several syntactic shortcuts make function definitions more concise:

- `fun x y -> e` stands for `fun x -> fun y -> e`. Note that `fun x -> fun y -> e` parses as `fun x -> (fun y -> e)`.
- `function A x -> e1 | B y -> e2` stands for `fun p -> match p with A x -> e1 | B y -> e2`. The general form is: `function PATTERN-MATCHING` stands for `fun v -> match v with PATTERN-MATCHING`.
- `let f ARGs = e` is a shorthand for `let f = fun ARGs -> e`.

2.4 Pattern Matching

Recall that we introduced `fst` and `snd` as means to access elements of a pair. But what about larger tuples? The fundamental way to access any tuple uses the `match` construct. In fact, `fst` and `snd` can easily be defined using pattern matching:

```
let fst = fun p -> match p with (a, b) -> a
let snd = fun p -> match p with (a, b) -> b
```

Matching on Records Pattern matching also works with records:

```
type person = {name: string; surname: string; age: int}

let greet_person () =
  match {name="Walker"; surname="Johnnie"; age=207}
  with {name=n; surname=sn; age=a} -> "Hi " ^ sn ^ "!"
```

Understanding Patterns The left-hand sides of `->` in `match` expressions are called **patterns**. Patterns describe the structure of values we want to match against.

Patterns can be nested, allowing us to match complex structures:

```
match Some (5, 7) with
| None -> "sum: nothing"
| Some (x, y) -> "sum: " ^ string_of_int (x+y)
```

Simple Patterns and Wildcards A pattern can simply bind the entire value without destructuring. Writing `match f x with v -> ...` is the same as `let v = f x in ...`

When we do not need a value in a pattern, it is good practice to use the underscore `_`, which is a wildcard (not a variable):

```
let fst (a, _) = a
let snd (_, b) = b
```

Pattern Linearity A variable can only appear once in a pattern. This property is called *linearity*. However, we can add conditions to patterns using `when`, so linearity is not a limitation:

```
let describe_point p =
  match p with
  | (x, y) when x = y -> "diag"
  | _ -> "off-diag"
```

Here is a more elaborate example:

```
let compare a b = match a, b with
  | (x, y) when x < y -> -1
  | (x, y) when x = y -> 0
  | _ -> 1
```

Partial Record Patterns We can skip unused fields of a record in a pattern. Only the fields we care about need to be mentioned.

Or-Patterns We can compress patterns by using `|` inside a single pattern to match multiple alternatives:

```
type month =
  | Jan | Feb | Mar | Apr | May | Jun
  | Jul | Aug | Sep | Oct | Nov | Dec

type weekday = Mon | Tue | Wed | Thu | Fri | Sat | Sun

type date =
```

```

{year: int; month: month; day: int; weekday: weekday}

let day =
  {year = 2012; month = Feb; day = 14; weekday = Wed};;

match day with
| {weekday = Sat | Sun} -> "Weekend!"
| _ -> "Work day"

```

Named Patterns with as We use (pattern as v) to name a nested pattern, binding the matched value to v:

```

match day with
| {weekday = (Mon | Tue | Wed | Thu | Fri as wday)}
  when not (day.month = Dec && day.day = 24) ->
  Some (work (get_plan wday))
| _ -> None

```

This example shows the as keyword binding the matched weekday to wday for use in the expression on the right side of the arrow.

2.5 Interpreting Algebraic Data Types as Polynomials

Let us explore a curious analogy between algebraic data types and polynomials. We translate data types to mathematical expressions by:

- Replacing | (variant choice) with +
- Replacing * (tuple product) with \times
- Treating record types as tuple types (erasing field names and translating ; as \times)

We also need translations for some special types:

- The **void type** (a type with no constructors, hence no values):

```
type void
```

(Yes, this is its complete definition, with no = something part.) Translate it as 0.

- The **unit type** translates as 1. Since variants without arguments behave like variants of unit, translate them as 1 as well.
- The **bool type** translates as 2.
- Types like **int**, **string**, **float**, and type parameters translate as variables.
- Defined types translate according to their definitions (substituting variables as necessary).

Give a name to the type being defined (representing a function of the introduced variables). Now interpret the result as an ordinary numeric polynomial! (Or a “rational function” if recursively defined.)

Let us have some fun with this translation.

Example: Date Type

```
type date = {year: int; month: int; day: int}
```

Translating to a polynomial (using x for `int`):

$$D = x \times x \times x = x^3$$

Example: Option Type The built-in option type is defined as:

```
type 'a option = None | Some of 'a
```

Translating:

$$O = 1 + x$$

Example: List Type

```
type 'a my_list = Empty | Cons of 'a * 'a my_list
```

Translating (where L represents the list type):

$$L = 1 + x \cdot L$$

Example: Binary Tree Type

```
type btree = Tip | Node of int * btree * btree
```

Translating:

$$T = 1 + x \cdot T \cdot T = 1 + x \cdot T^2$$

Type Isomorphisms When translations of two types are equal according to the laws of high-school algebra, the types are *isomorphic*. This means there exist bijective (one-to-one and onto) functions between them.

Let us manipulate the binary tree polynomial:

$$\begin{aligned}
T &= 1 + x \cdot T^2 \\
&= 1 + x \cdot T + x^2 \cdot T^3 \\
&= 1 + x + x^2 \cdot T^2 + x^2 \cdot T^3 \\
&= 1 + x + x^2 \cdot T^2 \cdot (1 + T) \\
&= 1 + x \cdot (1 + x \cdot T^2 \cdot (1 + T))
\end{aligned}$$

Now let us translate the resulting expression back to a type:

```
type repr =
  (int * (int * btree * btree * btree option) option) option
```

The challenge is to find isomorphism functions with signatures:

```
val iso1 : btree -> repr
val iso2 : repr -> btree
```

These functions should satisfy: for all trees t , $\text{iso2}(\text{iso1 } t) = t$, and for all representations r , $\text{iso1}(\text{iso2 } r) = r$.

A First Attempt Here is a first (failed) attempt:

```
# let iso1 (t : btree) : repr =
  match t with
  | Tip -> None
  | Node (x, Tip, Tip) -> Some (x, None)
  | Node (x, Node (y, t1, t2), Tip) ->
    Some (x, Some (y, t1, t2, None))
  | Node (x, Node (y, t1, t2), t3) ->
    Some (x, Some (y, t1, t2, Some t3));;
```

Warning 8: this pattern-matching is not exhaustive.

Here is an example of a value that is not matched:

```
Node (_, Tip, Node (_, _, _))
```

We forgot about one case! It seems difficult to guess the solution directly.

Breaking Down the Problem Let us divide the task into smaller steps corresponding to intermediate points in the polynomial transformation:

```
type ('a, 'b) choice = Left of 'a | Right of 'b

type interm1 =
  ((int * btree, int * int * btree * btree * btree) choice)
  option

type interm2 =
```

```
((int, int * int * btree * btree * btree option) choice)
option
```

Now we can define each step:

```
let step1r (t : btree) : interm1 =
  match t with
    | Tip -> None
    | Node (x, t1, Tip) -> Some (Left (x, t1))
    | Node (x, t1, Node (y, t2, t3)) ->
      Some (Right (x, y, t1, t2, t3))

let step2r (r : interm1) : interm2 =
  match r with
    | None -> None
    | Some (Left (x, Tip)) -> Some (Left x)
    | Some (Left (x, Node (y, t1, t2))) ->
      Some (Right (x, y, t1, t2, None))
    | Some (Right (x, y, t1, t2, t3)) ->
      Some (Right (x, y, t1, t2, Some t3))

let step3r (r : interm2) : repr =
  match r with
    | None -> None
    | Some (Left x) -> Some (x, None)
    | Some (Right (x, y, t1, t2, t3opt)) ->
      Some (x, Some (y, t1, t2, t3opt))

let iso1 (t : btree) : repr =
  step3r (step2r (step1r t))
```

Defining `step1l`, `step2l`, `step3l`, and `iso2` is now straightforward—each step is the inverse of its corresponding forward step.

Take-Home Lessons

- Design for validity:** Try to define data structures so that only meaningful information can be represented—as long as it does not overcomplicate the data structures. Avoid catch-all clauses when defining functions. The compiler will then tell you if you have forgotten about a case.
- Divide and conquer:** Break solutions into small steps so that each step can be easily understood and verified.

2.6 Differentiating Algebraic Data Types

The title might seem strange—we will differentiate the translated polynomials, not the types themselves. But what sense does this make?

It turns out that taking the partial derivative of a polynomial (translated from a data type), when translated back, gives a type representing how to change one occurrence of a value corresponding to the variable with respect to which we differentiated. In other words, the derivative represents a “context” or “hole” in the data structure.

Example: Differentiating the Date Type

```
type date = {year: int; month: int; day: int}
```

The translation:

$$D = x \cdot x \cdot x = x^3$$

$$\frac{\partial D}{\partial x} = 3x^2 = x \cdot x + x \cdot x + x \cdot x$$

We could have left it as $3 \cdot x \cdot x$, but expanding shows the structure more clearly. Translating back to a type:

```
type date_deriv =
  Year of int * int | Month of int * int | Day of int * int
```

Each variant represents a “hole” at a different position: `Year` means the year field is missing (and we have the month and day), and so on.

Now we can define functions to introduce and eliminate this derivative type:

```
let date_deriv {year=y; month=m; day=d} =
  [Year (m, d); Month (y, d); Day (y, m)]
```



```
let date_integr n = function
  | Year (m, d) -> {year=n; month=m; day=d}
  | Month (y, d) -> {year=y; month=n; day=d}
  | Day (y, m) -> {year=y; month=m; day=n}
;;

```



```
List.map (date_integr 7)
  (date_deriv {year=2012; month=2; day=14}))
```

The `date_deriv` function produces all contexts (one for each field), and `date_integr` fills in a hole with a new value.

Example: Differentiating Binary Trees Let us tackle the more challenging case of binary trees:

```
type btree = Tip | Node of int * btree * btree
```

The translation and differentiation:

$$T = 1 + x \cdot T^2$$

$$\frac{\partial T}{\partial x} = 0 + T^2 + 2 \cdot x \cdot T \cdot \frac{\partial T}{\partial x} = T \cdot T + 2 \cdot x \cdot T \cdot \frac{\partial T}{\partial x}$$

The derivative is recursive! This makes sense: a context in a tree is either at the current node ($T \cdot T$, the two subtrees) or somewhere below ($2 \cdot x \cdot T \cdot \frac{\partial T}{\partial x}$, choosing left or right, with the node value, the other subtree, and a deeper context).

Instead of translating 2 as `bool`, we introduce a more descriptive type:

```
type btree_dir = LeftBranch | RightBranch
```

```
type btree_deriv =
| Here of btree * btree
| Below of btree_dir * int * btree * btree_deriv
```

(You might someday hear about *zippers*—they are “inverted” relative to our type, with the hole coming first.)

The integration function fills the hole with a value:

```
let rec btree_integr n = function
| Here (ltree, rtree) -> Node (n, ltree, rtree)
| Below (LeftBranch, m, rtree, deriv) ->
  Node (m, btree_integr n deriv, rtree)
| Below (RightBranch, m, ltree, deriv) ->
  Node (m, ltree, btree_integr n deriv)
```

2.7 Exercises

Exercise 1 *Due to Yaron Minsky.*

Consider a datatype to store internet connection information. The time `when_initiated` marks the start of connecting and is not needed after the connection is established (it is only used to decide whether to give up trying to connect). The ping information is available for established connections but not straight away.

```
type connectionstate = Connecting | Connected | Disconnected

type connectioninfo = {
  state : connectionstate;
  server : Inetaddr.t;
  lastpingtime : Time.t option;
  lastpingid : int option;
  sessionid : string option;
  wheninitiated : Time.t option;
  whendisconnected : Time.t option;
}
```

(The types `Time.t` and `Inetaddr.t` come from the *Core* library. You can replace them with `float` and `Unix.inet_addr`. Load the Unix library in the interactive toplevel with `#load "unix.cma";;.`)

Rewrite the type definitions so that the datatype will contain only reasonable combinations of information.

Exercise 2 In OCaml, functions can have labeled arguments and optional arguments (parameters with default values that can be omitted). Labels can differ from the names of argument values:

```
let f ~meaningfulname:n = n + 1
let _ = f ~meaningfulname:5 (* We do not need the result so we ignore it. *)
```

When the label and value names are the same, the syntax is shorter:

```
let g ~pos ~len =
  StringLabels.sub "0123456789abcdefghijklmnopqrstuvwxyz" ~pos ~len

let () = (* A nicer way to mark computations that return unit. *)
  let pos = Random.int 26 in
  let len = Random.int 10 in
  print_string (g ~pos ~len)
```

When some function arguments are optional, the function must take non-optional arguments after the last optional argument. Optional parameters with default values:

```
let h ?(len=1) pos = g ~pos ~len
let () = print_string (h 10)
```

Optional arguments are implemented as parameters of an option type. This allows checking whether the argument was provided:

```
let foo ?bar n =
  match bar with
  | None -> "Argument = " ^ string_of_int n
  | Some m -> "Sum = " ^ string_of_int (m + n)
```

We can use it in various ways:

```
let _ = foo 5
let _ = foo ~bar:5 7
```

We can also provide the option value directly:

```
let test_foo () =
  let bar = if Random.int 10 < 5 then None else Some 7 in
  foo ?bar 7
```

1. Observe the types that functions with labeled and optional arguments have.
Come up with coding style guidelines for when to use labeled arguments.

2. Write a rectangle-drawing procedure that takes three optional arguments: left-upper corner, right-lower corner, and a width-height pair. It should draw a correct rectangle whenever two arguments are given, and raise an exception otherwise. Load the graphics library with `#load "graphics.cma";;`. Use `invalid_arg`, `Graphics.open_graph`, and `Graphics.draw_rect`.
3. Write a function that takes an optional argument of arbitrary type and a function argument, and passes the optional argument to the function without inspecting it.

Exercise 3 *From a past exam.*

1. Give the (most general) types of the following expressions, either by guessing or by inferring by hand:
 1. `let double f y = f (f y) in fun g x -> double (g x)`
 2. `let rec tails l = match l with [] -> [] | x::xs -> xs::tails xs in fun l -> List.combine l (tails l)`
2. Give example expressions that have the following types (without using type constraints):
 1. `(int -> int) -> bool`
 2. `'a option -> 'a list`

Exercise 4 We have seen that algebraic data types can be related to analytic functions (the subset definable from polynomials via recursion)—by literally interpreting sum types (variant types) as sums and product types (tuple and record types) as products. We can extend this interpretation to function types by interpreting $a \rightarrow b$ as b^a (i.e., b to the power of a). Note that the b^a notation is actually used to denote functions in set theory.

1. Translate a^{b+cd} and $a^b \cdot (a^c)^d$ into OCaml types, using any distinct types for a, b, c, d , and using type `('a,'b) choice = Left of 'a | Right of 'b` for $+$. Write the bijection function in both directions.
2. Come up with a type `'t exp` that shares with the exponential function the following property: $\frac{\partial \exp(t)}{\partial t} = \exp(t)$, where we translate a derivative of a type as a context (i.e., the type with a “hole”), as in this chapter. Explain why your answer is correct. *Hint:* in computer science, our logarithms are mostly base 2.

Further reading: Algebraic Type Systems - Combinatorial Species

Exercise 5 (Homework) Write a function `btree_deriv_at` that takes a predicate over integers (i.e., a function `f: int -> bool`) and a `btree`, and builds a `btree_deriv` whose “hole” is in the first position for which the predicate returns true. It should return a `btree_deriv option`, with `None` if the predicate does not hold for any node.

Chapter 3: Computation

Reduction semantics and operational reasoning

References:

- “Using, Understanding and Unraveling the OCaml Language” by Didier Remy, Chapter 1
- “The OCaml system” manual, the tutorial part, Chapter 1

3.1 Function Composition

The usual way function composition is defined in mathematics is “backward”—the notation follows the convention of mathematical function application:

$$(f \circ g)(x) = f(g(x))$$

This means that when we write $f \circ g$, we first apply g and then apply f to the result. Here is how this is expressed in different functional programming languages:

Language	Definition
Math	$(f \circ g)(x) = f(g(x))$
OCaml	<code>let (-) f g x = f (g x)</code>
F#	<code>let (<<) f g x = f (g x)</code>
Haskell	<code>(.) f g = \x -> f (g x)</code>

This backward composition looks like function application but needs fewer parentheses. Recall the functions `iso1` and `iso2` from the previous chapter on type isomorphisms. Using backward composition, we could write:

```
let iso2 = step1l -| step2l -| step3l
```

A more natural definition of function composition is “forward” composition, which follows the order in which computation actually proceeds:

Language	Definition
OCaml	<code>let (\ -) f g x = g (f x)</code>
F#	<code>let (>>) f g x = g (f x)</code>

With forward composition, data flows from left to right, matching how we typically read code:

```
let iso1 = step1r |- step2r |- step3r
```

Partial Application Both composition examples above use **partial application**. Recall from the previous chapter that `((+) 1)` is a function that adds 1 to its argument. Partial application occurs when we do not pass all the arguments a function needs; the result is a function that requires the remaining arguments.

In the composition `step1r |- step2r |- step3r`, each `stepNr` function is partially applied. The composition operator `(|-)` takes two functions `f` and `g` and returns a new function that first applies `f`, then applies `g` to the result.

Power Function Now we define iterated function composition:

$$f^n(x) := \underbrace{(f \circ \cdots \circ f)}_{n \text{ times}}(x)$$

In OCaml, we first define the backward composition operator, then use it in `power`:

```
let (-|) f g x = f (g x)

let rec power f n =
  if n <= 0 then (fun x -> x) else f -| power f (n-1)
```

When `n <= 0`, we return the identity function. Otherwise, we compose `f` with `power f (n-1)`, which gives us one more application of `f`.

Numerical Derivative Using `power`, we can define a numerical approximation of the derivative:

```
let derivative dx f = fun x -> (f(x +. dx) -. f(x)) /. dx
```

This definition emphasizes that `derivative dx f` is itself a function of `x`. We can write it more concisely as:

```
let derivative dx f x = (f(x +. dx) -. f(x)) /. dx
```

Note that OCaml uses different operators for floating-point arithmetic. We have `(+)`: `int` → `int` → `int` for integers, so we cannot use `+` with floating-point numbers. Instead, operators followed by a dot work on `float` values: `+.`, `-.`, `*.`, `/..`

Computing Higher-Order Derivatives With `power` and `derivative`, we can easily compute higher-order derivatives:

```
let pi = 4.0 *. atan 1.0
let sin''' = (power (derivative 1e-5) 3) sin;;
sin''' pi
```

Here `sin'''` is the third derivative of sine. The result should be approximately $-\cos(\pi) = 1$ (with some numerical error due to the finite difference approximation).

3.2 Evaluation Rules (Reduction Semantics)

To understand how OCaml programs compute their results, we need to formalize the evaluation process. This section presents **reduction semantics**, which describes computation as a series of rewriting steps.

Expressions Programs consist of **expressions**. Here is the grammar of expressions for a simplified version of OCaml:

$a :=$ <ul style="list-style-type: none"> x $\mid \text{fun } x \rightarrow a$ $\mid a\ a$ $\mid C^0$ $\mid C^n(a, \dots, a)$ $\mid f^n$ $\mid \text{let } x = a \text{ in } a$ $\mid \text{match } a \text{ with}$ <ul style="list-style-type: none"> $p \rightarrow a \mid \dots \mid p \rightarrow a$ 	variables (defined) functions applications value constructors of arity 0 value constructors of arity n built-in values (primitives) of arity n name bindings (local definitions) pattern matching
$p :=$ <ul style="list-style-type: none"> x $\mid (p, \dots, p)$ $\mid C^0$ $\mid C^n(p, \dots, p)$ 	pattern variables tuple patterns variant patterns of arity 0 variant patterns of arity n

Arity means how many arguments something requires (and for tuples, the length of the tuple).

The fix Primitive To simplify our presentation of recursion, we use a primitive `fix` to define a limited form of `let rec`:

```
let rec f x = e1 in e2 ≡ let f = fix (fun f x -> e1) in e2
```

The `fix` primitive captures the essence of recursion: it takes a function that expects to receive itself as an argument and produces a fixed point—a function that, when called, behaves as if it had access to itself.

Values Expressions evaluate (i.e., compute) to **values**. Values are expressions that cannot be reduced further:

$$\begin{aligned}
v := & \quad \text{fun } x \rightarrow a && (\text{defined}) \text{ functions} \\
& | & C^n(v_1, \dots, v_n) & \text{constructed values} \\
& | & f^n v_1 \dots v_k & k < n \text{ (partially applied primitives)}
\end{aligned}$$

Note that functions are values: `fun x -> x + 1` is already fully evaluated. Partially applied primitives like `(+) 3` are also values—they are waiting for more arguments.

Substitution To **substitute** a value v for a variable x in expression a , we write $a[x := v]$. This notation means that every occurrence of x in a is replaced by v .

In the actual implementation, the value v is not duplicated in memory. Instead, OCaml uses references or closures to share the value efficiently.

Reduction Rules (Redexes) Reduction (i.e., computation) proceeds by finding reducible expressions called **redexes** and applying reduction rules. Here are the fundamental redexes:

Function application (beta reduction):

$$(\text{fun } x \rightarrow a) v \rightsquigarrow a[x := v]$$

When we apply a function to a value, we substitute the value for the parameter in the function body.

Let binding:

$$\text{let } x = v \text{ in } a \rightsquigarrow a[x := v]$$

A let binding with a value substitutes that value into the body.

Primitive application:

$$f^n v_1 \dots v_n \rightsquigarrow f(v_1, \dots, v_n)$$

When a primitive receives all its arguments, it computes the result. Here $f(v_1, \dots, v_n)$ denotes the actual result of the primitive operation.

Pattern matching with a variable pattern:

$$\text{match } v \text{ with } x \rightarrow a \mid \dots \rightsquigarrow a[x := v]$$

Pattern matching with a non-matching constructor:

$$\frac{C_1 \neq C_2}{\text{match } C_1^n(v_1, \dots, v_n) \text{ with } C_2^k(p_1, \dots, p_k) \rightarrow a \mid pm \rightsquigarrow \text{match } C_1^n(v_1, \dots, v_n) \text{ with } pm}$$

If the constructor does not match, we try the next pattern.

Pattern matching with a matching constructor:

`match C1n(v1, ..., vn) with C1n(x1, ..., xn) -> a | ... ~> a[x1 := v1; ...; xn := vn]`

If the constructor matches, we substitute all the bound values.

If $n = 0$, then $C_1^n(v_1, \dots, v_n)$ stands for simply C_1^0 , a constructor with no arguments. We omit the more complex cases of nested pattern matching.

Rule Variables In these rules, the metavariables have specific meanings: - x matches any expression or pattern variable - a, a_1, \dots, a_n match any expression - v, v_1, \dots, v_n match any value

To apply a rule, find substitutions for these metavariables that make the left-hand side match your expression. The right-hand side (with the same substitutions) is the reduced expression.

Evaluation Context Rules The rules above only apply when the arguments are already values. We also need rules that allow evaluation of subexpressions. If $a_i \rightsquigarrow a'_i$, then:

$$\begin{array}{ll} a_1 a_2 & \rightsquigarrow a'_1 a_2 \\ a_1 a_2 & \rightsquigarrow a_1 a'_2 \\ C^n(a_1, \dots, a_i, \dots, a_n) & \rightsquigarrow C^n(a_1, \dots, a'_i, \dots, a_n) \\ \text{let } x = a_1 \text{ in } a_2 & \rightsquigarrow \text{let } x = a'_1 \text{ in } a_2 \\ \text{match } a_1 \text{ with } pm & \rightsquigarrow \text{match } a'_1 \text{ with } pm \end{array}$$

These rules say that: - In an application, either the function or the argument can be evaluated (in arbitrary order) - In a constructor, any argument can be evaluated - In a let binding, the bound expression is evaluated before the body - In a match, the scrutinee is evaluated before matching

The fix Rule Finally, the rule for the `fix` primitive, which enables recursion:

$$\text{fix}^2 v_1 v_2 \rightsquigarrow v_1 (\text{fix}^2 v_1) v_2$$

Because `fix` is a binary primitive (arity 2), the expression $(\text{fix}^2 v_1)$ is already a value (a partially applied primitive). This means it will not be further evaluated until it is applied inside v_1 . This delayed evaluation is what makes recursion work without infinite loops.

Practice Exercise: Compute some simple programs by hand using these rules. For example, trace the evaluation of:

```
let double x = x + x in double 3
```

3.3 Symbolic Derivation Example

Let us see the reduction rules in action with a more complex example. Consider the symbolic expression evaluator from `Lec3.ml`:

```

type expression =
| Const of float
| Var of string
| Sum of expression * expression (* e1 + e2 *)
| Diff of expression * expression (* e1 - e2 *)
| Prod of expression * expression (* e1 * e2 *)
| Quot of expression * expression (* e1 / e2 *)

exception Unbound_variable of string

let rec eval env exp =
  match exp with
  | Const c -> c
  | Var v ->
    (try List.assoc v env with Not_found -> raise (Unbound_variable v))
  | Sum(f, g) -> eval env f +. eval env g
  | Diff(f, g) -> eval env f -. eval env g
  | Prod(f, g) -> eval env f *. eval env g
  | Quot(f, g) -> eval env f /. eval env g

```

We can also define symbolic differentiation:

```

let rec deriv exp dv =
  match exp with
  | Const c -> Const 0.0
  | Var v -> if v = dv then Const 1.0 else Const 0.0
  | Sum(f, g) -> Sum(deriv f dv, deriv g dv)
  | Diff(f, g) -> Diff(deriv f dv, deriv g dv)
  | Prod(f, g) -> Sum(Prod(deriv f dv, deriv g dv), Prod(deriv f dv, g))
  | Quot(f, g) -> Quot(Diff(Prod(deriv f dv, g), Prod(f, deriv g dv)),
                        Prod(g, g))

```

For convenience, let us define some operators and variables:

```

let x = Var "x"
let y = Var "y"
let (+:) f g = Sum (f, g)
let (-:) f g = Diff (f, g)
let (*:) f g = Prod (f, g)
let (/:) f g = Quot (f, g)
let (!:) i = Const i

```

Now consider evaluating the expression $3x + 2y + x^2$ at $x = 1, y = 2$:

```

let example = !:3.0 *: x +: !:2.0 *: y +: x *: x *: y
let env = ["x", 1.0; "y", 2.0]

```

When we trace the evaluation, we can see the recursive structure of the computation:

```

eval_1_2 <- 3.00 * x + 2.00 * y + x * x * y
  eval_1_2 <- x * x * y
    eval_1_2 <- y
    eval_1_2 --> 2.
      eval_1_2 <- x * x
        eval_1_2 <- x
        eval_1_2 --> 1.
          eval_1_2 <- x
          eval_1_2 --> 1.
            eval_1_2 --> 1.
            eval_1_2 --> 1.
            eval_1_2 --> 1.
          eval_1_2 --> 2.
        eval_1_2 <- 3.00 * x + 2.00 * y
          eval_1_2 <- 2.00 * y
            eval_1_2 <- y
            eval_1_2 --> 2.
              eval_1_2 <- 2.00
              eval_1_2 --> 2.
              eval_1_2 --> 4.
            eval_1_2 <- 3.00 * x
              eval_1_2 <- x
              eval_1_2 --> 1.
                eval_1_2 <- 3.00
                eval_1_2 --> 3.
                  eval_1_2 --> 3.
                  eval_1_2 --> 7.
                eval_1_2 --> 9.
              - : float = 9.

```

The indentation levels in this trace correspond to **stack frames**—the runtime structures that store the state of each function call. This brings us to an important optimization technique.

3.4 Tail Calls and Tail Recursion

Computers normally evaluate programs by creating **stack frames** on the call stack for each function call. The trace above illustrates this: each level of indentation represents a new stack frame.

What is a Tail Call? A **tail call** is a function call that is performed last when computing a function—there is nothing more to do after the call returns. For example, in:

```
let f x = g (x + 1)
```

The call to `g` is a tail call because after `g` returns, `f` immediately returns that value.

In contrast, in:

```
let f x = 1 + g x
```

The call to `g` is *not* a tail call because after `g` returns, we still need to add 1 to the result.

Tail Call Optimization Functional language compilers (including OCaml's) recognize tail calls and optimize them. Instead of creating a new stack frame, they reuse the current frame by performing a “jump” to the called function. This means tail calls use constant stack space.

Tail Recursive Functions A function is **tail recursive** if it calls itself (and any mutually recursive functions it depends on) only using tail calls.

Tail recursive functions often use special **accumulator** arguments that store intermediate computation results. In a non-tail-recursive function, these intermediate results would be values of subexpressions stored on the stack.

The key insight is that the accumulated result is computed in “reverse order”—while climbing up the recursion (making calls) rather than while descending (returning from calls).

Example: Counting Compare these two counting functions:

```
let rec count n =
  if n <= 0 then 0 else 1 + (count (n-1))
```

This is *not* tail recursive because after the recursive call returns, we still need to add 1.

```
let rec count_tcall acc n =
  if n <= 0 then acc else count_tcall (acc+1) (n-1)
```

This *is* tail recursive: the recursive call is the last thing the function does.

Example: Building Lists Let us see a more dramatic example:

```
let rec unfold n = if n <= 0 then [] else n :: unfold (n-1)
```

This function builds a list counting down from `n`. It is not tail recursive because after the recursive call, we must cons `n` onto the result.

```
# unfold 100000;;
- : int list = [100000; 99999; 99998; 99997; ...]
```

```
# unfold 1000000;;
Stack overflow during evaluation (looping recursion?).
```

With a million elements, we run out of stack space! Now consider the tail-recursive version:

```
let rec unfold_tcall acc n =
  if n <= 0 then acc else unfold_tcall (n::acc) (n-1)
```

The accumulator `acc` collects the list as we go. Note that the list is built in reverse order:

```
# unfold_tcall [] 100000;;
- : int list = [1; 2; 3; 4; 5; 6; 7; 8; 9; 10; 11; 12; ...]

# unfold_tcall [] 1000000;;
- : int list = [1; 2; 3; 4; 5; 6; 7; 8; 9; 10; 11; 12; ...]
```

The tail-recursive version handles a million elements with no problem.

A Challenge: Tree Depth Is it possible to find the depth of a tree using a tail-recursive function?

```
type btree = Tip | Node of int * btree * btree
```

The naive approach:

```
let rec depth tree = match tree with
| Tip -> 0
| Node(_, left, right) -> 1 + max (depth left) (depth right)
```

This is not tail recursive: after both recursive calls, we still need to compute `1 + max ...`. The challenge is that we have *two* recursive calls, and we cannot simply use an accumulator.

Note on Lazy Languages The issue of tail recursion is more complex for **lazy** programming languages like Haskell. In a lazy language, the `(:)` operation does not immediately evaluate its arguments, so building a list with `n :: unfold (n-1)` does not consume stack space in the same way.

3.5 First Encounter of Continuation Passing Style

We can solve the tree depth problem using **Continuation Passing Style (CPS)**. The key idea is to postpone doing actual work until the very last moment by passing around a “continuation”—a function that represents “what to do next.”

```
let rec depth_cps tree k = match tree with
| Tip -> k 0
| Node(_, left, right) ->
  depth_cps left (fun dleft ->
    depth_cps right (fun dright ->
```

```

k (1 + (max dleft dright)))
let depth tree = depth_cps tree (fun d -> d)

```

Let us understand how this works:

1. The function takes an extra parameter `k`, called the **continuation**. It represents what to do with the final result.
2. In the `Tip` case, we call the continuation with the depth 0.
3. In the `Node` case, we recursively compute the depth of the left subtree, passing a continuation that:
 - Receives the left depth `dleft`
 - Then recursively computes the depth of the right subtree, passing a continuation that:
 - Receives the right depth `dright`
 - Finally calls the original continuation with `1 + max dleft dright`
4. The wrapper function passes the identity function `fun d -> d` as the initial continuation.

The magic is that each recursive call is now a tail call! The “work” of computing `1 + max dleft dright` is captured in the continuation closures, which are allocated on the heap rather than the stack.

However, this does not completely solve the stack overflow problem—we are trading stack space for heap space (storing the continuation closures). For very deep trees, we might still run out of memory. True solutions involve trampolining or iterative approaches with explicit stacks.

CPS is a powerful technique that appears throughout functional programming. We will encounter it again when studying monads and advanced control flow.

3.6 Exercises

Exercise 1: By “traverse a tree” below we mean: write a function that takes a tree and returns a list of values in the nodes of the tree.

1. Write a function (of type `btree -> int list`) that traverses a binary tree in **prefix order**—first the value stored in a node, then values in all nodes to the left, then values in all nodes to the right.
2. Write a traversal in **infix order**—first values in all nodes to the left, then the value stored in the node, then values in all nodes to the right (so it is “left-to-right” order).
3. Write a traversal in **breadth-first order**—first values in shallower nodes before deeper nodes.

Exercise 2: Turn the function from Exercise 1 (prefix or infix traversal) into continuation passing style.

Exercise 3: Do the homework from the end of Chapter 2: write `btree_deriv_at` that takes a predicate over integers and a `btree`, and builds a `btree_deriv` whose “hole” is in the first position for which the predicate returns true.

Exercise 4: Write a function `simplify`: `expression -> expression` that simplifies symbolic expressions, so that for example the result of `simplify (deriv exp dv)` looks more like what a human would get computing the derivative of `exp` with respect to `dv`.

- Write a `simplify_once` function that performs a single step of simplification.
- Wrap it using a general `fixpoint` function that performs an operation until a **fixed point** is reached: given f and x , it computes $f^n(x)$ such that $f^n(x) = f^{n+1}(x)$.

Exercise 5: Write two sorting algorithms working on lists: merge sort and quicksort.

1. **Merge sort** splits the list roughly in half, sorts the parts recursively, and merges the sorted parts into the sorted result.
2. **Quicksort** splits the list into elements smaller than and greater than (or equal to) the first element, sorts the parts recursively, and concatenates them.

Chapter 4: Functions

Programming in untyped lambda-calculus

This chapter explores the theoretical foundations of functional programming through the untyped lambda-calculus. We begin with a review of computation by hand using our reduction semantics, then introduce the lambda-calculus notation and show how to encode fundamental data types—booleans, pairs, and natural numbers—using only functions. The chapter concludes with an examination of recursion through fixpoint combinators and practical considerations for avoiding infinite loops in eager evaluation.

References: - “Introduction to Lambda Calculus” by Henk Barendregt and Erik Barendsen - “Lecture Notes on the Lambda Calculus” by Peter Selinger

4.1 Review: Computation by Hand

Before diving into the lambda-calculus, let us work through a complete example of evaluation using the reduction rules from Chapter 3. This exercise reinforces our understanding of how computation proceeds and prepares us for the more abstract setting of lambda-calculus.

Recall that we use `fix` instead of `let rec` to simplify rules for recursion. Also remember our syntactic conventions: `fun x y -> e` stands for `fun x -> (fun y -> e)`, and so forth.

Consider the following recursive `length` function applied to a two-element list:

```
let rec fix f x = f (fix f) x

type int_list = Nil | Cons of int * int_list

let length =
  fix (fun f l ->
    match l with
    | Nil -> 0
    | Cons (x, xs) -> 1 + f xs) in
length (Cons (1, (Cons (2, Nil))))
```

Let us trace through this computation step by step. First, we eliminate the `let` binding:

$$\text{let } x = v \text{ in } a \Downarrow a[x := v]$$

This gives us:

```
fix (fun f l ->
  match l with
  | Nil -> 0
  | Cons (x, xs) -> 1 + f xs) (Cons (1, (Cons (2, Nil))))
```

Next, we apply the `fix` rule:

$$\text{fix}^2 v_1 v_2 \Downarrow v_1 (\text{fix}^2 v_1) v_2$$

This unfolds to:

```
(fun f l ->
  match l with
  | Nil -> 0
  | Cons (x, xs) -> 1 + f xs)
(fix (fun f l ->
  match l with
  | Nil -> 0
  | Cons (x, xs) -> 1 + f xs))
(Cons (1, (Cons (2, Nil))))
```

Function application reduces according to:

$$(\text{fun } x \rightarrow a) v \rightsquigarrow a[x := v]$$

After substituting both `f` and `l`, we get:

```
(match Cons (1, (Cons (2, Nil))) with
  | Nil -> 0
  | Cons (x, xs) -> 1 + (fix (fun f l ->
    match l with
      | Nil -> 0
      | Cons (x, xs) -> 1 + f xs)) xs)
```

Pattern matching against a non-matching constructor moves to the next branch:

```
match  $C_1^n(v_1, \dots, v_n)$  with
 $C_2^n(p_1, \dots, p_k) \rightarrow a \mid pm \Downarrow \text{match } C_1^n(v_1, \dots, v_n) \text{ with } pm$ 
```

Pattern matching against a matching constructor performs substitution:

```
match  $C_1^n(v_1, \dots, v_n)$  with
 $C_1^n(x_1, \dots, x_n) \rightarrow a \mid \dots \Downarrow a[x_1 := v_1; \dots; x_n := v_n]$ 
```

After matching and substitution:

```
1 + (fix (fun f l ->
  match l with
    | Nil -> 0
    | Cons (x, xs) -> 1 + f xs)) (Cons (2, Nil))
```

Continuing the evaluation, we apply `fix` again and work through the pattern match for `Cons (2, Nil)`, eventually reaching:

```
1 + (1 + (fix (fun f l ->
  match l with
    | Nil -> 0
    | Cons (x, xs) -> 1 + f xs)) Nil)
```

One more unfolding and pattern match against `Nil` gives:

```
1 + (1 + 0)
```

Finally, applying the built-in addition:

$$f^n v_1 \dots v_n \Downarrow f(v_1, \dots, v_n)$$

We obtain the result: 2.

4.2 Language and Rules of the Untyped Lambda-Calculus

The lambda-calculus, introduced by Alonzo Church in the 1930s, is a minimal formal system for expressing computation. To work with it, we first simplify our language:

1. **Forget about types.** In pure lambda-calculus, there is no type system constraining which terms can be combined.
2. **Introduce notation.** We write $\lambda x.a$ for `fun x -> a`, and $\lambda xy.a$ for `fun x y -> a`, and so forth.
3. **Reduce to essentials.** We keep only functions (lambda abstractions) and variables—no constructors, no built-in primitives.

The core reduction rule of lambda-calculus is called **β -reduction**:

$$(\text{fun } x \rightarrow a_1) a_2 \rightsquigarrow a_1[x := a_2]$$

Note that this rule is more general than the one we use for OCaml evaluation. In our OCaml semantics, we require the argument to be a value: $(\text{fun } x \rightarrow a) v \rightsquigarrow a[x := v]$. The general β -reduction rule allows substituting any expression, not just values.

Lambda-calculus also uses **α -conversion** (bound variable renaming), or equivalent techniques, to avoid **variable capture**—the unintended binding of free variables during substitution. We will explore β -reduction further in the chapter on laziness.

Why is β -reduction more general than our evaluation rule? Consider the expression $(\lambda x.x)((\lambda y.y)z)$. With β -reduction, we could reduce the outer application first, obtaining $((\lambda y.y)z)$. Our evaluation rule would require first reducing the argument to a value.

4.3 Booleans

Alonzo Church introduced lambda-calculus to encode logic. There are multiple ways to encode various sorts of data in lambda-calculus, though not all of them work well in a typed setting—the straightforward encode/decode functions may not type-check.

The **Church encoding** of booleans represents truth values as selector functions:

- **True** selects the first argument: `c_true = $\lambda xy.x$`
- **False** selects the second argument: `c_false = $\lambda xy.y$`

In OCaml syntax:

```
let c_true = fun x y -> x (* "True" is projection on the first argument *)
let c_false = fun x y -> y (* And "false" on the second argument *)
```

Logical conjunction can be defined as:

$$c_{\text{and}} = \lambda xy.x y c_{\text{false}}$$

The logic is: if x is true, return y (so the result is true only if both are true); if x is false, return false immediately.

```
let c_and = fun x y -> x y c_false (* If one is false, then return false *)
```

Let us verify this works. For c_{and} c_{true} c_{false} :

$$(\lambda xy.x y c_{\text{false}}) (\lambda xy.x) (\lambda xy.x)$$

reduces to:

$$(\lambda xy.x) (\lambda xy.x) c_{\text{false}}$$

which gives us $\lambda xy.x = c_{\text{true}}$. For any other combination involving c_{false} , the result is c_{false} .

To verify our encodings in OCaml, we need encode and decode functions:

```
let encode_bool b = if b then c_true else c_false
let decode_bool c = c true false (* Test the functions in the toplevel *)
```

Exercise: Define c_{or} and c_{not} yourself!

4.4 If-then-else and Pairs

From now on, we will use OCaml syntax for our lambda-calculus programs. An important observation is that our encoded booleans already implement conditional selection:

```
let if_then_else = fun b -> b (* Booleans select the argument! *)
```

Since c_{true} returns its first argument and c_{false} returns its second, `if_then_else b then_branch else_branch` simply applies `b` to the two branches. Remember to play with these functions in the toplevel to build intuition.

Pairs Pairs (ordered tuples of two elements) can be encoded similarly:

```
let c_pair m n = fun x -> x m n (* We couple things *)
let c_first = fun p -> p c_true (* by passing them together *)
let c_second = fun p -> p c_false (* Check that it works! *)
```

A pair is a function that, when given a selector, applies that selector to both components. To extract the first component, we pass c_{true} (which selects the first argument); to extract the second, we pass c_{false} .

For verification:

```

let encode_pair enc_fst enc_snd (a, b) =
  c_pair (enc_fst a) (enc_snd b)
let decode_pair de_fst de_snd c = c (fun x y -> de_fst x, de_snd y)
let decode_bool_pair c = decode_pair decode_bool decode_bool c

```

We can define larger tuples in the same manner:

```
let c_triple l m n = fun x -> x l m n
```

4.5 Pair-Encoded Natural Numbers

Our first encoding of natural numbers uses nested pairs. The representation is based on the depth of nested pairs whose rightmost leaf is the identity function $\lambda x.x$ and whose left elements are `c_false`.

```

let pn0 = fun x -> x          (* Start with the identity function *)
let pn_succ n = c_pair c_false n (* Stack another pair *)

let pn_pred = fun x -> x c_false (* Extract the nested number *)
let pn_is_zero = fun x -> x c_true (* Check if it's the base case *)

```

The number 0 is represented as the identity function. The number 1 is `c_pair c_false pn0`, the number 2 is `c_pair c_false (c_pair c_false pn0)`, and so on. The `pn_is_zero` function works because: - For `pn0`, applying it to `c_true` gives `c_true` (since `pn0` is the identity). - For any successor, applying `c_pair c_false n` to `c_true` applies the pair to `c_true`, which selects `c_false`.

We program in untyped lambda-calculus as an exercise, and we need encoding/decoding to verify our exercises. Using `Obj.magic` to bypass the type system for encoding/decoding is “fair game”:

```

let rec encode_pnat n =                      (* We use Obj.magic to forget types *)
  if n <= 0 then Obj.magic pn0
  else pn_succ (Obj.magic (encode_pnat (n-1))) (* Disregarding types, *)
let rec decode_pnat pn =                      (* these functions are straightforward! *)
  if decode_bool (pn_is_zero pn) then 0
  else 1 + decode_pnat (pn_pred (Obj.magic pn))

```

4.6 Church Numerals

Do you remember our function `power f n` from Chapter 3? We will use a similar idea for a different representation of numbers. **Church numerals** represent a natural number n as a function that applies its first argument n times to its second argument:

```

let cn0 = fun f x -> x          (* The same as c_false *)
let cn1 = fun f x -> f x        (* Behaves like identity when f = id *)
let cn2 = fun f x -> f (f x)
let cn3 = fun f x -> f (f (f x))

```

This is the original Alonzo Church encoding. The number n is represented as $\lambda f x. f^n(x)$, where f^n denotes n -fold composition.

The successor function adds one more application of f :

```
let cn_succ = fun n f x -> f (n f x)
```

Exercise: Define addition, multiplication, comparing to zero, and the predecessor function “-1” for Church numerals.

It turns out even Alonzo Church could not define predecessor right away! His student Stephen Kleene eventually found it. Try to make some progress before looking at the solution below.

```
let (-|) f g x = f (g x) (* Backward composition operator *)

let rec encode_cnat n f =
  if n <= 0 then (fun x -> x) else f -| encode_cnat (n-1) f
let decode_cnat n = n ((+) 1) 0
let cn7 f x = encode_cnat 7 f x (* We need to eta-expand these definitions *)
let cn13 f x = encode_cnat 13 f x (* for type-system reasons *)
(* (because OCaml allows side-effects) *)
let cn_add = fun n m f x -> n f (m f x) (* Put n of f in front *)
let cn_mult = fun n m f -> n (m f) (* Repeat n times *)
(* putting m of f in front *)
let cn_prev n =
  fun f x ->
    n
      (fun g v -> v (g f)) (* This is the "Church numeral signature" *)
      (* The only thing we have is an n-step loop *)
      (fun z -> x) (* We need sth that operates on f *)
      (fun z -> z) (* We need to ignore the innermost step *)
(* We've built a "machine" not results -- start the machine *)
```

The predecessor function is ingenious. It builds up a chain of functions that, when “started” with the identity, yields $n - 1$ applications of f . The key insight is to delay the actual application of f and skip the first one.

`cn_is_zero` is left as an exercise.

Tracing `cn_prev cn3` Let us trace through `decode_cnat (cn_prev cn3)`:

```
↓
(cn_prev cn3) ((+) 1) 0
↓
(fun f x ->
  cn3
    (fun g v -> v (g f))
```

```
(fun z -> x)
  (fun z -> z)) ((+) 1) 0
```

↓

```
((fun f x -> f (f (f x)))
  (fun g v -> v (g ((+) 1)))
  (fun z -> 0)
  (fun z -> z))
```

↓

```
((fun g v -> v (g ((+) 1)))
  ((fun g v -> v (g ((+) 1)))
    ((fun g v -> v (g ((+) 1)))
      (fun z -> 0))))
  (fun z -> z))
```

↓

```
((fun z -> z)
  (((fun g v -> v (g ((+) 1)))
    ((fun g v -> v (g ((+) 1)))
      (fun z -> 0)))) ((+) 1)))
```

↓

```
(fun g v -> v (g ((+) 1)))
  ((fun g v -> v (g ((+) 1)))
    (fun z -> 0)) ((+) 1))
```

↓

```
((+) 1) ((fun g v -> v (g ((+) 1)))
  (fun z -> 0)) ((+) 1))
```

↓

```
((+) 1) (((+) 1) ((fun z -> 0) ((+) 1)))
```

↓

```
((+) 1) (((+) 1) (0))
```

↓

((+) 1) 1

↓

2

4.7 Recursion: Fixpoint Combinators

In lambda-calculus, recursion is achieved through **fixpoint combinators**—lambda terms that compute fixed points of functions.

Turing's Fixpoint Combinator

$$\Theta = (\lambda xy.y(x x y))(\lambda xy.y(x x y))$$

Let us verify it computes fixed points. Define $N = \Theta F$:

$$\begin{aligned} N &= \Theta F \\ &= (\lambda xy.y(x x y))(\lambda xy.y(x x y))F \\ &\stackrel{=} {\rightarrow} F((\lambda xy.y(x x y))(\lambda xy.y(x x y))F) \\ &= F(\Theta F) = FN \end{aligned}$$

So $N = FN$, meaning N is a fixed point of F .

Curry's Fixpoint Combinator (Y Combinator)

$$Y = \lambda f.(\lambda x.f(x x))(\lambda x.f(x x))$$

$$\begin{aligned} N &= YF \\ &= (\lambda f.(\lambda x.f(x x))(\lambda x.f(x x)))F \\ &\stackrel{=} {\rightarrow} (\lambda x.F(x x))(\lambda x.F(x x)) \\ &\stackrel{=} {\rightarrow} F((\lambda x.F(x x))(\lambda x.F(x x))) \\ &\stackrel{=} {\leftarrow} F((\lambda f.(\lambda x.f(x x))(\lambda x.f(x x)))F) \\ &= F(YF) = FN \end{aligned}$$

Call-by-Value Fixpoint Combinator

$$\text{fix} = \lambda f'.(\lambda fx.f'(f f)x)(\lambda fx.f'(f f)x)$$

$$\begin{aligned}
N &= \text{fix } F \\
&= (\lambda f'.(\lambda fx.f' (f f) x) (\lambda fx.f' (f f) x)) F \\
&\Rightarrow (\lambda fx.F (f f) x) (\lambda fx.F (f f) x) \\
&\Rightarrow \lambda x.F ((\lambda fx.F (f f) x) (\lambda fx.F (f f) x)) x \\
&\Leftarrow \lambda x.F ((\lambda f'.(\lambda fx.f' (f f) x) (\lambda fx.f' (f f) x)) F) x \\
&= \lambda x.F (\text{fix } F) x = \lambda x.F N x \\
&=_\eta F N
\end{aligned}$$

The Problem with the First Two Combinators What is the problem with Turing's and Curry's combinators? Consider what happens when we try to evaluate ΘF :

$$\begin{aligned}
\Theta F &\rightsquigarrow F ((\lambda xy.y (x x y)) (\lambda xy.y (x x y)) F) \\
&\rightsquigarrow F (F ((\lambda xy.y (x x y)) (\lambda xy.y (x x y)) F)) \\
&\rightsquigarrow F (F (F ((\lambda xy.y (x x y)) (\lambda xy.y (x x y)) F))) \\
&\rightsquigarrow \dots
\end{aligned}$$

Recall the distinction between *expressions* and *values* from Chapter 3 on Computation. The reduction rule for lambda-calculus is meant to determine which expressions are considered “equal”—it is highly *non-deterministic*, while on a computer, computation needs to go one way or another.

Using the general reduction rule of lambda-calculus, for a recursive definition, it is always possible to find an infinite reduction sequence. This means a naive lambda-calculus compiler could legitimately generate infinite loops for all recursive definitions!

Therefore, we need more specific rules. Most languages use **call-by-value** (also called **eager** evaluation):

$$(\text{fun } x \rightarrow a) v \rightsquigarrow a[x := v]$$

The program *eagerly* computes arguments before starting to compute the function body. This is exactly the rule we introduced in the Computation chapter.

Call-by-Value Fixpoint Combinator in Action What happens with the call-by-value fixpoint combinator?

$$\begin{aligned}
\text{fix } F &\rightsquigarrow (\lambda fx.F (f f) x) (\lambda fx.F (f f) x) \\
&\rightsquigarrow \lambda x.F ((\lambda fx.F (f f) x) (\lambda fx.F (f f) x)) x
\end{aligned}$$

The computation stops because we use the rule $(\text{fun } x \rightarrow a) v \rightsquigarrow a[x := v]$ rather than $(\text{fun } x \rightarrow a_1) a_2 \rightsquigarrow a_1[x := a_2]$. The expression inside the lambda is not evaluated until the function is applied.

Let us compute the function on some input:

$$\begin{aligned} \text{fix } F v &\rightsquigarrow (\lambda f x. F(f f) x) (\lambda f x. F(f f) x) v \\ &\rightsquigarrow (\lambda x. F((\lambda f x. F(f f) x) (\lambda f x. F(f f) x))) x) v \\ &\rightsquigarrow F((\lambda f x. F(f f) x) (\lambda f x. F(f f) x)) v \\ &\rightsquigarrow F(\lambda x. F((\lambda f x. F(f f) x) (\lambda f x. F(f f) x))) x) v \\ &\rightsquigarrow \text{depends on } F \end{aligned}$$

Why “Fixpoint”? If you examine our derivations, you will see they establish $x = f(x)$. Such values x are called **fixpoints** of f . An arithmetic function can have several fixpoints—for example, $f(x) = x^2$ has fixpoints 0 and 1—or no fixpoints, such as $f(x) = x + 1$.

When you define a function (or another object) by recursion, it has similar meaning: the name appears on both sides of the equality. In lambda-calculus, functions like Θ and \mathbf{Y} take *any* function as an argument and return its fixpoint.

We turn a specification of a recursive object into a definition by solving it with respect to the recurring name: deriving $x = f(x)$ where x is the recurring name. We then have $x = \text{fix}(f)$.

Deriving Factorial Let us walk through this for the factorial function. We omit the prefix `cn_` (could be `pn_` if using pair-encoded numbers) and shorten `if_then_else` to `if_t_e`:

```
fact n = if_t_e (is_zero n) cn1 (mult n (fact (pred n)))
fact = λn.if_t_e (is_zero n) cn1 (mult n (fact (pred n)))
fact = (λfn.if_t_e (is_zero n) cn1 (mult n (f (pred n)))) fact
fact = fix (λfn.if_t_e (is_zero n) cn1 (mult n (f (pred n))))
```

The last line is a valid definition: we simply give a name to a *ground* (also called *closed*) expression—one with no free variables.

Exercise: Compute `fact cn2`.

Exercise: What does `fix (fun x → cn_succ x)` mean?

4.8 Encoding Lists and Trees

A **list** is either empty (often called `Empty` or `Nil`) or consists of an element followed by another list (the “tail”), called `Cons`.

Define: - **nil** = $\lambda xy.y$ - **cons** $H T = \lambda xy.x H T$

To add numbers stored inside a list:

```
addlist  $l = l (\lambda ht.cn\_add h (addlist t)) cn0$ 
```

To make a proper definition, we apply **fix** to the solution of the above equation:

```
addlist = fix ( $\lambda fl.l (\lambda ht.cn\_add h (f t)) cn0$ )
```

For **trees**, let us use a different form of binary trees: instead of keeping elements in inner nodes, we keep elements in leaves.

Define: - **leaf** $n = \lambda xy.x n$ - **node** $L R = \lambda xy.y L R$

To add numbers stored inside a tree:

```
addtree  $t = t (\lambda n.n) (\lambda lr.cn\_add (addtree l) (addtree r))$ 
```

And in solved form:

```
addtree = fix ( $\lambda ft.t (\lambda n.n) (\lambda lr.cn\_add (f l) (f r))$ )
let rec fix f x = f (fix f) x
let nil = fun x y -> y
let cons h t = fun x y -> x h t
let addlist l =
  fix (fun f l -> l (fun h t -> cn_add h (f t)) cn0) l
;;
decode_cnat
  (addlist (cons cn1 (cons cn2 (cons cn7 nil))));;
let leaf n = fun x y -> x n
let node l r = fun x y -> y l r
let addtree t =
  fix (fun f t ->
    t (fun n -> n) (fun l r -> cn_add (f l) (f r)))
  ) t
;;
decode_cnat
  (addtree (node (node (leaf cn3) (leaf cn7))
    (leaf cn1))));;
```

The General Pattern Observe a regularity: when we encode a variant type with n variants, for each variant we define a function that takes n arguments.

If the k th variant C_k has m_k parameters, then the function c_k that encodes it has the form:

$$C_k(v_1, \dots, v_{m_k}) \sim c_k v_1 \dots v_{m_k} = \lambda x_1 \dots x_n. x_k v_1 \dots v_{m_k}$$

The encoded variants serve as shallow pattern matching with guaranteed exhaustiveness: the k th argument corresponds to the k th branch of pattern matching.

4.9 Looping Recursion

Let us return to pair-encoded numbers and define addition:

```
let pn_add m n =
  fix (fun f m n ->
    if_then_else (pn_is_zero m)
      n (pn_succ (f (pn_pred m) n))
    ) m n;;
decode_pnat (pn_add pn3 pn3);;
```

Oops... OCaml says: Stack overflow during evaluation (looping recursion?).

What went wrong? Nothing as far as lambda-calculus is concerned. But OCaml (and F#) always compute arguments before calling a function. By definition of `fix`, `f` corresponds to recursively calling `pn_add`. Therefore, `(pn_succ (f (pn_pred m) n))` will be evaluated regardless of what `(pn_is_zero m)` returns!

Why do `addlist` and `addtree` work? Because their recursive calls are “guarded” by corresponding `fun`. What is inside of `fun` is not computed immediately—only when the function is applied to argument(s).

To avoid looping recursion, you need to guard all recursive calls. Besides putting them inside `fun`, in OCaml or F# you can also put them in branches of a `match` clause, as long as one of the branches does not have unguarded recursive calls.

The trick for functions like `if_then_else` is to guard their arguments with `fun x ->`, where `x` is not used, and apply the *result* of `if_then_else` to some dummy value:

```
let id x = x
let rec fix f x = f (fix f) x
let pn1 x = pn_succ pn0 x
let pn2 x = pn_succ pn1 x
let pn3 x = pn_succ pn2 x
let pn7 x = encode_pnat 7 x
let pn_add m n =
```

```

fix (fun f m n ->
  (if_then_else (pn_is_zero m)
    (fun x -> n) (fun x -> pn_succ (f (pn_pred m) n)))
  id
) m n;;
decode_pnat (pn_add pn3 pn3);;
decode_pnat (pn_add pn3 pn7);;

```

In OCaml or F# we would typically guard by `fun () ->` and then apply to `()`, but we do not have datatypes like `unit` in pure lambda-calculus.

4.10 Exercises

Exercise 1: Define (implement) and test on a couple of examples functions corresponding to or computing:

1. `c_or` and `c_not`;
2. exponentiation for Church numerals;
3. is-zero predicate for Church numerals;
4. even-number predicate for Church numerals;
5. multiplication for pair-encoded natural numbers;
6. factorial $n!$ for pair-encoded natural numbers;
7. the length of a list (in Church numerals);
8. `cn_max` – maximum of two Church numerals;
9. the depth of a tree (in Church numerals).

Exercise 2: Construct lambda-terms m_0, m_1, \dots such that for all n one has:

$$\begin{aligned} m_0 &= x \\ m_{n+1} &= m_{n+2} m_n \end{aligned}$$

(where equality is after performing β -reductions).

Exercise 3: Representing side-effects as an explicitly “passed around” state value, write (higher-order) functions that represent the imperative constructs:

1. `for...to...`
2. `for...downto...`
3. `while...do...`
4. `do...while...`
5. `repeat...until...`

Rather than writing a lambda-term using the encodings that we have learnt, just implement the functions in OCaml / F#, using built-in `int` and `bool` types. You can use `let rec` instead of `fix`.

- For example, in exercise (a), write a function `let rec for_to f beg_i end_i s = ...` where `f` takes arguments `i` ranging from `beg_i` to `end_i`,

state **s** at given step, and returns state **s** at next step; the **for_to** function returns the state after the last step.

- And in exercise (c), write a function **let rec while_do p f s = ...** where both **p** and **f** take state **s** at given step, and if **p s** returns true, then **f s** is computed to obtain state at next step; the **while_do** function returns the state after the last step.

Do not use the imperative features of OCaml and F#!

Although we will not cover imperative features in this course, it is instructive to see the implementation using them, to better understand what is actually required of a solution to Exercise 3:

```
(* (a) *)
let for_to f beg_i end_i s =
  let s = ref s in
  for i = beg_i to end_i do
    s := f i !s
  done;
!s

(* (b) *)
let for_downto f beg_i end_i s =
  let s = ref s in
  for i = beg_i downto end_i do
    s := f i !s
  done;
!s

(* (c) *)
let while_do p f s =
  let s = ref s in
  while p !s do
    s := f !s
  done;
!s

(* (d) *)
let do_while p f s =
  let s = ref (f s) in
  while p !s do
    s := f !s
  done;
!s

(* (e) *)
let repeat_until p f s =
```

```

let s = ref (f s) in
  while not (p !s) do
    s := f !s
  done;
!s

```

Chapter 5: Polymorphism and Abstract Data Types

This chapter explores how OCaml’s type system supports generic programming through parametric polymorphism, and how abstract data types provide clean interfaces for data structures. We examine the formal mechanics of type inference and then apply these concepts to build progressively more sophisticated implementations of the map (dictionary) data structure, culminating in red-black trees.

5.1 Type Inference

We have seen the rules that govern the assignment of types to expressions, but how does OCaml guess what types to use, and when no correct types exist? The answer is that it solves equations.

5.1.1 Variables: Unknowns and Parameters Variables in type inference play two distinct roles: they can be *unknowns* (standing for a specific but not-yet-determined type) or *parameters* (standing for any type whatsoever).

Consider the following example:

```

# let f = List.hd;;
val f : 'a list -> 'a = <fun>

```

Here '`'a`' is a parameter: it can become any type. Mathematically we write: $f : \forall \alpha. \alpha \text{ list} \rightarrow \alpha$ – the quantified type is called a *type scheme*.

In contrast:

```

# let x = ref [];;
val x : '_weak1 list ref = {contents = []}

```

Here '`'_a`' is an unknown. It stands for a particular type like `float` or `int` → `int`, but OCaml just does not yet know which type. OCaml only reports unknowns like '`'_a`' in inferred types for reasons related to mutable state (the “value restriction”), which are not central to functional programming.

When unknowns appear in inferred types against our expectations, η -*expansion* may help: writing `let f x = expr x` instead of `let f = expr`. For example:

```

# let f = List.append [];;
val f : '_weak2 list -> '_weak2 list = <fun>
# let f l = List.append [] l;;
val f : 'a list -> 'a list = <fun>

```

In the second definition, the eta-expanded form allows full generalization.

5.1.2 Type Environments A *type environment* specifies what names (corresponding to parameters and definitions) are available for an expression because they were introduced above it, and it specifies their types.

5.1.3 Solving Type Equations Type inference solves equations over unknowns. The central question is: “What has to hold so that $e : \tau$ in type environment Γ ?”

The process works as follows:

- If, for example, $f : \forall \alpha. \alpha \text{ list} \rightarrow \alpha \in \Gamma$, then for $f : \tau$ we introduce $\gamma \text{ list} \rightarrow \gamma = \tau$ for some fresh unknown γ .
- For function application $e_1 e_2 : \tau$, we introduce $\beta = \tau$ and ask for $e_1 : \gamma \rightarrow \beta$ and $e_2 : \gamma$, for some fresh unknowns β, γ .
- For a function $\text{fun } x \rightarrow e : \tau$, we introduce $\beta \rightarrow \gamma = \tau$ and ask for $e : \gamma$ in environment $\{x : \beta\} \cup \Gamma$, for some fresh unknowns β, γ .
- The case $\text{let } x = e_1 \text{ in } e_2 : \tau$ is different. One approach is to *first* solve the equations that we get by asking for $e_1 : \beta$, for some fresh unknown β . Let us say a solution $\beta = \tau_\beta$ has been found, $\alpha_1 \dots \alpha_n, \beta_1 \dots \beta_m$ are the remaining unknowns in τ_β , and $\alpha_1 \dots \alpha_n$ are all that do not appear in Γ . Then we ask for $e_2 : \tau$ in environment $\{x : \forall \alpha_1 \dots \alpha_n. \tau_\beta\} \cup \Gamma$.
- Remember that whenever we establish a solution $\beta = \tau_\beta$ to an unknown β , it takes effect everywhere!
- To find a type for e (in environment Γ), we pick a fresh unknown β and ask for $e : \beta$ (in Γ).

5.1.4 Polymorphism The “top-level” definitions for which the system infers types with variables are called *polymorphic*, which informally means “working with different shapes of data”. This kind of polymorphism is called *parametric polymorphism*, since the types have parameters. A different kind of polymorphism is provided by object-oriented programming languages (sometimes called *subtype polymorphism* or *ad-hoc polymorphism*).

5.2 Parametric Types

Polymorphic functions shine when used with polymorphic data types. Consider:

```
type 'a my_list = Empty | Cons of 'a * 'a my_list
```

We define lists that can store elements of any type ' a '. Now:

```
# let tail l =
  match l with
```

```

| Empty -> invalid_arg "tail"
| Cons (_, tl) -> tl;;
val tail : 'a my_list -> 'a my_list = <fun>

```

This is a polymorphic function: it works for lists with elements of any type.

A *parametric type* like `'a my_list` is *not* itself a data type but a family of data types: `bool my_list`, `int my_list`, etc. are different types. We say that the type `int my_list` *instantiates* the parametric type `'a my_list`.

5.2.1 Multiple Type Parameters In OCaml, the syntax is a bit confusing: type parameters precede the type name. For example:

```
type ('a, 'b) choice = Left of 'a | Right of 'b
```

This type has two parameters. Mathematically we would write $\text{choice}(\alpha, \beta)$.

Functions do not have to be polymorphic:

```

# let get_int c =
  match c with
  | Left i -> i
  | Right b -> if b then 1 else 0;;
val get_int : (int, bool) choice -> int = <fun>

```

5.2.2 Syntax in Other Languages In F#, we provide parameters (when more than one) after the type name:

```
type choice<'a,'b> = Left of 'a | Right of 'b
```

In Haskell, we provide type parameters similarly to function arguments:

```
data Choice a b = Left a | Right b
```

5.3 Type Inference, Formally

A statement that an expression has a type in an environment is called a *type judgement*. For environment $\Gamma = \{x : \forall \alpha_1 \dots \alpha_n. \tau_x; \dots\}$, expression e and type τ we write:

$$\Gamma \vdash e : \tau$$

We will derive the equations in one go using $\llbracket \cdot \rrbracket$, to be solved later. Besides equations we will need to manage introduced variables, using existential quantification.

For local definitions we require remembering what constraints should hold when the definition is used. Therefore we extend *type schemes* in the environment to: $\Gamma = \{x : \forall \beta_1 \dots \beta_m [\exists \alpha_1 \dots \alpha_n. D]. \tau_x; \dots\}$ where D are equations – keeping the

variables $\alpha_1 \dots \alpha_n$ introduced while deriving D in front. A simpler form would be enough: $\Gamma = \{x : \forall \beta[\exists \alpha_1 \dots \alpha_n.D].\beta; \dots\}$

The formal constraint generation rules are:

$$[\![\Gamma \vdash x : \tau]\!] = \exists \overline{\beta' \alpha'}. (D[\overline{\beta \alpha} := \overline{\beta' \alpha'}] \wedge \tau_x[\overline{\beta \alpha} := \overline{\beta' \alpha'}] \doteq \tau)$$

where $\Gamma(x) = \forall \overline{\beta}[\exists \overline{\alpha}.D].\tau_x, \overline{\beta' \alpha'} \# \text{FV}(\Gamma, \tau)$

$$[\![\Gamma \vdash \mathbf{fun} \ x \rightarrow e : \tau]\!] = \exists \alpha_1 \alpha_2. ([\![\Gamma\{x : \alpha_1\} \vdash e : \alpha_2]\!] \wedge \alpha_1 \rightarrow \alpha_2 \doteq \tau)$$

where $\alpha_1 \alpha_2 \# \text{FV}(\Gamma, \tau)$

$$[\![\Gamma \vdash e_1 \ e_2 : \tau]\!] = \exists \alpha. ([\![\Gamma \vdash e_1 : \alpha \rightarrow \tau]\!] \wedge [\![\Gamma \vdash e_2 : \alpha]\!]), \alpha \# \text{FV}(\Gamma, \tau)$$

$$[\![\Gamma \vdash K \ e_1 \dots e_n : \tau]\!] = \exists \overline{\alpha'}. (\bigwedge_i [\![\Gamma \vdash e_i : \tau_i[\overline{\alpha} := \overline{\alpha'}]]\!] \wedge \varepsilon(\overline{\alpha'}) \doteq \tau)$$

where $K : \forall \overline{\alpha}. \tau_1 \times \dots \times \tau_n \rightarrow \varepsilon(\overline{\alpha}), \overline{\alpha'} \# \text{FV}(\Gamma, \tau)$

For let-expressions:

$$[\![\Gamma \vdash \mathbf{let} \ x = e_1 \ \mathbf{in} \ e_2 : \tau]\!] = (\exists \beta.C) \wedge [\![\Gamma\{x : \forall \beta[C].\beta\} \vdash e_2 : \tau]\!]$$

where $C = [\![\Gamma \vdash e_1 : \beta]\!]$

For recursive let-expressions:

$$[\![\Gamma \vdash \mathbf{letrec} \ x = e_1 \ \mathbf{in} \ e_2 : \tau]\!] = (\exists \beta.C) \wedge [\![\Gamma\{x : \forall \beta[C].\beta\} \vdash e_2 : \tau]\!]$$

where $C = [\![\Gamma\{x : \beta\} \vdash e_1 : \beta]\!]$

For match expressions:

$$[\![\Gamma \vdash \mathbf{match} \ e_v \ \mathbf{with} \ \bar{c} : \tau]\!] = \exists \alpha_v. [\![\Gamma \vdash e_v : \alpha_v]\!] \bigwedge_i [\![\Gamma \vdash p_i.e_i : \alpha_v \rightarrow \tau]\!]$$

where $\bar{c} = p_1.e_1 | \dots | p_n.e_n, \alpha_v \# \text{FV}(\Gamma, \tau)$

For pattern clauses:

$$[\![\Gamma, \Sigma \vdash p.e : \tau_1 \rightarrow \tau_2]\!] = [\![\Sigma \vdash p \downarrow \tau_1]\!] \wedge \exists \overline{\beta}. [\![\Gamma \Gamma' \vdash e : \tau_2]\!]$$

where $\exists \overline{\beta} \Gamma'$ is $[\![\Sigma \vdash p \uparrow \tau_1]\!], \overline{\beta} \# \text{FV}(\Gamma, \tau_2)$

The notation $\llbracket \Sigma \vdash p \downarrow \tau_1 \rrbracket$ derives constraints on the type of the matched value, while $\llbracket \Sigma \vdash p \uparrow \tau_1 \rrbracket$ derives the environment for pattern variables.

By $\overline{\alpha}$ or $\overline{\alpha_i}$ we denote a sequence of some length: $\alpha_1 \dots \alpha_n$. By $\bigwedge_i \varphi_i$ we denote a conjunction of $\overline{\varphi_i}$: $\varphi_1 \wedge \dots \wedge \varphi_n$.

5.3.1 Polymorphic Recursion Note the limited polymorphism of `let rec f = ...` – we cannot use `f` polymorphically in its definition. In modern OCaml we can bypass the problem if we provide the type of `f` upfront:

```
let rec f : 'a. 'a -> 'a list = ...
```

where `'a. 'a -> 'a list` stands for $\forall \alpha. \alpha \rightarrow \alpha$ list.

Using the recursively defined function with different types in its definition is called *polymorphic recursion*. It is most useful together with *irregular recursive datatypes* where the recursive use has different type arguments than the actual parameters.

Example: A List Alternating Between Two Types of Elements

```
type ('x, 'o) alternating =
| Stop
| One of 'x * ('o, 'x) alternating

let rec to_list :
  'x 'o 'a. ('x -> 'a) -> ('o -> 'a) ->
  ('x, 'o) alternating -> 'a list =
fun x2a o2a ->
  function
  | Stop -> []
  | One (x, rest) -> x2a x :: to_list o2a x2a rest

let to_choice_list alt =
  to_list (fun x -> Left x) (fun o -> Right o) alt

let it = to_choice_list
  (One (1, One ("o", One (2, One ("oo", Stop)))))
```

Notice how the recursive call to `to_list` swaps `o2a` and `x2a` – this is necessary because the alternating structure swaps the type parameters at each level.

Example: Data-Structural Bootstrapping

```
type 'a seq =
| Nil
| Zero of ('a * 'a) seq
| One of 'a * ('a * 'a) seq
```

We store a list of elements in exponentially increasing chunks:

```

let example =
  One (0, One ((1,2), Zero (One (((3,4),(5,6)), ((7,8),(9,10))), Nil)))

let rec cons : 'a. 'a -> 'a seq -> 'a seq = (* Appending an element to the *)
  fun x -> function (* datastructure is like *)
    | Nil -> One (x, Nil) (* adding one to a binary number: 1+0=1 *)
    | Zero ps -> One (x, ps) (* 1+...0=...1 *)
    | One (y, ps) -> Zero (cons (x,y) ps) (* 1+...1=[...+1]0 *)

let rec lookup : 'a. int -> 'a seq -> 'a =
  fun i s -> match i, s with
    | _, Nil -> raise Not_found (* Rather than returning None : 'a option *)
    | 0, One (x, _) -> x (* we raise exception, for convenience. *)
    | i, One (_, ps) -> lookup (i-1) (Zero ps)
    | i, Zero ps ->
        let x, y = lookup (i / 2) ps in (* Random-Access lookup works *)
          if i mod 2 = 0 then x else y (* in logarithmic time -- much faster than *)
                                         (* in standard lists. *)

```

5.4 Algebraic Specification

The way we introduce a data structure, like complex numbers or strings, in mathematics is by specifying an *algebraic structure*.

Algebraic structures consist of a set (or several sets, for so-called *multisorted* algebras) and a bunch of functions (also known as operations) over this set (or sets).

A *signature* is a rough description of an algebraic structure: it provides *sorts* – names for the sets (in the multisorted case) – and names of the functions-operations together with their arity (and what sorts of arguments they take).

We select a class of algebraic structures by providing axioms that have to hold. We will call such classes *algebraic specifications*. In mathematics, a rusty name for some algebraic specifications is a *variety*; a more modern name is *algebraic category*.

Algebraic structures correspond to “implementations” and signatures to “interfaces” in programming languages. We will say that an algebraic structure *implements* an algebraic specification when all axioms of the specification hold in the structure. All algebraic specifications are implemented by multiple structures!

We say that an algebraic structure does not have *junk* when all its elements (i.e., elements in the sets corresponding to sorts) can be built using operations in its signature.

We allow parametric types as sorts. In that case, strictly speaking, we define a family of algebraic specifications (a different specification for each instantiation of the parametric type).

5.4.1 Algebraic Specifications: Examples An algebraic specification can also use an earlier specification. In “impure” languages like OCaml and F# we allow that the result of any operation be an error. In Haskell we could use `Maybe`.

Specification nat_p (bounded natural numbers):

nat_p

$0 : \text{nat}_p$
 $\text{succ} : \text{nat}_p \rightarrow \text{nat}_p$
 $+ : \text{nat}_p \rightarrow \text{nat}_p \rightarrow \text{nat}_p$
 $* : \text{nat}_p \rightarrow \text{nat}_p \rightarrow \text{nat}_p$
 Variables: $n, m : \text{nat}_p$
 Axioms:
 $0 + n = n, n + 0 = n$
 $m + \text{succ}(n) = \text{succ}(m + n)$
 $0 * n = 0, n * 0 = 0$
 $m * \text{succ}(n) = m + (m * n)$
 $\underbrace{\text{succ}(\dots \text{succ}(0))}_{\substack{\text{less than } p \text{ times}}} \neq 0$
 $\underbrace{\text{succ}(\dots \text{succ}(0))}_{p \text{ times}} = 0$

Specification string_p (bounded strings):

string_p

uses `char`, nat_p
 $"" : \text{string}_p$
 $"c" : \text{char} \rightarrow \text{string}_p$
 $^ : \text{string}_p \rightarrow \text{string}_p \rightarrow \text{string}_p$
 $\cdot[] : \text{string}_p \rightarrow \text{nat}_p \rightarrow \text{char}$
 Variables: $s : \text{string}_p, c, c_1, \dots, c_p : \text{char}, n : \text{nat}_p$
 Axioms:
 $"" ^ s = s, s ^ "" = s$
 $\underbrace{"c_1" ^ (\dots ^ "c_p")}_{p \text{ times}} = \text{error}$
 $r^s = (r^s)^t$
 $(c^s)[0] = c$
 $(c^s)[\text{succ}(n)] = s[n]$
 $""[n] = \text{error}$

5.5 Homomorphisms

Homomorphisms are mappings between algebraic structures with the same signature that preserve operations.

A *homomorphism* from algebraic structure $(A, \{f^A, g^A, \dots\})$ to $(B, \{f^B, g^B, \dots\})$ is a function $h : A \rightarrow B$ such that: - $h(f^A(a_1, \dots, a_{n_f})) = f^B(h(a_1), \dots, h(a_{n_f}))$ for all (a_1, \dots, a_{n_f}) - $h(g^A(a_1, \dots, a_{n_g})) = g^B(h(a_1), \dots, h(a_{n_g}))$ for all (a_1, \dots, a_{n_g}) - and so on for all operations.

Two algebraic structures are *isomorphic* if there are homomorphisms $h_1 : A \rightarrow B$, $h_2 : B \rightarrow A$ from one to the other and back, that when composed in any order form identity: $\forall(b \in B) h_1(h_2(b)) = b$ and $\forall(a \in A) h_2(h_1(a)) = a$.

An algebraic specification whose all implementations without junk are isomorphic is called “*monomorphic*”. We usually only add axioms that really matter to us to the specification, so that the implementations have room for optimization. For this reason, the resulting specifications will often not be monomorphic in the above sense.

5.6 Example: Maps

A *map* (also called dictionary or associative array) associates keys with values. Here is an algebraic specification:

(α, β) map

uses bool, type parameters α, β

empty : (α, β) map

member : $\alpha \rightarrow (\alpha, \beta)$ map \rightarrow bool

add : $\alpha \rightarrow \beta \rightarrow (\alpha, \beta)$ map $\rightarrow (\alpha, \beta)$ map

remove : $\alpha \rightarrow (\alpha, \beta)$ map $\rightarrow (\alpha, \beta)$ map

find : $\alpha \rightarrow (\alpha, \beta)$ map $\rightarrow \beta$

Variables: $k, k_2 : \alpha, v, v_2 : \beta, m : (\alpha, \beta)$ map

Axioms:

member($k, \text{add}(k, v, m)$) = true

member($k, \text{remove}(k, m)$) = false

member($k, \text{add}(k_2, v, m)$) = true $\wedge k \neq k_2 \Leftrightarrow$ member(k, m) = true $\wedge k \neq k_2$

member($k, \text{remove}(k_2, m)$) = true $\wedge k \neq k_2 \Leftrightarrow$ member(k, m) = true $\wedge k \neq k_2$

find($k, \text{add}(k, v, m)$) = v

find($k, \text{remove}(k, m)$) = error, find(k, empty) = error

find($k, \text{add}(k_2, v_2, m)$) = v $\wedge k \neq k_2 \Leftrightarrow$ find(k, m) = v $\wedge k \neq k_2$

find($k, \text{remove}(k_2, m)$) = v $\wedge k \neq k_2 \Leftrightarrow$ find(k, m) = v $\wedge k \neq k_2$

remove(k, empty) = empty

5.7 Modules and Interfaces (Signatures): Syntax

In the ML family of languages, structures are given names by **module** bindings, and signatures are types of modules. From outside of a structure or signature, we refer to the values or types it provides with a dot notation: `Module.value`.

Module (and module type) names have to start with a capital letter (in ML languages). Since modules and module types have names, there is a tradition to name the central type of a signature (the one that is “specified” by the signature), for brevity, `t`. Module types are often named with “all-caps” (all letters upper case).

```
module type MAP = sig
  type ('a, 'b) t
  val empty : ('a, 'b) t
  val member : 'a -> ('a, 'b) t -> bool
  val add : 'a -> 'b -> ('a, 'b) t -> ('a, 'b) t
  val remove : 'a -> ('a, 'b) t -> ('a, 'b) t
  val find : 'a -> ('a, 'b) t -> 'b
end

module ListMap : MAP = struct
  type ('a, 'b) t = ('a * 'b) list
  let empty = []
  let member = List.mem_assoc
  let add k v m = (k, v)::m
  let remove = List.remove_assoc
  let find = List.assoc
end
```

5.8 Implementing Maps: Association Lists

Let us now build an implementation of maps from the ground up. The most straightforward implementation... might not be what you expected:

```
module TrivialMap : MAP = struct
  type ('a, 'b) t =
    | Empty
    | Add of 'a * 'b * ('a, 'b) t
    | Remove of 'a * ('a, 'b) t

  let empty = Empty

  let rec member k m =
    match m with
    | Empty -> false
    | Add (k2, _, _) when k = k2 -> true
    | Remove (k2, _) when k = k2 -> false
```

```

| Add (_, _, m2) -> member k m2
| Remove (_, m2) -> member k m2

let add k v m = Add (k, v, m)
let remove k m = Remove (k, m)

let rec find k m =
  match m with
  | Empty -> raise Not_found
  | Add (k2, v, _) when k = k2 -> v
  | Remove (k2, _) when k = k2 -> raise Not_found
  | Add (_, _, m2) -> find k m2
  | Remove (_, m2) -> find k m2
end

```

This “trivial” implementation simply records all operations as a log. The `add` and `remove` operations are $O(1)$, but `member` and `find` must traverse the entire history.

Here is an implementation based on association lists, i.e., on lists of key-value pairs:

```

module MyListMap : MAP = struct
  type ('a, 'b) t = Empty | Add of 'a * 'b * ('a, 'b) t

  let empty = Empty

  let rec member k m =
    match m with
    | Empty -> false
    | Add (k2, _, _) when k = k2 -> true
    | Add (_, _, m2) -> member k m2

  let rec add k v m =
    match m with
    | Empty -> Add (k, v, Empty)
    | Add (k2, _, m) when k = k2 -> Add (k, v, m)
    | Add (k2, v2, m) -> Add (k2, v2, add k v m)

  let rec remove k m =
    match m with
    | Empty -> Empty
    | Add (k2, _, m) when k = k2 -> m
    | Add (k2, v, m) -> Add (k2, v, remove k m)

  let rec find k m =
    match m with

```

```

| Empty -> raise Not_found
| Add (k2, v, _) when k = k2 -> v
| Add (_, _, m2) -> find k m2
end

```

5.9 Implementing Maps: Binary Search Trees

Binary search trees are binary trees with elements stored at the interior nodes, such that elements to the left of a node are smaller than, and elements to the right bigger than, elements within a node.

For maps, we store key-value pairs as elements in binary search trees, and compare the elements by keys alone.

On average, binary search trees are fast because they use “divide-and-conquer” to search for the value associated with a key ($O(\log n)$ complexity). In the worst case, however, they reduce to association lists.

The simple polymorphic signature for maps is only possible with implementations based on some total order of keys because OCaml has polymorphic comparison (and equality) operators. These operators work on elements of most types, but not on functions. They may not work in a way you would want though! Our signature for polymorphic maps is not the standard approach because of the problem of needing the order of keys; it is just to keep things simple.

```

module BTreemap : MAP = struct
  type ('a, 'b) t = Empty | T of ('a, 'b) t * 'a * 'b * ('a, 'b) t

  let empty = Empty

  let rec member k m =
    match m with
    | Empty -> false
    | T (_, k2, _, _) when k = k2 -> true
    | T (m1, k2, _, _) when k < k2 -> member k m1
    | T (_, _, _, m2) -> member k m2

  let rec add k v m =
    match m with
    | Empty -> T (Empty, k, v, Empty)
    | T (m1, k2, _, m2) when k = k2 -> T (m1, k, v, m2)
    | T (m1, k2, v2, m2) when k < k2 -> T (add k v m1, k2, v2, m2)
    | T (m1, k2, v2, m2) -> T (m1, k2, v2, add k v m2)

  let rec split_rightmost m =
    match m with
    | Empty -> raise Not_found
    | T (Empty, k, v, Empty) -> k, v, Empty

```

(* A helper function, it does not belong *
(* to the "exported" signature. *)
(* We remove one element, *)

```

| T (m1, k, v, m2) ->                                (* the one that is on the bottom right. *)
  let rk, rv, rm = split_rightmost m2 in
  rk, rv, T (m1, k, v, rm)

let rec remove k m =
  match m with
  | Empty -> Empty
  | T (m1, k2, _, Empty) when k = k2 -> m1
  | T (Empty, k2, _, m2) when k = k2 -> m2
  | T (m1, k2, _, m2) when k = k2 ->
    let rk, rv, rm = split_rightmost m1 in
    T (rm, rk, rv, m2)
  | T (m1, k2, v, m2) when k < k2 -> T (remove k m1, k2, v, m2)
  | T (m1, k2, v, m2) -> T (m1, k2, v, remove k m2)

let rec find k m =
  match m with
  | Empty -> raise Not_found
  | T (_, k2, v, _) when k = k2 -> v
  | T (m1, k2, _, _) when k < k2 -> find k m1
  | T (_, _, _, m2) -> find k m2
end

```

5.10 Implementing Maps: Red-Black Trees

This section is based on Wikipedia's Red-black tree article, Chris Okasaki's "Purely Functional Data Structures" and Matt Might's excellent blog post on red-black tree deletion.

Binary search trees are good when we encounter keys in random order, because the cost of operations is limited by the depth of the tree which is small relative to the number of nodes... unless the tree grows unbalanced achieving large depth (which means there are sibling subtrees of vastly different sizes on some path).

To remedy this, we *rebalance* the tree while building it – i.e., while adding elements.

In *red-black trees* we achieve balance by: 1. Remembering one of two colors with each node 2. Keeping the same length of each root-to-leaf path if only black nodes are counted 3. Not allowing a red node to have a red child

This way the depth is at most twice the depth of a perfectly balanced tree with the same number of nodes.

5.10.1 B-trees of Order 4 (2-3-4 Trees) How can we have perfectly balanced trees without worrying about having $2^k - 1$ elements? **2-3-4 trees** can store

from 1 to 3 elements in each node and have 2 to 4 subtrees correspondingly. Lots of freedom!

A 2-node contains one element and has two children. A 3-node contains two elements and has three children. A 4-node contains three elements and has four children.

To insert “25” into a 2-3-4 tree, we descend right, but if we encounter a full node (4-node), we move the middle element up and split the remaining elements. This maintains balance at all times.

To represent a 2-3-4 tree as a binary tree with one element per node, we color the middle element of a 4-node, or the first element of a 2-/3-node, black and make it the parent of its neighbor elements, and make them parents of the original subtrees. This correspondence provides the intuition behind red-black trees.

5.10.2 Red-Black Trees, Without Deletion Red-black trees maintain two invariants:

Invariant 1. No red node has a red child.

Invariant 2. Every path from the root to an empty node contains the same number of black nodes.

First we implement red-black tree based sets without deletion. The implementation proceeds almost exactly like for unbalanced binary search trees; we only need to restore invariants.

By keeping balance at each step of constructing a node, it is enough to check locally (around the root of the subtree). For an understandable implementation of deletion, we need to introduce more colors – see Matt Might’s post.

```

type color = R | B
type 'a t = E | T of color * 'a t * 'a * 'a t

let empty = E

let rec member x m =                                     (* Like in unbalanced binary search tree. *)
  match m with
  | E -> false
  | T (_ , _ , y , _) when x = y -> true
  | T (_ , a , y , _) when x < y -> member x a
  | T (_ , _ , b) -> member x b

let balance = function                                (* Restoring the invariants. *)
  | B , T (R , T (R,a,x,b) , y , c) , z , d      (* On next figure: left, *)
  | B , T (R , a , x , T (R,b,y,c)) , z , d       (* top, *)
  | B , a , x , T (R , T (R,b,y,c) , z , d)        (* bottom, *)
  | B , a , x , T (R , b , y , T (R,c,z,d))        (* right, *)

```

```

    -> T (R, T (B,a,x,b), y, T (B,c,z,d))      (* center tree. *)
| color, a, x, b -> T (color, a, x, b)          (* We allow red-red violation for now. *)

let insert x s =
  let rec ins = function
    | E -> T (R, E, x, E)                         (* Like in unbalanced binary search tree, *)
    | T (color, a, y, b) as s ->
        if x < y then balance (color, ins a, y, b)
        else if x > y then balance (color, a, y, ins b)
        else s
    in
    match ins s with
    | T (_, a, y, b) -> T (B, a, y, b)             (* We could still have red-red violation at *)
    | E -> failwith "insert: impossible"            (* fixed by coloring it black. *)

```

The `balance` function handles four cases where a red-red violation occurs (a red node with a red child). In each case, we restructure the tree to eliminate the violation while maintaining the binary search tree property. All four cases produce the same balanced result: a red root with two black children.

Exercises

Exercise 1. Derive the equations and solve them to find the type for:

```
let cadr l = List.hd (List.tl l) in cadr (1::2::[]), cadr (true::false::[])
```

in environment $\Gamma = \{\text{List.hd} : \forall \alpha. \alpha \text{ list} \rightarrow \alpha; \text{List.tl} : \forall \alpha. \alpha \text{ list} \rightarrow \alpha \text{ list}\}$. You can take “shortcuts” if it is too many equations to write down.

Exercise 2. Terms $t_1, t_2, \dots \in T(\Sigma, X)$ are built out of variables $x, y, \dots \in X$ and function symbols $f, g, \dots \in \Sigma$ the way you build values out of functions:

- $X \subset T(\Sigma, X)$ – variables are terms; usually an infinite set,
- for terms $t_1, \dots, t_n \in T(\Sigma, X)$ and a function symbol $f \in \Sigma_n$ of arity n , $f(t_1, \dots, t_n) \in T(\Sigma, X)$ – bigger terms arise from applying function symbols to smaller terms; $\Sigma = \bigcup_n \Sigma_n$ is called a signature.

In OCaml, we can define terms as: `type term = V of string | T of string * term list`, where for example `V("x")` is a variable x and `T("f", [V("x"); V("y")])` is the term $f(x, y)$.

By *substitutions* σ, ρ, \dots we mean finite sets of variable-term pairs which we can write as $\{x_1 \mapsto t_1, \dots, x_k \mapsto t_k\}$ or $[x_1 := t_1; \dots; x_k := t_k]$, but also functions from terms to terms $\sigma : T(\Sigma, X) \rightarrow T(\Sigma, X)$ related to the pairs as follows: if $\sigma = \{x_1 \mapsto t_1, \dots, x_k \mapsto t_k\}$, then

- $\sigma(x_i) = t_i$ for $x_i \in \{x_1, \dots, x_k\}$,
- $\sigma(x) = x$ for $x \in X \setminus \{x_1, \dots, x_k\}$,
- $\sigma(f(t_1, \dots, t_n)) = f(\sigma(t_1), \dots, \sigma(t_n))$.

In OCaml, we can define substitutions σ as: `type subst = (string * term) list`, together with a function `apply : subst -> term -> term` which computes $\sigma(\cdot)$.

We say that a substitution σ is *more general* than all substitutions $\rho \circ \sigma$, where $(\rho \circ \sigma)(x) = \rho(\sigma(x))$. In type inference, we are interested in most general solutions.

A *unification problem* is a finite set of equations $S = \{s_1 =? t_1, \dots, s_n =? t_n\}$. A solution, or *unifier* of S , is a substitution σ such that $\sigma(s_i) = \sigma(t_i)$ for $i = 1, \dots, n$. A *most general unifier*, or *MGU*, is a most general such substitution.

1. Implement an algorithm that, given a set of equations represented as a list of pairs of terms, computes an idempotent most general unifier of the equations.
2. (Ex. 4.22 in Franz Baader and Tobias Nipkow “Term Rewriting and All That”, p. 82.) Modify the implementation of unification to achieve linear space complexity by working with what could be called iterated substitutions.

Exercise 3.

1. What does it mean that an implementation has junk (as an algebraic structure for a given signature)? Is it bad?
2. Define a monomorphic algebraic specification (other than, but similar to, nat_p or string_p , some useful data type).
3. Discuss an example of a (monomorphic) algebraic specification where it would be useful to drop some axioms (giving up monomorphicity) to allow more efficient implementations.

Exercise 4.

1. Does the example `ListMap` meet the requirements of the algebraic specification for maps? Hint: here is the definition of `List.remove_assoc`; `compare a x` equals 0 if and only if `a = x`.

```
let rec remove_assoc x = function
| [] -> []
| (a, b as pair) :: l ->
  if compare a x = 0 then l else pair :: remove_assoc x l
```

2. Trick question: what is the computational complexity of `ListMap` or `TrivialMap`?
3. (*) The implementation `MyListMap` is inefficient: it performs a lot of copying and is not tail-recursive. Optimize it (without changing the type definition).
4. Add (and specify) `isEmpty : (α, β) map → bool` to the example algebraic specification of maps without increasing the burden on its implementations.

Hint: equational reasoning might be not enough; consider an equivalence relation \approx meaning “have the same keys”.

Exercise 5. Design an algebraic specification and write a signature for first-in-first-out queues. Provide two implementations: one straightforward using a list, and another one using two lists: one for freshly added elements providing efficient queueing of new elements, and “reversed” one for efficient popping of old elements.

Exercise 6. Design an algebraic specification and write a signature for sets. Provide two implementations: one straightforward using a list, and another one using a map into the unit type.

Exercise 7.

1. (Ex. 2.2 in Chris Okasaki “Purely Functional Data Structures”) In the worst case, `member` performs approximately $2d$ comparisons, where d is the depth of the tree. Rewrite `member` to take no more than $d + 1$ comparisons by keeping track of a candidate element that *might* be equal to the query element (say, the last element for which `<` returned false) and checking for equality only when you hit the bottom of the tree.
2. (Ex. 3.10 in Chris Okasaki “Purely Functional Data Structures”) The `balance` function currently performs several unnecessary tests: when e.g. `ins` recurses on the left child, there are no violations on the right child.
 - Split `balance` into `lbalance` and `rbalance` that test for violations of left resp. right child only. Replace calls to `balance` appropriately.
 - One of the remaining tests on grandchildren is also unnecessary. Rewrite `ins` so that it never tests the color of nodes not on the search path.

Exercise 8. (*) Implement maps (i.e. write a module for the map signature) based on AVL trees. See http://en.wikipedia.org/wiki/AVL_tree.

Chapter 6: Folding and Backtracking

This chapter explores two fundamental programming paradigms in functional programming: **folding** (also known as reduction) and **backtracking**. We begin with the classic `map` and `fold` higher-order functions, examine how they generalize to trees and other data structures, then move on to solving puzzles using backtracking with lists.

6.1 Basic Generic List Operations

Functional programming emphasizes identifying common patterns and abstracting them into reusable higher-order functions. Let us see how this works in practice.

The map Function Consider the problem of printing a comma-separated list of integers. The `String` module provides:

```
val concat : string -> string list -> string
```

First, we need to convert numbers into strings:

```
let rec strings_of_ints = function
| [] -> []
| hd::tl -> string_of_int hd :: strings_of_ints tl

let comma_sep_ints = String.concat ", " -| strings_of_ints
```

Similarly, to sort strings from shortest to longest, we first compute lengths:

```
let rec strings_lengths = function
| [] -> []
| hd::tl -> (String.length hd, hd) :: strings_lengths tl

let by_size = List.sort compare -| strings_lengths
```

Notice the common structure in `strings_of_ints` and `strings_lengths`: both transform each element of a list independently. We can extract this pattern into a generic function called `map`:

```
let rec list_map f = function
| [] -> []
| hd::tl -> f hd :: list_map f tl
```

Now we can rewrite our functions more concisely:

```
let comma_sep_ints =
  String.concat ", " -| list_map string_of_int

let by_size =
  List.sort compare -| list_map (fun s -> String.length s, s)
```

The fold Function Consider summing elements of a list:

```
let rec balance = function
| [] -> 0
| hd::tl -> hd + balance tl
```

Or multiplying elements:

```
let rec total_ratio = function
| [] -> 1.
| hd::tl -> hd *. total_ratio tl
```

The pattern is the same: we combine each element with the result of processing the rest of the list. This is the `fold` operation:

```
let rec list_fold f base = function
| [] -> base
| hd::tl -> f hd (list_fold f base tl)
```

Important: Note that `list_fold f base l` equals `List.fold_right f l base`. The OCaml standard library uses a different argument order.

The key insight is that `map` alters the *contents* of data without changing its structure, while `fold` computes a value using the structure as scaffolding. Visually:

- `map` transforms: `[a; b; c; d]` becomes `[f a; f b; f c; f d]`
- `fold` collapses: `[a; b; c; d]` becomes `f a (f b (f c (f d accu)))`

6.2 Making Fold Tail-Recursive

Let us investigate tail-recursive functions. Consider reversing a list:

```
let rec list_rev acc = function
| [] -> acc
| hd::tl -> list_rev (hd::acc) tl
```

Or computing an average:

```
let rec average (sum, tot) = function
| [] when tot = 0. -> 0.
| [] -> sum /. tot
| hd::tl -> average (hd +. sum, 1. +. tot) tl
```

The pattern here is different from `fold_right`. We process elements from left to right, accumulating a result:

```
let rec fold_left f accu = function
| [] -> accu
| a::l -> fold_left f (f accu a) l
```

With `fold_left`, hiding the accumulator is straightforward:

```
let list_rev l =
  fold_left (fun t h -> h::t) [] l

let average =
  fold_left (fun (sum, tot) e -> sum +. e, 1. +. tot) (0., 0.)
```

The naming convention for `fold_right` and `fold_left` reflects associativity:

- `fold_right f` makes `f` **right associative**, like the list constructor `:::`
`List.fold_right f [a1; ...; an] b` is `f a1 (f a2 (... (f an b) ...))`
- `fold_left f` makes `f` **left associative**, like function application:
`List.fold_left f a [b1; ...; bn]` is `f (... (f (f a b1) b2) ...) bn`

The “backward” structure of `fold_left`: - Input: [a; b; c; d] - Result: f (f (f (f accu a) b) c) d

Useful Derived Functions List filtering is naturally expressed using `fold_right`:

```
let list_filter p l =
  List.fold_right (fun h t -> if p h then h :: t else t) l []
```

A tail-recursive map that returns elements in reverse order:

```
let list_rev_map f l =
  List.fold_left (fun t h -> f h :: t) [] l
```

6.3 Map and Fold for Trees and Other Structures

Binary Trees Mapping binary trees is straightforward:

```
type 'a btree = Empty | Node of 'a * 'a btree * 'a btree

let rec bt_map f = function
  | Empty -> Empty
  | Node (e, l, r) -> Node (f e, bt_map f l, bt_map f r)

let test = Node
  (3, Node (5, Empty, Empty), Node (7, Empty, Empty))
let _ = bt_map ((+) 1) test
```

The `map` and `fold` functions we consider here preserve/respect the structure of data. They do **not** correspond to `map` and `fold` of abstract data type containers (which are like `List.rev_map` and `List.fold_left` over container elements in arbitrary order). Here we generalize `List.map` and `List.fold_right` to other structures.

The most general form of `fold` for binary trees processes each element together with partial results from subtrees:

```
let rec bt_fold f base = function
  | Empty -> base
  | Node (e, l, r) ->
    f e (bt_fold f base l) (bt_fold f base r)
```

Examples:

```
let sum_els = bt_fold (fun i l r -> i + l + r) 0
let depth t = bt_fold (fun _ l r -> 1 + max l r) 1 t
```

More Complex Structures: Expressions To demonstrate map and fold for more complex structures, we recall the expression type from Chapter 3:

```

type expression =
  Const of float
  | Var of string
  | Sum of expression * expression (* e1 + e2 *)
  | Diff of expression * expression (* e1 - e2 *)
  | Prod of expression * expression (* e1 * e2 *)
  | Quot of expression * expression (* e1 / e2 *)

```

The multitude of cases makes the datatype harder to work with. Fortunately, *or-patterns* help:

```

let rec vars = function
  | Const _ -> []
  | Var x -> [x]
  | Sum (a,b) | Diff (a,b) | Prod (a,b) | Quot (a,b) ->
    vars a @ vars b

```

Mapping and folding must be specialized for each case. We pack behaviors into records:

```

type expression_map = {
  map_const : float -> expression;
  map_var : string -> expression;
  map_sum : expression -> expression -> expression;
  map_diff : expression -> expression -> expression;
  map_prod : expression -> expression -> expression;
  map_quot : expression -> expression -> expression;
}

(* Note: 'a replaces expression because fold produces values of arbitrary type *)
type 'a expression_fold = {
  fold_const : float -> 'a;
  fold_var : string -> 'a;
  fold_sum : 'a -> 'a -> 'a;
  fold_diff : 'a -> 'a -> 'a;
  fold_prod : 'a -> 'a -> 'a;
  fold_quot : 'a -> 'a -> 'a;
}

```

We define standard behaviors that can be tailored for specific uses:

```

let identity_map = {
  map_const = (fun c -> Const c);
  map_var = (fun x -> Var x);
  map_sum = (fun a b -> Sum (a, b));
  map_diff = (fun a b -> Diff (a, b));
  map_prod = (fun a b -> Prod (a, b));
  map_quot = (fun a b -> Quot (a, b));
}

```

```

let make_fold op base = {
  fold_const = (fun _ -> base);
  fold_var = (fun _ -> base);
  fold_sum = op; fold_diff = op;
  fold_prod = op; fold_quot = op;
}

```

The actual `map` and `fold` functions:

```

let rec expr_map emap = function
| Const c -> emap.map_const c
| Var x -> emap.map_var x
| Sum (a,b) -> emap.map_sum (expr_map emap a) (expr_map emap b)
| Diff (a,b) -> emap.map_diff (expr_map emap a) (expr_map emap b)
| Prod (a,b) -> emap.map_prod (expr_map emap a) (expr_map emap b)
| Quot (a,b) -> emap.map_quot (expr_map emap a) (expr_map emap b)

let rec expr_fold efold = function
| Const c -> efold.fold_const c
| Var x -> efold.fold_var x
| Sum (a,b) -> efold.fold_sum (expr_fold efold a) (expr_fold efold b)
| Diff (a,b) -> efold.fold_diff (expr_fold efold a) (expr_fold efold b)
| Prod (a,b) -> efold.fold_prod (expr_fold efold a) (expr_fold efold b)
| Quot (a,b) -> efold.fold_quot (expr_fold efold a) (expr_fold efold b)

```

Using the `{record with field = value}` syntax to customize behaviors:

```

let prime_vars = expr_map
  {identity_map with map_var = fun x -> Var (x ^ "''")}

let subst s =
  let apply x = try List.assoc x s with Not_found -> Var x in
  expr_map {identity_map with map_var = apply}

let vars =
  expr_fold {(make_fold (@) []) with fold_var = fun x -> [x]}

let size = expr_fold (make_fold (fun a b -> 1 + a + b) 1)

let eval env = expr_fold {
  fold_const = id;
  fold_var = (fun x -> List.assoc x env);
  fold_sum = (+.); fold_diff = (-.);
  fold_prod = (*.); fold_quot = (/.);
}

```

6.4 Point-Free Programming

In 1977/78, John Backus designed **FP**, the first *function-level programming* language. Over the next decade it evolved into the **FL** language.

“Clarity is achieved when programs are written at the function level—that is, by putting together existing programs to form new ones, rather than by manipulating objects and then abstracting from those objects to produce programs.” – *The FL Project: The Design of a Functional Language*

For function-level programming, we need combinators like these from *OCaml Batteries*:

```
let const x _ = x
let ( |- ) f g x = g (f x)           (* forward composition *)
let ( -| ) f g x = f (g x)           (* backward composition *)
let flip f x y = f y x
let ( *** ) f g = fun (x,y) -> (f x, g y)
let ( &&& ) f g = fun x -> (f x, g x)
let first f x = fst (f x)
let second f x = snd (f x)
let curry f x y = f (x,y)
let uncurry f (x,y) = f x y
```

The flow of computation can be viewed as a circuit where results of nodes (functions) connect to further nodes as inputs. We represent cross-sections of the circuit as tuples of intermediate values.

```
let print2 c i =
  let a = Char.escaped c in
  let b = string_of_int i in
  a ^ b
```

In point-free style:

```
let print2 = curry
  ((Char.escaped *** string_of_int) |- uncurry (^))
```

Since we usually pass arguments one at a time rather than in tuples, we need `uncurry` to access multi-argument functions. Converting a C/Pascal-like function to one that takes arguments one at a time is called *currying*, after logician Haskell Brooks Curry.

Another approach uses function composition, `flip`, and the **S** combinator:

```
let s x y z = x z (y z)
```

Example: transforming a filter-map function step by step:

```
let func2 f g l = List.filter f (List.map g l)
(* Using composition: *)
```

```

let func2 f g = (-|) (List.filter f) (List.map g)
let func2 f = (-|) (List.filter f) -| List.map
(* Eliminating f: *)
let func2 f = (-|) ((-|) (List.filter f)) List.map
let func2 f = flip (-|) List.map ((-|) (List.filter f))
let func2 f = (((|-) List.map) -| ((-|) -| List.filter)) f
let func2 = (-|) List.map -| ((-|) -| List.filter)

```

6.5 Reductions and More Higher-Order Functions

Mathematics has notation for sums over intervals: $\sum_{n=a}^b f(n)$.

In OCaml, we do not have a universal addition operator:

```

let rec i_sum_fromto f a b =
  if a > b then 0
  else f a + i_sum_fromto f (a+1) b

let rec f_sum_fromto f a b =
  if a > b then 0.
  else f a +. f_sum_fromto f (a+1) b

let pi2_over6 =
  f_sum_fromto (fun i -> 1. /. float_of_int (i*i)) 1 5000

```

The natural generalization:

```

let rec op_fromto op base f a b =
  if a > b then base
  else op (f a) (op_fromto op base f (a+1) b)

```

Collecting Results: concat_map Let us collect results of a multifunction (set-valued function) for a set of arguments. In mathematical notation:

$$f(A) = \bigcup_{p \in A} f(p)$$

This translates to a useful list operation with union as append:

```

let rec concat_map f = function
  | [] -> []
  | a::l -> f a @ concat_map f l

```

More efficiently (tail-recursive):

```

let concat_map f l =
  let rec cmap_f accu = function
    | [] -> accu

```

```

| a::l -> cmap_f (List.rev_append (f a) accu) l in
List.rev (cmap_f [] 1)

```

All Subsequences of a List

```

let rec subseqs l =
  match l with
  | [] -> []
  | x::xs ->
    let pxs = subseqs xs in
    List.map (fun px -> x::px) pxs @ pxs

```

Tail-recursively:

```

let rec rmap_append f accu = function
  | [] -> accu
  | a::l -> rmap_append f (f a :: accu) l

let rec subseqs l =
  match l with
  | [] -> []
  | x::xs ->
    let pxs = subseqs xs in
    rmap_append (fun px -> x::px) pxs pxs

```

6.6 Grouping and Map-Reduce

It is often useful to organize values by some property.

Collecting by Key First, we collect elements from an association list by key:

```

let collect l =
  match List.sort (fun x y -> compare (fst x) (fst y)) l with
  | [] -> []
(* Start with associations sorted by key *)
  | (k0, v0)::tl ->
    let k0, vs, l = List.fold_left
      (fun (k0, vs, l) (kn, vn) ->
        (* Collect values for current key *)
        if k0 = kn then k0, vn::vs, l
        (* and when the key changes, *)
        else kn, [vn], (k0, List.rev vs)::l) (* stack the collected values *)
      (k0, [v0], []) tl in
      (* Why reverse? To preserve order *)
      List.rev ((k0, List.rev vs)::l)

```

Now we can group by an arbitrary property:

```
let group_by p l = collect (List.map (fun e -> p e, e) l)
```

Reduction (Aggregation) To process results like SQL aggregate operations, we add **reduction**:

```
let aggregate_by p red base l =
  let ags = group_by p l in
  List.map (fun (k, vs) -> k, List.fold_right red vs base) ags
```

Using the **feed-forward** (pipe) operator `let (|>) x f = f x:`

```
let aggregate_by p redf base l =
  group_by p l
  |> List.map (fun (k, vs) -> k, List.fold_right redf vs base)
```

Often it is easier to extract the property upfront. Since we first map elements into key-value pairs, we call this `map_reduce`:

```
let map_reduce mapf redf base l =
  List.map mapf l
  |> collect
  |> List.map (fun (k, vs) -> k, List.fold_right redf vs base)
```

Map-Reduce Examples Sometimes we have multiple sources of information:

```
let concat_reduce mapf redf base l =
  concat_map mapf l
  |> collect
  |> List.map (fun (k, vs) -> k, List.fold_right redf vs base)
```

Computing a merged histogram of documents:

```
let histogram documents =
  let mapf doc =
    Str.split (Str.regexp "[ \t.,;]+") doc
    |> List.map (fun word -> word, 1) in
  concat_reduce mapf (+) 0 documents
```

Computing an inverted index:

```
let cons hd tl = hd::tl

let inverted_index documents =
  let mapf (addr, doc) =
    Str.split (Str.regexp "[ \t.,;]+") doc
    |> List.map (fun word -> word, addr) in
  concat_reduce mapf cons [] documents
```

A simple “search engine”:

```
let search index words =
  match List.map (flip List.assoc index) words with
  | [] -> []
  | idx::idcs -> List.fold_left intersect idx idcs
```

where `intersect` computes intersection of sets represented as sorted lists:

```

let intersect xs ys = (* Sets as sorted lists *)
  let rec aux acc = function
    | [], _ | _, [] -> acc
    | (x::xs' as xs), (y::ys' as ys) ->
      let c = compare x y in
      if c = 0 then aux (x::acc) (xs', ys')
      else if c < 0 then aux acc (xs', ys)
      else aux acc (xs, ys') in
  List.rev (aux [] (xs, ys))

```

6.7 Higher-Order Functions for the Option Type

Operating on optional values:

```

let map_option f = function
| None -> None
| Some e -> f e

```

Mapping over a list and filtering out failures:

```

let rec map_some f = function
| [] -> []
| e::l -> match f e with
| None -> map_some f l
| Some r -> r :: map_some f l

```

Tail-recursively:

```

let map_some f l =
  let rec maps_f accu = function
    | [] -> accu
    | a::l -> maps_f (match f a with None -> accu
    | Some r -> r :: accu) l in
  List.rev (maps_f [] l)

```

6.8 The Countdown Problem Puzzle

The Countdown Problem is a classic puzzle:

- Using a given set of numbers and arithmetic operators +, -, *, /, construct an expression with a given value.
- All numbers, including intermediate results, must be positive integers.
- Each source number can be used at most once.

Example: - Numbers: 1, 3, 7, 10, 25, 50 - Target: 765 - Possible solution: $(25-10) * (50+1)$

There are 780 solutions for this example. Changing the target to 831 gives an example with no solutions.

Data Types

```
type op = Add | Sub | Mul | Div

let apply op x y =
  match op with
  | Add -> x + y
  | Sub -> x - y
  | Mul -> x * y
  | Div -> x / y

let valid op x y =
  match op with
  | Add -> true
  | Sub -> x > y
  | Mul -> true
  | Div -> x mod y = 0

type expr = Val of int | App of op * expr * expr

let rec eval = function
| Val n -> if n > 0 then Some n else None
| App (o, l, r) ->
  eval l |> map_option (fun x ->
    eval r |> map_option (fun y ->
      if valid o x y then Some (apply o x y)
      else None))

let rec values = function
| Val n -> [n]
| App (_, l, r) -> values l @ values r

let solution e ns n =
  list_diff (values e) ns = [] && is_unique (values e) &&
  eval e = Some n
```

Brute Force Solution Splitting a list into two non-empty parts:

```
let split l =
  let rec aux lhs acc = function
  | [] | [_] -> []
  | [y; z] -> (List.rev (y::lhs), [z])::acc
  | hd::rhs ->
    let lhs = hd::lhs in
    aux lhs ((List.rev lhs, rhs)::acc) rhs in
  aux [] [] l
```

We introduce an operator for working with multiple data sources:

```
let ( |-> ) x f = concat_map f x
```

Generating all expressions from a list of numbers:

```
let combine l r =
  List.map (fun o -> App (o, l, r)) [Add; Sub; Mul; Div] (* Combine two expressions using each operator *)

let rec exprs = function
| [] -> []
| [n] -> [Val n]
| ns ->
  split ns |-> (fun (ls, rs) ->
    exprs ls |-> (fun l ->
      exprs rs |-> (fun r ->
        combine l r))) (* For each split ls,rs of numbers *)
(* for each expression l over ls *)
(* and expression r over rs *)
(* produce all l ? r expressions *)
```

Finding solutions:

```
let guard n =
  List.filter (fun e -> eval e = Some n)

let solutions ns n =
  choices ns |-> (fun ns' ->
    exprs ns' |> guard n)
```

Optimization: Fuse Generation with Testing We memorize values with expressions as pairs (`e`, `eval e`), so only valid subexpressions are generated:

```
let combine' (l, x) (r, y) =
  [Add; Sub; Mul; Div]
  |> List.filter (fun o -> valid o x y)
  |> List.map (fun o -> App (o, l, r), apply o x y)

let rec results = function
| [] -> []
| [n] -> if n > 0 then [Val n, n] else []
| ns ->
  split ns |-> (fun (ls, rs) ->
    results ls |-> (fun lx ->
      results rs |-> (fun ry ->
        combine' lx ry)))

let solutions' ns n =
  choices ns |-> (fun ns' ->
    results ns'
    |> List.filter (fun (e, m) -> m = n))
```

```
|> List.map fst) (* Discard memorized values *)
```

Eliminating Symmetric Cases Strengthening the validity predicate to account for commutativity and identity:

```
let valid op x y =
  match op with
  | Add -> x <= y
  | Sub -> x > y
  | Mul -> x <= y && x <> 1 && y <> 1
  | Div -> x mod y = 0 && y <> 1
```

This eliminates repeating symmetrical solutions on the semantic level (values) rather than syntactic level (expressions)—both easier and more effective.

6.9 The Honey Islands Puzzle

The Honey Islands puzzle: Find cells to eat honey from so that the least amount of honey becomes sour (assuming sourness spreads through contact).

Given a honeycomb with some cells initially marked black, mark additional cells so that unmarked cells form `num_islands` disconnected components, each with `island_size` cells.

Representing the Honeycomb

```
type cell = int * int (* Cartesian coordinates *)
module CellSet = Set.Make (struct type t = cell let compare = compare end)
(* Store cells in sets *)

type task = {
  board_size : int; (* For board size N, coordinates *)
  num_islands : int; (* range from (-2N, -N) to (2N, N) *)
  island_size : int; (* Required number of islands *)
  empty_cells : CellSet.t; (* Required cells per island *)
  (* Initially empty cells *)
}

let cellset_of_list l = (* List to set, inverse of CellSet.elements *)
  List.fold_right CellSet.add l CellSet.empty
```

Neighborhood: Each cell (x, y) has up to 6 neighbors:

```
let neighbors n eaten (x, y) =
  List.filter
    (inside_board n eaten)
    [x-1,y-1; x+1,y-1; x+2,y;
     x+1,y+1; x-1,y+1; x-2,y]
```

Building the honeycomb:

```
let even x = x mod 2 = 0

let inside_board n eaten (x, y) =
  even x = even y && abs y <= n &&
  abs x + abs y <= 2*n &&
  not (CellSet.mem (x, y) eaten)

let honey_cells n eaten =
  fromto (-2*n) (2*n) |-> (fun x ->
    fromto (-n) n |-> (fun y ->
      guard (inside_board n eaten)
      (x, y)))
```

Testing Correctness We walk through each island counting cells depth-first:

```
let check_correct n island_size num_islands empty_cells =
  let honey = honey_cells n empty_cells in

  let rec check_board been_islands unvisited visited =
    match unvisited with
    | [] -> been_islands = num_islands
    | cell::remaining when CellSet.mem cell visited ->
        check_board been_islands remaining visited (* Keep looking *)
    | cell::remaining (* when not visited *) ->
        let (been_size, unvisited, visited) =
          check_island cell (* Visit another island *)
          (1, remaining, CellSet.add cell visited) in
        been_size = island_size
        && check_board (been_islands+1) unvisited visited

  and check_island current state =
    neighbors n empty_cells current
    |> List.fold_left (* Walk into each direction *)
        (fun (been_size, unvisited, visited as state)
         neighbor ->
          if CellSet.mem neighbor visited then state
          else
            let unvisited = remove neighbor unvisited in
            let visited = CellSet.add neighbor visited in
            let been_size = been_size + 1 in
            check_island neighbor
              (been_size, unvisited, visited))
        state in (* Initial been_size is 1 *)
```

```
check_board 0 honey empty_cells
```

Multiple Results per Step: concat_fold When processing lists with potentially multiple results per step, we need `concat_fold`:

```
let rec concat_fold f a = function
| [] -> [a]
| x::xs ->
  f x a |-> (fun a' -> concat_fold f a' xs)
```

Generating Solutions We transform the testing code into generation code by:

- Passing around the current solution `eaten`
- Returning results in a list (empty list = no solutions)
- At each neighbor, trying both eating and keeping

```
let find_to_eat n island_size num_islands empty_cells =
  let honey = honey_cells n empty_cells in

  let rec find_board been_islands unvisited visited eaten =
    match unvisited with
    | [] ->
      if been_islands = num_islands then [eaten] else []
    | cell::remaining when CellSet.mem cell visited ->
      find_board been_islands remaining visited eaten
    | cell::remaining (* when not visited *) ->
      find_island cell
      (1, remaining, CellSet.add cell visited, eaten)
    |-> (* Concatenate solutions *)
      (fun (been_size, unvisited, visited, eaten) ->
        if been_size = island_size
        then find_board (been_islands+1)
              unvisited visited eaten
        else []))

  and find_island current state =
    neighbors n empty_cells current
    |> concat_fold (* Multiple results *)
      (fun neighbor
        (been_size, unvisited, visited, eaten as state) ->
        if CellSet.mem neighbor visited then [state]
        else
          let unvisited = remove neighbor unvisited in
          let visited = CellSet.add neighbor visited in
          (been_size, unvisited, visited,
           neighbor::eaten):::
          (* solutions where neighbor is honey *)
          find_island neighbor
```

```

        (been_size+1, unvisited, visited, eaten))
state in

find_board 0 honey empty_cells []

```

Optimizations The main rule: **fail (drop solution candidates) as early as possible.**

We guard both choices (eating and keeping) and track how much honey needs to be eaten:

```

type state = {
  been_size: int;                                     (* Honey cells in current island *)
  been_islands: int;                                  (* Islands visited so far *)
  unvisited: cell list;                               (* Cells to visit *)
  visited: CellSet.t;                                (* Already visited *)
  eaten: cell list;                                   (* Current solution candidate *)
  more_to_eat: int;                                   (* Remaining cells to eat *)
}

let rec visit_cell s =
  match s.unvisited with
  | [] -> None
  | c::remaining when CellSet.mem c s.visited ->
    visit_cell {s with unvisited=remaining}
  | c::remaining (* when c not visited *) ->
    Some (c, {s with
      unvisited=remaining;
      visited = CellSet.add c s.visited})

let eat_cell c s =
  {s with eaten = c::s.eaten;
        visited = CellSet.add c s.visited;
        more_to_eat = s.more_to_eat - 1}

let keep_cell c s =                                     (* c is actually unused *)
  {s with been_size = s.been_size + 1;
        visited = CellSet.add c s.visited}

let fresh_island s =                                 (* Increment been_size at start of find_island *)
  {s with been_size = 0;
        been_islands = s.been_islands + 1}

let init_state unvisited more_to_eat = {
  been_size = 0; been_islands = 0;
  unvisited; visited = CellSet.empty;
}

```

```

    eaten = [] ; more_to_eat;
}

```

The optimized island loop only tries actions that make sense:

```

and find_island current s =
  let s = keep_cell current s in
  neighbors n empty_cells current
  |> concat_fold
  (fun neighbor s ->
    if CellSet.mem neighbor s.visited then [s]
    else
      let choose_eat = (* Guard against failed actions *)
        if s.more_to_eat = 0 then []
        else [eat_cell neighbor s]
      and choose_keep =
        if s.been_size >= island_size then []
        else find_island neighbor s in
      choose_eat @ choose_keep)
  s in
(* Finally, compute the required eaten cells and start searching *)
let cells_to_eat =
  List.length honey - island_size * num_islands in
find_board (init_state honey cells_to_eat)

```

6.10 Constraint-Based Puzzles

Puzzles can be presented by providing: 1. The general form of solutions 2. Additional requirements (constraints) that solutions must meet

For many puzzles, solutions decompose into a fixed number of **variables**: - A **domain** is the set of possible values a variable can have - In Honey Islands, variables are cells with domain {Honey, Empty} - **Constraints** specify relationships: cells that must be empty, number and size of connected components, neighborhood graph

Finite Domain Constraint Programming A general and often efficient scheme:

1. With each variable, associate a set of values (initially the full domain). The singleton containing this association is the initial set of partial solutions.
2. While there is a solution with more than one value for some variable:
 - (a) If some value for a variable fails for all possible assignments to other variables, remove it
 - (b) If a variable has an empty set of possible values, remove that solution

- (c) Select the variable with the smallest non-singleton set. Split into similarly-sized parts. Replace the solution with two solutions for each part.
3. Build final solutions by assigning each variable its single remaining value.

Simplifications: In step (2c), instead of equal-sized splits, we can partition into singleton and remainder, or partition completely into singletons.

6.11 Exercises

1. Recall how we generated all subsequences of a list. Find (generate) all:
 - permutations of a list
 - ways of choosing without repetition from a list
 - combinations of K distinct objects chosen from N elements of a list
2. Using folding for the `expression` data type, compute the degree of the corresponding polynomial.
3. Implement simplification of expressions using mapping for the `expression` data type.
4. Express in terms of `fold_left` or `fold_right`:
 - `indexed : 'a list -> (int * 'a) list`, which pairs elements with their indices
 - `concat_fold` as used in Honey Islands
 - Run-length encoding of a list: `encode ['a;'a;'a;'a;'b;'c;'c;'a;'a;'d] = [4,'a; 1,'b; 2,'c; 2,'a; 1,'d]`
5. Write more efficient variants:
 - `list_diff` computing difference of sets represented as sorted lists
 - `is_unique` in constant stack space
6. Write functions `compose` and `perform` that take a list of functions and return their composition:
 - `compose [f1; ...; fn] = x -> f1 (... (fn x)...)`
 - `perform [f1; ...; fn] = x -> fn (... (f1 x)...)`
7. Write a solver for the *Tents Puzzle*.
8. **Robot Squad** (harder): Given a map with walls and lidar readings (8 directions: E, NE, N, NW, W, SW, S, SE) for multiple robots, determine possible robot positions.
9. Write a solver for the *Plinx Puzzle* (does not need to solve all levels, but should handle initial ones).

Chapter 7: Laziness

This chapter explores lazy evaluation and stream processing in OCaml. We examine different evaluation strategies, implement streams and lazy lists, apply them to power series computation and differential equations, build circular data structures, and develop a sophisticated pipe-based pretty-printer.

7.1 Evaluation Strategies and Parameter Passing

Evaluation strategy is the order in which expressions are computed – primarily, when arguments are computed. Recall our problems with using *flow control* expressions like `if_then_else` in examples from the lambda-calculus lecture. There are many technical terms describing various evaluation strategies:

Strict evaluation: Arguments are always evaluated completely before the function is applied.

Non-strict evaluation: Arguments are not evaluated unless they are actually used in the evaluation of the function body.

Eager evaluation: An expression is evaluated as soon as it gets bound to a variable.

Lazy evaluation: Non-strict evaluation which avoids repeating computation.

Call-by-value: The argument expression is evaluated, and the resulting value is bound to the corresponding variable in the function (frequently by copying the value into a new memory region).

Call-by-reference: A function receives an implicit reference to a variable used as argument, rather than a copy of its value. In purely functional languages there is no difference between the two strategies, so they are typically described as call-by-value even though implementations use call-by-reference internally for efficiency. Call-by-value languages like C and OCaml support explicit references (objects that refer to other objects), and these can be used to simulate call-by-reference.

Normal order: Start computing function bodies before evaluating their arguments. Do not even wait for arguments if they are not needed.

Call-by-name: Arguments are substituted directly into the function body and then left to be evaluated whenever they appear in the function.

Call-by-need: If the function argument is evaluated, that value is stored for subsequent uses.

Almost all languages do not compute inside the body of an un-applied function, but with curried functions you can pre-compute data before all arguments are provided (recall the `search_bible` example from earlier lectures).

In eager / call-by-value languages we can simulate call-by-name by taking a function to compute the value as an argument instead of the value directly.

“Our” languages have a `unit` type with a single value () specifically for use as throw-away arguments. Scala has built-in support for call-by-name (i.e. direct, without the need to build argument functions).

ML languages have built-in support for lazy evaluation, while Haskell has built-in support for eager evaluation (to override the default laziness).

7.2 Call-by-name: Streams

Call-by-name is useful not only for implementing flow control:

```
let if_then_else cond e1 e2 =
  match cond with
  | true -> e1 ()
  | false -> e2 ()
```

but also for arguments of value constructors, i.e. for data structures.

Streams are lists with call-by-name tails:

```
type 'a stream = SNil | SCons of 'a * (unit -> 'a stream)
```

Reading from a stream into a list:

```
let rec stake n = function
  | SCons (a, s) when n > 0 -> a :: (stake (n-1) (s()))
  | _ -> []
```

Streams can easily be infinite:

```
let rec s_ones = SCons (1, fun () -> s_ones)

let rec s_from n =
  SCons (n, fun () -> s_from (n+1))
```

7.2.1 Stream Operations Streams admit list-like operations:

```
let rec smap f = function
  | SNil -> SNil
  | SCons (a, s) -> SCons (f a, fun () -> smap f (s()))

let rec szip = function
  | SNil, SNil -> SNil
  | SCons (a1, s1), SCons (a2, s2) ->
    SCons ((a1, a2), fun () -> szip (s1(), s2()))
  | _ -> raise (Invalid_argument "szip")
```

Streams can provide scaffolding for recursive algorithms. Consider the Fibonacci sequence:

```
let rec sfib =
  SCons (1, fun () -> smap (fun (a,b) -> a+b)
    (szip (sfib, SCons (1, fun () -> sfib))))
```

This definition creates a stream where each element is computed by adding pairs from the current stream and itself shifted by one position:

sfib	1	2	3	5	8	13	...
sfib	1	2	3	5	8	13	...
shifted	1	1	2	3	5	8	...

The + operation between corresponding elements produces the next values.

7.2.2 Streams and Input-Output Streams are less functional than could be expected in the context of input-output effects:

```
let file_stream name =
  let ch = open_in name in
  let rec ch_read_line () =
    try SCons (input_line ch, ch_read_line)
    with End_of_file -> SNil in
    ch_read_line ()
```

OCaml Batteries uses a stream type `enum` for interfacing between various sequence-like data types. The safest way to use streams is in a *linear / ephemeral* manner: every value used only once. Streams minimize space consumption at the expense of time for recomputation.

7.3 Lazy Values

Lazy evaluation is more general than call-by-need as any value can be lazy, not only a function parameter.

A *lazy value* is a value that “holds” an expression until its result is needed, and from then on it “holds” the result. It is also called a *suspension*. If it holds the expression (not yet evaluated), it is called a *thunk*.

In OCaml, we build lazy values explicitly. In Haskell, all values are lazy but functions can have call-by-value parameters which “need” the argument.

To create a lazy value: `lazy expr` – where `expr` is the suspended computation.

Two ways to use a lazy value (be careful when the result is computed!): - In expressions: `Lazy.force l_expr` - In patterns: `match l_expr with lazy v -> ...` - Syntactically `lazy` behaves like a data constructor.

7.3.1 Lazy Lists Lazy lists are defined as:

```
type 'a llist = LNil | LCons of 'a * 'a llist Lazy.t
```

Reading from a lazy list into a list:

```
let rec ltake n = function
| LCons (a, lazy l) when n > 0 -> a :: (ltake (n-1) l)
| _ -> []
```

Lazy lists can easily be infinite:

```
let rec l_ones = LCons (1, lazy l_ones)

let rec l_from n = LCons (n, lazy (l_from (n+1)))
```

Read once, access multiple times (unlike streams):

```
let file_llist name =
  let ch = open_in name in
  let rec ch_read_line () =
    try LCons (input_line ch, lazy (ch_read_line ()))
    with End_of_file -> LNil in
  ch_read_line ()
```

7.3.2 Lazy List Operations

```
let rec lzip = function
| LNil, LNil -> LNil
| LCons (a1, l1), LCons (a2, l2) ->
  LCons ((a1, a2), lazy (
    lzip (Lazy.force l1, Lazy.force l2)))
| _ -> raise (Invalid_argument "lzip")

let rec lmap f = function
| LNil -> LNil
| LCons (a, l) ->
  LCons (f a, lazy (lmap f (Lazy.force l)))
```

Using these operations, we can define the factorial sequence elegantly:

```
let posnums = l_from 1

let rec lfact =
  LCons (1, lazy (lmap (fun (a,b) -> a*b)
    (lzip (lfact, posnums))))
```

This produces: 1, 1, 2, 6, 24, 120, ... where each element is the product of the previous factorial and the corresponding positive integer:

lfact	1	1	2	6	24	120	...
lfact	1	1	2	6	24	120	...
posnums	1	2	3	4	5	6	...

The $*$ operation between corresponding elements produces the next values.

7.4 Power Series and Differential Equations

This section presents an application of lazy lists to power series computation and solving differential equations through power series. The differential equations idea is due to Henning Thieleemann.

The expression $P(x) = \sum_{i=0}^n a_i x^i$ defines a polynomial for $n < \infty$ and a power series for $n = \infty$.

If we define:

```
let rec lfold_right f l base =
  match l with
  | LNil -> base
  | LCons (a, lazy l) -> f a (lfold_right f l base)
```

then we can compute polynomials using Horner's method:

```
let horner x l =
  lfold_right (fun c sum -> c +. x *. sum) l 0.
```

But this will not work for infinite power series! Does it make sense to compute the value at x of a power series? Does it make sense to fold an infinite list?

If the power series converges for $x > 1$, then when the elements a_n get small, the remaining sum $\sum_{i=n}^{\infty} a_i x^i$ is also small.

`lfold_right` falls into an infinite loop on infinite lists. We need call-by-name / call-by-need semantics for the argument function `f`:

```
let rec lazy_foldr f l base =
  match l with
  | LNil -> base
  | LCons (a, ll) ->
    f a (lazy_foldr f (Lazy.force ll) base))
```

We need a stopping condition in the Horner algorithm step:

```
let lhorner x l =                                     (* This is a bit of a hack, *)
  let upd c sum =                                     (* we hope to "hit" the interval (0, epsilon]. *)
    if c = 0. || abs_float c > epsilon_float
    then c +. x *. Lazy.force sum
    else 0. in
  lazy_foldr upd l 0.
```

```

let inv_fact = lmap (fun n -> 1. /. float_of_int n) lfact
let e = lhorner 1. inv_fact

7.4.1 Power Series / Polynomial Operations

let rec add xs ys =
  match xs, ys with
  | LNil, _ -> ys
  | _, LNil -> xs
  | LCons (x,xs), LCons (y,ys) ->
    LCons (x +. y, lazy (add (Lazy.force xs) (Lazy.force ys)))

let rec sub xs ys =
  match xs, ys with
  | LNil, _ -> lmap (fun x -> -.x) ys
  | _, LNil -> xs
  | LCons (x,xs), LCons (y,ys) ->
    LCons (x -. y, lazy (add (Lazy.force xs) (Lazy.force ys)))

let scale s = lmap (fun x -> s *. x)

let rec shift n xs =
  if n = 0 then xs
  else if n > 0 then LCons (0., lazy (shift (n-1) xs))
  else match xs with
    | LNil -> LNil
    | LCons (0., lazy xs) -> shift (n+1) xs
    | _ -> failwith "shift: fractional division"

let rec mul xs = function
  | LNil -> LNil
  | LCons (y, ys) ->
    add (scale y xs) (LCons (0., lazy (mul xs (Lazy.force ys)))))

let rec div xs ys =
  match xs, ys with
  | LNil, _ -> LNil
  | LCons (0., xs'), LCons (0., ys') ->
    div (Lazy.force xs') (Lazy.force ys')
  | LCons (x, xs'), LCons (y, ys') ->
    let q = x /. y in
    LCons (q, lazy (div (sub (Lazy.force xs')
                           (scale q (Lazy.force ys')))) ys'))
  | LCons _, LNil -> failwith "div: division by zero"

```

```

let integrate c xs =
  LCons (c, lazy (lmap (uncurry (/ .)) (lzip (xs, posnums)))))

let ltail = function
  | LNil -> invalid_arg "ltail"
  | LCons (_, lazy tl) -> tl

let differentiate xs =
  lmap (uncurry ( * .)) (lzip (ltail xs, posnums))

```

7.4.2 Differential Equations Consider the differential equations for sine and cosine:

$$\frac{d \sin x}{dx} = \cos x, \quad \frac{d \cos x}{dx} = -\sin x, \quad \sin 0 = 0, \quad \cos 0 = 1$$

We will solve the corresponding integral equations. We cannot define the integral by direct recursion like this:

```

let (~-:) = lmap (fun x -> -.x) (* Unary negation for series *)

let rec sin = integrate (of_int 0) cos
and cos = integrate (of_int 1) (~-:sin)

```

Unfortunately this fails with: `Error: This kind of expression is not allowed as right-hand side of 'let rec'`

Even changing the second argument of `integrate` to call-by-need does not help, because OCaml cannot represent the values that `sin` and `cos` refer to at the point of their definition.

We need to inline a bit of `integrate` so that OCaml knows how to start building the recursive structure:

```

let integ xs = lmap (uncurry (/ .)) (lzip (xs, posnums))

let rec sin = LCons (of_int 0, lazy (integ cos))
and cos = LCons (of_int 1, lazy (integ (~-:sin)))

```

The complete example would look much more elegant in Haskell, where all values are lazy by default.

Although this approach is not limited to linear equations, equations like Lotka-Volterra or Lorentz are not “solvable” this way – computed coefficients quickly grow instead of quickly falling.

Drawing functions work like in the previous lecture, but with open curves:

```

let plot_1D f ~w ~scale ~t_beg ~t_end =
  let dt = (t_end -. t_beg) /. of_int w in

```

```
Array.init w (fun i ->
  let y = lhorner (dt *. of_int i) f in
  i, to_int (scale *. y))
```

7.5 Arbitrary Precision Computation

Putting together the power series computation with floating-point numbers reveals drastic numerical errors for large x . Floating-point numbers have limited precision, and we break out of Horner method computations too quickly.

For infinite precision on rational numbers we use the `nums` library – but it does not help by itself.

We need to generate a sequence of approximations to the power series limit at x :

```
let infhorner x l =
  let upd c sum =
    LCons (c, lazy (lmap (fun apx -> c +. x *. apx)
                           (Lazy.force sum))) in
  lazy_foldr upd l (LCons (of_int 0, lazy LNil))
```

Find where the series converges – as far as a given test is concerned:

```
let rec exact f = function
  (* We arbitrarily decide that convergence is *)
  | LNil -> assert false
  (* when three consecutive results are the same. *)
  | LCons (x0, lazy (LCons (x1, lazy (LCons (x2, _))))))
    when f x0 = f x1 && f x0 = f x2 -> f x0
  | LCons (_, lazy t1) -> exact f t1
```

Draw the pixels of the graph at exact coordinates:

```
let plot_1D f ~w ~h0 ~scale ~t_beg ~t_end =
  let dt = (t_end -. t_beg) /. of_int w in
  let eval = exact (fun y -> to_int (scale *. y)) in
  Array.init w (fun i ->
    let y = infhorner (t_beg +. dt *. of_int i) f in
    i, h0 + eval y)
```

If a power series had every third term contributing we would have to check three terms in the function `exact`. We could also test for `f x0 = f x1 && not (x0 =. x1)` like in `lhorner`.

7.5.1 Example: Nuclear Chain Reaction Consider a nuclear chain reaction where substance A decays into B, which decays into C. The differential equations are:

$$\frac{dN_A}{dt} = -\lambda_A N_A, \quad \frac{dN_B}{dt} = \lambda_A N_A - \lambda_B N_B$$

```
let n_chain ~nA0 ~nB0 ~lA ~lB =
```

```

let rec nA =
  LCons (nA0, lazy (integ (~-.1A *:. nA)))
and nB =
  LCons (nB0, lazy (integ (~-.1B *:. nB +: 1A *:. nA))) in
nA, nB

```

(See Radioactive decay chain processes for more information.)

7.6 Circular Data Structures: Double-Linked Lists

Without delayed computation, the ability to define data structures with referential cycles is very limited.

Double-linked lists contain such cycles between any two nodes even if they are not cyclic when following only *forward* or *backward* links:

```
+-----+ +-----+ +-----+ +-----+ +-----+
| DLNil | <-> | a1 | <-> | a2 | <-> | a3 | <-> | DLNil |
+-----+ +-----+ +-----+ +-----+ +-----+
```

We need to “break” the cycles by making some links lazy:

```

type 'a dllist =
  DLNil | DLCons of 'a dllist Lazy.t * 'a * 'a dllist

let rec dldrop n l =
  match l with
  | DLCons (_, x, xs) when n > 0 ->
    dldrop (n-1) xs
  | _ -> l

```

Creating a double-linked list from a regular list:

```

let dllist_of_list l =
  let rec dllist prev l =
    match l with
    | [] -> DLNil
    | x::xs ->
      let rec cell =
        lazy (DLCons (prev, x, dllist cell xs)) in
        Lazy.force cell in
      dllist (lazy DLNil) l

```

Taking elements going forward:

```

let rec dltake n l =
  match l with
  | DLCons (_, x, xs) when n > 0 ->
    x :: dltake (n-1) xs
  | _ -> []

```

Taking elements going backward:

```
let rec dlbackwards n l =
  match l with
  | DLCons (lazy xs, x, _) when n > 0 ->
    x :: dlbackwards (n-1) xs
  | _ -> []
```

7.7 Input-Output Streams

The stream type used a throwaway argument to make a suspension:

```
type 'a stream = SNil | SCons of 'a * (unit -> 'a stream)
```

What if we take a real argument?

```
type ('a, 'b) iostream =
  EOS | More of 'b * ('a -> ('a, 'b) iostream)
```

This is a stream that for a single input value produces an output value.

```
type 'a istream = (unit, 'a) iostream (* Input stream produces output when "asked". *)
type 'a ostream = ('a, unit) iostream (* Output stream consumes provided input. *)
```

(The confusion arises from adapting the *input file / output file* terminology, also used for streams.)

We can compose streams: directing output of one to input of another.

```
let rec compose sf sg =
  match sg with
  | EOS -> EOS                                     (* No more output. *)
  | More (z, g) ->
    match sf with
    | EOS -> More (z, fun _ -> EOS)             (* No more input "processing power". *)
    | More (y, f) ->
      let update x = compose (f x) (g y) in
      More (z, update)
```

Every box has one incoming and one outgoing wire. Notice how the output stream is ahead of the input stream.

7.8 Pipes

We need a more flexible input-output stream definition: - Consume several inputs to produce a single output. - Produce several outputs after a single input (or even without input). - No need for a dummy when producing output requires input.

After Haskell, we call the data structure **pipe**:

```

type ('a, 'b) pipe =
  EOP
  | Yield of 'b * ('a, 'b) pipe
  | Await of ('a -> ('a, 'b) pipe)

```

Again, we can have producing output only *input pipes* and consuming input only *output pipes*:

```

type 'a ipipe = (unit, 'a) pipe
type void
type 'a opipe = ('a, void) pipe

```

Why **void** rather than **unit**, and why only for **opipe**? Because an output pipe never yields values – if it used **unit** as the output type, it could still yield () values, but with the abstract **void** type, it cannot yield anything.

7.8.1 Pipe Composition Composition of pipes is like “concatenating them in space” or connecting boxes:

```

let rec compose pf pg =
  match pg with
    | EOP -> EOP
      (* Done producing results. *)
    | Yield (z, pg') -> Yield (z, compose pf pg') (* Ready result. *)
    | Await g ->
      match pf with
        | EOP -> EOP
          (* End of input. *)
        | Yield (y, pf') -> compose pf' (g y) (* Compute next result. *)
        | Await f ->
          let update x = compose (f x) pg in
          Await update
          (* Wait for more input. *)

let (>->) pf pg = compose pf pg

```

Appending pipes means “concatenating them in time” or adding more fuel to a box:

```

let rec append pf pg =
  match pf with
    | EOP -> pg
      (* When pf runs out, use pg. *)
    | Yield (z, pf') -> Yield (z, append pf' pg)
    | Await f ->
      let update x = append (f x) pg in
      Await update
      (* If pf awaits input, continue when it comes. *)

```

Append a list of ready results in front of a pipe:

```

let rec yield_all l tail =
  match l with

```

```

| [] -> tail
| x::xs -> Yield (x, yield_all xs tail)

```

Iterate a pipe (**not functional** – performs side effects):

```

let rec iterate f : 'a opipe =
  Await (fun x -> let () = f x in iterate f)

```

7.9 Example: Pretty-Printing

Print a hierarchically organized document with a limited line width.

```

type doc =
  Text of string | Line | Cat of doc * doc | Group of doc

let (++) d1 d2 = Cat (d1, Cat (Line, d2))
let (!) s = Text s

let test_doc =
  Group (!"Document" ++
         Group (!"First part" ++ !"Second part"))

```

Example output with different widths:

```

# let () = print_endline (pretty 30 test_doc);;
Document
First part Second part

# let () = print_endline (pretty 20 test_doc);;
Document
First part
Second part

# let () = print_endline (pretty 60 test_doc);;
Document First part Second part

```

7.9.1 Straightforward Solution

```

let pretty w d =                                     (* Allowed width of line w. *)
  let rec width = function                         (* Total length of subdocument. *)
    | Text z -> String.length z
    | Line -> 1
    | Cat (d1, d2) -> width d1 + width d2
    | Group d -> width d in
  let rec format f r = function                   (* Remaining space r. *)
    | Text z -> z, r - String.length z
    | Line when f -> " ", r-1                    (* If not f then line breaks. *)
    | Line -> "\n", w
    | Cat (d1, d2) ->

```

```

let s1, r = format f r d1 in
let s2, r = format f r d2 in
s1 ^ s2, r                                     (* If following group fits, then without line b
| Group d -> format (f || width d <= r) r d in
fst (format false w d)

```

7.9.2 Stream-Based Solution

Working with a stream of nodes:

```

type ('a, 'b) doc_e =                         (* Annotated nodes, special for group beginning
    TE of 'a * string | LE of 'a | GBeg of 'b | GEnd of 'a

```

Normalize a subdocument – remove empty groups:

```

let rec norm = function
| Group d -> norm d
| Text "" -> None
| Cat (Text "", d) -> norm d
| d -> Some d

```

Generate the stream by infix traversal:

```

let rec gen = function
| Text z -> Yield (TE (((),z), EOP)
| Line -> Yield (LE (), EOP)
| Cat (d1, d2) -> append (gen d1) (gen d2)
| Group d ->
    match norm d with
    | None -> EOP
    | Some d ->
        Yield (GBeg (), append (gen d) (Yield (GEnd (), EOP)))

```

Compute lengths of document prefixes, i.e. the position of each node counting by characters from the beginning of document:

```

let rec docpos curpos =
    Await (function                                (* We input from a doc_e pipe *)
    | TE (_, z) ->
        Yield (TE (curpos, z),                      (* and output doc_e annotated with position. *)
                docpos (curpos + String.length z))
    | LE _ ->                                     (* Space and line breaks increase position by 1 *)
        Yield (LE curpos, docpos (curpos + 1))
    | GBeg _ ->                                    (* Groups do not increase position. *)
        Yield (GBeg curpos, docpos curpos)
    | GEnd _ ->
        Yield (GEnd curpos, docpos curpos))

let docpos = docpos 0                           (* The whole document starts at 0. *)

```

Put the end position of the group into the group beginning marker, so that we can know whether to break it into multiple lines:

```
let rec grends grstack =
  Await (function
  | TE _ | LE _ as e ->
    (match grstack with
    | [] -> Yield (e, grends [])
      (* We can yield only when *)
    | gr::grs -> grends ((e::gr)::grs))
      (* no group is waiting. *)
    | GBeg _ -> grends ([]::grstack)
      (* Wait for end of group. *)
    | GEnd endp ->
      match grstack with
        (* End the group on top of stack. *)
      | [] -> failwith "grends: unmatched group end marker"
      | [gr] ->
        (* Top group -- we can yield now. *)
        yield_all
        (GBeg endp)::List.rev (GEnd endp::gr)
        (grends [])
      | gr::par::grs ->
        let par = GEnd endp::gr @ [GBeg endp] @ par in
        grends (par::grs)
        (* Could use catenable lists above. *)
```

That's waiting too long! We can stop waiting when the width of a group exceeds the line limit. GBeg will not store end of group when it is irrelevant:

```
type grp_pos = Pos of int | Too_far

let rec grends w grstack =
  let flush tail =
    yield_all
    (rev_concat_map ~prep:(GBeg Too_far) snd grstack)
    tail in
    (* Above: concatenate in rev. with prep before *)
  Await (function
  | TE (curp, _) | LE curp as e ->
    (match grstack with
    | [] -> Yield (e, grends w [])
      (* Remember beginning of groups in the stack. *)
    | (begp, _)::_ when curp-begp > w ->
      flush (Yield (e, grends w [])))
    | (begp, gr)::grs -> grends w ((begp, e::gr)::grs))
  | GBeg begp -> grends w ((begp, [])::grstack)
  | GEnd endp as e ->
    match grstack with
    | [] -> Yield (e, grends w [])
      (* No longer fail when the stack is empty -- *)
      (* could have been flushed. *)
    | (begp, _)::_ when endp-begp > w ->
      flush (Yield (e, grends w [])))
    | _, gr] ->
      yield_all
      (GBeg (Pos endp)::List.rev (GEnd endp::gr))
```

```

(grends w [])
| (_, gr)::(par_begp, par)::grs ->
  let par =
    GEnd endp::gr @ [GBeg (Pos endp)] @ par in
    grends w ((par_begp, par)::grs))

let grends w = grends w []                      (* Initial stack is empty. *)

```

Finally we produce the resulting stream of strings:

```

let rec format w (inline, endlpos as st) = (* State: the stack of *)
  Await (function                                (* "group fits in line"; position where end of *)
  | TE (_, z) -> Yield (z, format w st)
  | LE p when List.hd inline ->
    Yield (" ", format w st)                  (* After return, line has w free space. *)
  | LE p -> Yield ("\n", format w (inline, p+w))
  | GBeg Too_far ->                         (* Group with end too far is not inline. *)
    format w (false::inline, endlpos)
  | GBeg (Pos p) ->                          (* Group is inline if it ends soon enough. *)
    format w ((p<=endlpos)::inline, endlpos)
  | GEnd _ -> format w (List.tl inline, endlpos))

let format w = format w ([false], w)          (* Break lines outside of groups. *)

```

Put the pipes together:

```

+-----+ +-----+ +-----+ +-----+ +-----+
| gen doc| --> |docpos| --> |grends w| --> |format w| --> |iterate print_s|
+-----+ +-----+ +-----+ +-----+ +-----+

```

7.9.3 Factored Solution Factorize `format` so that various line breaking styles can be plugged in:

```

let rec breaks w (inline, endlpos as st) =
  Await (function
  | TE _ as e -> Yield (e, breaks w st)
  | LE p when List.hd inline ->
    Yield (TE (p, " "), breaks w st)
  | LE p as e -> Yield (e, breaks w (inline, p+w))
  | GBeg Too_far as e ->
    Yield (e, breaks w (false::inline, endlpos))
  | GBeg (Pos p) as e ->
    Yield (e, breaks w ((p<=endlpos)::inline, endlpos))
  | GEnd _ as e ->
    Yield (e, breaks w (List.tl inline, endlpos)))

let breaks w = breaks w ([false], w)

```

```

let rec emit =
  Await (function
  | TE (_, z) -> Yield (z, emit)
  | LE _ -> Yield ("\n", emit)
  | GBeg _ | GEnd _ -> emit)

let pretty_print w doc =
  gen doc >-> docpos >-> grends w >-> breaks w >->
  emit >-> iterate print_string

```

7.10 Exercises

Exercise 1: My first impulse was to define lazy list functions as follows:

```

let rec wrong_lzip = function
| LNil, LNil -> LNil
| LCons (a1, lazy l1), LCons (a2, lazy l2) ->
  LCons ((a1, a2), lazy (wrong_lzip (l1, l2)))
| _ -> raise (Invalid_argument "lzip")

let rec wrong_lmap f = function
| LNil -> LNil
| LCons (a, lazy l) -> LCons (f a, lazy (wrong_lmap f l))

```

What is wrong with these definitions – for which edge cases do they not work as intended?

Exercise 2: Cyclic lazy lists.

1. Implement a function `cycle : 'a list -> 'a llist` that creates a lazy list with elements from a standard list, and the whole list as the tail after the last element from the input list: `[a1; a2; ...; aN]` maps to a cyclic structure where `aN` points back to `a1`. Your function `cycle` can either return `LNil` or fail for an empty list as argument.
2. Note that `inv_fact` from the lecture defines the power series for the $\exp(\cdot)$ function ($\exp(x) = e^x$). Using `cycle` and `inv_fact`, define the power series for $\sin(\cdot)$ and $\cos(\cdot)$, and draw their graphs using helper functions from the lecture script `Lec7.ml`.

Exercise 3: Modify one of the puzzle solving programs (either from the previous lecture or from your previous homework) to work with lazy lists. Implement the necessary higher-order lazy list functions. Check that indeed displaying only the first solution when there are multiple solutions in the result takes shorter than computing solutions by the original program.

Exercise 4: *Hamming's problem.* Generate in increasing order the numbers of the form $2^{a_1}3^{a_2}5^{a_3}\dots p_k^{a_k}$, that is numbers not divisible by prime numbers greater than the k th prime number.

In the original Hamming's problem posed by Dijkstra, $k = 3$, which is related to regular numbers.

Starter code is available in the lecture script Lec7.ml:

```

let rec lfilter f = function
| LNil -> LNil
| LCons (n, ll) ->
    if f n then LCons (n, lazy (lfilter f (Lazy.force ll)))
    else lfilter f (Lazy.force ll)

let primes =
  let rec sieve = function
  | LCons(p, nf) ->
      LCons(p, lazy (sieve (sift p (Lazy.force nf))))
  | LNil -> failwith "Impossible! Internal error."
  and sift p = lfilter (fun n -> n mod p <> 0)
  in sieve (l_from 2)

let times ll n = lmap (fun i -> i * n) ll

let rec merge xs ys =
  match xs, ys with
  | LCons (x, lazy xr), LCons (y, lazy yr) ->
      if x < y then LCons (x, lazy (merge xr ys))
      else if x > y then LCons (y, lazy (merge xs yr))
      else LCons (x, lazy (merge xr yr))
  | r, LNil | LNil, r -> r

let hamming k =
  let pr = ltake k primes in
  let rec h = LCons (1, lazy (
    (* TODO *)
  )) in h

```

Exercise 5: Modify `format` and/or `breaks` to use just a single number instead of a stack of booleans to keep track of what groups should be inlined.

Exercise 6: Add `indentation` to the pretty-printer for groups: if a group does not fit in a single line, its consecutive lines are indented by a given amount `tab` of spaces deeper than its parent group lines would be. For comparison, let's do several implementations.

1. Modify the straightforward implementation of `pretty`.
2. Modify the first pipe-based implementation of `pretty` by modifying the `format` function.
3. Modify the second pipe-based implementation of `pretty` by modifying the `breaks` function. Recover the positions of elements – the number of

characters from the beginning of the document – by keeping track of the growing offset.

4. (Harder) Modify a pipe-based implementation to provide a different style of indentation: indent the first line of a group, when the group starts on a new line, at the same level as the consecutive lines (rather than at the parent level of indentation).

Exercise 7: Write a pipe that takes document elements annotated with linear position, and produces document elements annotated with (line, column) coordinates.

Write another pipe that takes so annotated elements and adds a line number indicator in front of each line. Do not update the column coordinate. Test the pipes by plugging them before the `emit` pipe.

```
1: first line
2: second line, etc.
```

Exercise 8: Write a pipe that consumes document elements `doc_e` and yields the toplevel subdocuments `doc` which would generate the corresponding elements.

You can modify the definition of documents to allow annotations, so that the element annotations are preserved (`gen` should ignore annotations to keep things simple):

```
type 'a doc =
  Text of 'a * string | Line of 'a | Cat of doc * doc | Group of 'a * doc
```

Exercise 9: (Harder) Design and implement a way to duplicate arrows outgoing from a pipe-box, that would memoize the stream, i.e. not recompute everything “upstream” for the composition of pipes. Such duplicated arrows would behave nicely with pipes reading from files.

Chapter 8: Monads

This chapter explores one of functional programming’s most powerful abstractions: monads. We begin with list comprehensions, introduce monadic concepts, examine monad laws and the monad-plus extension, then work through various monad instances including state, exception, and probability monads. We conclude with monad transformers and cooperative lightweight threads.

8.1 List Comprehensions

Recall the somewhat awkward syntax we used in the Countdown Problem example from earlier chapters. The brute-force generation of expressions looked like this:

```
let combine l r =
  List.map (fun o -> App (o, l, r)) [Add; Sub; Mul; Div]
```

```

let rec exprs = function
| [] -> []
| [n] -> [Val n]
| ns ->
    split ns |-> (fun (ls, rs) ->
        exprs ls |-> (fun l ->
            exprs rs |-> (fun r ->
                combine l r)))

```

And the generate-and-test scheme used:

```

let guard p e = if p e then [e] else []

let solutions ns n =
    choices ns |-> (fun ns' ->
        exprs ns' |->
        guard (fun e -> eval e = Some n))

```

We introduced the operator |-> defined as:

```
let ( |-> ) x f = concat_map f x
```

We can do much better with list comprehensions syntax extension. In older versions of OCaml with Camlp4, this was loaded via:

```
#load "dynlink.cma";;
#load "camlp4o.cma";;
#load "Camlp4Parsers/Camlp4ListComprehension.cmo";;
```

With list comprehensions, we can write:

```
let test = [i * 2 | i <- from_to 2 22; i mod 3 = 0]
```

The translation rules for list comprehensions are:

- [expr |] translates to [expr]
- [expr | v <- generator; more] translates to generator |-> (fun v -> [expr | more])
- [expr | condition; more] translates to if condition then [expr | more] else []

Revisiting Countdown with List Comprehensions The brute-force generation becomes cleaner:

```

let rec exprs = function
| [] -> []
| [n] -> [Val n]
| ns ->
    [App (o, l, r) | (ls, rs) <- split ns;
     l <- exprs ls; r <- exprs rs;
     o <- [Add; Sub; Mul; Div]]

```

And the generate-and-test scheme simplifies to:

```
let solutions ns n =
  [e | ns' <- choices ns;
    e <- exprs ns'; eval e = Some n]
```

More List Comprehension Examples Computing subsequences using list comprehensions (with some garbage generation):

```
let rec subseqs l =
  match l with
  | [] -> []
  | x::xs -> [ys | px <- subseqs xs; ys <- [px; x::px]]
```

Computing permutations via insertion:

```
let rec insert x = function
  | [] -> [[x]]
  | y::ys' as ys ->
    (x::ys) :: [y::zs | zs <- insert x ys']

let rec ins_perms = function
  | [] -> []
  | x::xs -> [zs | ys <- ins_perms xs; zs <- insert x ys]
```

And via selection:

```
let rec select = function
  | [x] -> [x, []]
  | x::xs -> (x, xs) :: [y, x::ys | y, ys <- select xs]

let rec sel_perms = function
  | [] -> []
  | xs -> [x::ys | x, xs' <- select xs; ys <- sel_perms xs']
```

8.2 Generalized Comprehensions: Binding Operators

OCaml 5 introduced **binding operators** that provide a clean, native syntax for monadic computations. Instead of external syntax extensions like the old `pa_monad`, we can define custom `let*` and `let+` operators that integrate naturally with the language.

For the list monad, we define these binding operators:

```
let ( let* ) x f = concat_map f x      (* bind *)
let ( let+ ) x f = List.map f x        (* map/fmap *)
let ( and* ) x y = concat_map (fun a -> List.map (fun b -> (a, b)) y) x
let ( and+ ) = ( and* )
let return x = [x]
let fail = []
```

With these operators, the expression generation code becomes:

```
let rec exprs = function
| [] -> []
| [n] -> [Val n]
| ns ->
  let* (ls, rs) = split ns in
  let* l = exprs ls in
  let* r = exprs rs in
  let* o = [Add; Sub; Mul; Div] in
  [App (o, l, r)]
```

Note that unlike the old `perform` syntax where we used `<--` for binding, we now use `let*` followed by `=` and must explicitly write `in` before the continuation.

The `let*` syntax does not directly support guards. If we try to write:

```
let solutions ns n =
  let* ns' = choices ns in
  let* e = exprs ns' in
  eval e = Some n; (* Error! *)
  e
```

We get an error because it expects a list, not a boolean. We can work around this by deciding whether to return anything:

```
let solutions ns n =
  let* ns' = choices ns in
  let* e = exprs ns' in
  if eval e = Some n then [e] else []
```

For a general guard check function, we define:

```
let guard p = if p then [] else []
```

And then:

```
let solutions ns n =
  let* ns' = choices ns in
  let* e = exprs ns' in
  let* () = guard (eval e = Some n) in
  [e]
```

8.3 Monads

A monad is a polymorphic type `'a monad` (or `'a Monad.t`) that supports at least two operations:

- `bind` : `'a monad -> ('a -> 'b monad) -> 'b monad`
- `return` : `'a -> 'a monad`
- The infix `>>=` is commonly used for `bind`: `let (>>=) a b = bind a b`

With OCaml 5's binding operators, we define `let*` as an alias for `bind`:

```
let bind a b = concat_map b a
let return x = [x]
let ( let* ) = bind

let solutions ns n =
  let* ns' = choices ns in
  let* e = exprs ns' in
  let* () = guard (eval e = Some n) in
  return e
```

Why does `guard` look this way? Let us examine:

```
let fail = []
let guard p = if p then return () else fail
```

Steps in monadic computation are composed with `let*` (or `>>=`, like `|->` for lists). The key insight is:

- `let* _ = [] in ...` does not produce anything – as needed by guarding
- `let* _ = [()` in ... becomes `(fun _ -> ...) ()` which simply continues the computation unchanged

Throwing away the binding argument is common. With binding operators, we can use `let* () = ...` or `let* _ = ...`:

```
let (>>=) a b = bind a b
let (>>) m f = m >>= (fun _ -> f)
```

The Binding Operator Syntax OCaml 5's binding operators translate as follows:

Source	Translation
<code>let* x = exp in body</code>	<code>bind exp (fun x -> body)</code>
<code>let+ x = exp in body</code>	<code>map (fun x -> body) exp</code>
<code>let* () = exp in body</code>	<code>bind exp (fun () -> body)</code>
<code>let* x = e1 and* y = e2 in body</code>	<code>bind (and* e1 e2) (fun (x, y) -> body)</code>

The binding operators `let*`, `let+`, `and*`, and `and+` must be defined in scope. These are regular OCaml operators and require no syntax extensions.

For pattern matching in bindings, if the pattern is refutable (can fail to match), the monadic operation should handle the failure appropriately.

8.4 Monad Laws

A parametric data type is a monad only if its `bind` and `return` operations meet these axioms:

$$\begin{aligned} \text{bind } (\text{return } a) f &\approx f a && \text{(left identity)} \\ \text{bind } a (\lambda x. \text{return } x) &\approx a && \text{(right identity)} \\ \text{bind } (\text{bind } a (\lambda x. b)) (\lambda y. c) &\approx \text{bind } a (\lambda x. \text{bind } b (\lambda y. c)) && \text{(associativity)} \end{aligned}$$

You should verify that these laws hold for our list monad:

```
let bind a b = concat_map b a
let return x = [x]
```

8.5 Monoid Laws and Monad-Plus

A monoid is a type with at least two operations:

- `mzero` : '`a` monoid
- `mplus` : '`a` monoid \rightarrow '`a` monoid

that meet these laws:

$$\begin{aligned} \text{mplus mzero } a &\approx a && \text{(left identity)} \\ \text{mplus } a \text{ mzero} &\approx a && \text{(right identity)} \\ \text{mplus } a (\text{mplus } b \text{ } c) &\approx \text{mplus } (\text{mplus } a \text{ } b) \text{ } c && \text{(associativity)} \end{aligned}$$

We define `fail` as a synonym for `mzero` and infix `++` for `mplus`.

Fusing monads and monoids gives the most popular general flavor of monads, which we call **monad-plus** after Haskell. Monad-plus requires additional axioms relating its “addition” and “multiplication”:

$$\begin{aligned} \text{bind mzero } f &\approx \text{mzero} \\ \text{bind } m (\lambda x. \text{mzero}) &\approx \text{mzero} \end{aligned}$$

Using infix notation with `\oplus` for `mplus`, `0` for `mzero`, `\triangleright` for `bind`, and `1` for `return`, the complete monad-plus axioms are:

$$\begin{aligned}
\mathbf{0} \oplus a &\approx a \\
a \oplus \mathbf{0} &\approx a \\
a \oplus (b \oplus c) &\approx (a \oplus b) \oplus c \\
\mathbf{1} x \triangleright f &\approx f x \\
a \triangleright \lambda x. \mathbf{1} x &\approx a \\
(a \triangleright \lambda x. b) \triangleright \lambda y. c &\approx a \triangleright (\lambda x. b \triangleright \lambda y. c) \\
\mathbf{0} \triangleright f &\approx \mathbf{0} \\
a \triangleright (\lambda x. \mathbf{0}) &\approx \mathbf{0}
\end{aligned}$$

The list type has a natural monad and monoid structure:

```

let mzero = []
let mplus = (@)
let bind a b = concat_map b a
let return a = [a]

```

We can define in any monad-plus:

```

let fail = mzero
let failwith _ = fail
let (++) = mplus
let (>>=) a b = bind a b
let guard p = if p then return () else fail

```

8.6 Backtracking: Computation with Choice

We have seen `mzero` (i.e., `fail`) in the countdown problem. What about `mplus`? Here is an example from a puzzle solver:

```

let find_to_eat n island_size num_islands empty_cells =
  let honey = honey_cells n empty_cells in

  let rec find_board s =
    match visit_cell s with
    | None ->
        let* () = guard (s.been_islands = num_islands) in
        return s.eaten
    | Some (cell, s) ->
        let* s = find_island cell (fresh_island s) in
        let* () = guard (s.been_size = island_size) in
        find_board s

  and find_island current s =
    let s = keep_cell current s in
    neighbors n empty_cells current

```

```

|> foldM
  (fun neighbor s -
    if CellSet.mem neighbor s.visited then return s
    else
      let choose_eat =
        if s.more_to_eat <= 0 then fail
        else return (eat_cell neighbor s)
      and choose_keep =
        if s.been_size >= island_size then fail
        else find_island neighbor s in
      mplus choose_eat choose_keep) (* Choice point! *)
  s in

let cells_to_eat =
  List.length honey - island_size * num_islands in
find_board (init_state honey cells_to_eat)

```

The `mplus choose_eat choose_keep` creates a choice point: either eat the cell or keep it as part of the island. The monad-plus structure handles backtracking automatically.

8.7 Monad Flavors

Monads “wrap around” a type, but some monads need an additional type parameter. Usually the additional type does not change while within a monad, so we stick to `'a monad` rather than `('s, 'a) monad`.

As monad-plus shows, things get interesting when we add more operations to a basic monad. Here are some common monad flavors:

Monads with access:

```
access : 'a monad -> 'a
```

Example: the lazy monad.

Monad-plus (non-deterministic computation):

```
mzero : 'a monad
mplus : 'a monad -> 'a monad -> 'a monad
```

Monads with state (parameterized by type store):

```
get : store monad
put : store -> unit monad
```

There is a “canonical” state monad. Similar monads include the writer monad (with `get` called `listen` and `put` called `tell`) and the reader monad, without `put`, but with `get` (called `ask`) and:

```
local : (store -> store) -> 'a monad -> 'a monad
```

Exception/error monads (parameterized by type `excn`):

```
throw : excn -> 'a monad
catch : 'a monad -> (excn -> 'a monad) -> 'a monad
```

Continuation monad:

```
callCC : (('a -> 'b monad) -> 'a monad) -> 'a monad
```

We will not cover continuations in detail here.

Probabilistic computation:

```
choose : float -> 'a monad -> 'a monad -> 'a monad
```

satisfying the laws with $a \oplus_p b$ for `choose p a b` and $p \cdot q$ for `p *. q`, where $0 \leq p, q \leq 1$:

$$\begin{aligned} a \oplus_0 b &\approx b \\ a \oplus_p b &\approx b \oplus_{1-p} a \\ a \oplus_p (b \oplus_q c) &\approx (a \oplus_{\frac{p}{p+q-pq}} b) \oplus_{p+q-pq} c \\ a \oplus_p a &\approx a \end{aligned}$$

Parallel computation (monad with access and parallel bind):

```
parallel : 'a monad -> 'b monad -> ('a -> 'b -> 'c monad) -> 'c monad
```

Example: lightweight threads.

8.8 Interlude: The Module System

OCaml's module system provides the infrastructure for defining monads in a reusable way. Here is a brief overview of the key concepts.

Modules collect related type definitions and operations together. Module values are introduced with `struct ... end` (structures), and module types with `sig ... end` (signatures). A structure is a package of definitions; a signature is an interface for packages.

A source file `source.ml` defines a module `Source`. A file `source.mli` defines its type.

In the module level, modules are defined with `module ModuleName = ...` or `module ModuleName : MODULE_TYPE = ...`, and module types with `module type MODULE_TYPE = ...`.

Locally in expressions, modules are defined with `let module M = ... in ...`.

The content of a module is made visible with `open Module`. Module `Pervasives` (now `Stdlib`) is initially visible.

Content of a module is included into another module with `include Module`.

Functors are module functions – functions from modules to modules:

```
module Funct = functor (Arg : sig ... end) -> struct ... end
(* Or equivalently: *)
module Funct (Arg : sig ... end) = struct ... end
```

Functors can return functors, and modules can be parameterized by multiple modules. Functor application always uses parentheses: `Funct (struct ... end)`.

A signature `MODULE_TYPE` with type `t_name = ...` is like `MODULE_TYPE` but with `t_name` made more specific. We can also include signatures with `include MODULE_TYPE`.

Finally, we can pass around modules in normal functions using first-class modules:

```
module type T = sig val g : int -> int end

let f mod_v x =
  let module M = (val mod_v : T) in
  M.g x
(* val f : (module T) -> int -> int = <fun> *)

let test = f (module struct let g i = i*i end : T)
(* val test : int -> int = <fun> *)
```

8.9 The Two Metaphors

Monads can be understood through two complementary metaphors.

Monads as Containers A monad is a **quarantine container**:

- We can put something into the container with `return`
- We can operate on it, but the result needs to stay in the container

```
let lift f m =
  let* x = m in
  return (f x)
(* val lift : ('a -> 'b) -> 'a monad -> 'b monad *)
```

- We can deactivate-unwrap the quarantine container but only when it is in another container so the quarantine is not broken

```
let join m =
  let* x = m in
  x
(* val join : ('a monad) monad -> 'a monad *)
```

The quarantine container for a **monad-plus** is more like other containers: it can be empty, or contain multiple elements.

Monads with access allow us to extract the resulting element from the container; other monads provide a `run` operation that exposes “what really happened behind the quarantine.”

Monads as Computation To compute the result, use `let*` bindings to sequence instructions, naming partial results. The physical metaphor is an **assembly line**:

```
let assemblyLine w =
  let* c = makeChopsticks w in      (* Worker makes chopsticks *)
  let* c' = polishChopsticks c in (* Worker polishes them *)
  let* c'' = wrapChopsticks c' in (* Worker wraps them *)
  return c''                      (* Loader returns the result *)
```

Any expression can be spread over a monad. For lambda-terms:

$\llbracket N \rrbracket = \text{return } N$	(constant)
$\llbracket x \rrbracket = \text{return } x$	(variable)
$\llbracket \lambda x. a \rrbracket = \text{return } (\lambda x. \llbracket a \rrbracket)$	(function)
$\llbracket \text{let } x = a \text{ in } b \rrbracket = \text{bind } \llbracket a \rrbracket (\lambda x. \llbracket b \rrbracket)$	(local definition)
$\llbracket a b \rrbracket = \text{bind } \llbracket a \rrbracket (\lambda v_a. \text{bind } \llbracket b \rrbracket (\lambda v_b. v_a v_b))$	(application)

When an expression is spread over a monad, its computation can be monitored or affected without modifying the expression.

8.10 Monad Classes and Instances

To implement a monad, we need to provide the implementation type, `return`, and `bind` operations.

```
module type MONAD = sig
  type 'a t
  val return : 'a -> 'a t
  val bind : 'a t -> ('a -> 'b t) -> 'b t
end
```

Alternatively, we could start from `return`, `lift`, and `join` operations.

Based on just these two operations, we can define a suite of general-purpose functions:

```
module type MONAD_OPS = sig
  type 'a monad
  include MONAD with type 'a t := 'a monad
  val ( let* ) : 'a monad -> ('a -> 'b monad) -> 'b monad
  val ( let+ ) : 'a monad -> ('a -> 'b) -> 'b monad
  val ( >>= ) : 'a monad -> ('a -> 'b monad) -> 'b monad
```

```

val foldM : ('a -> 'b -> 'a monad) -> 'a -> 'b list -> 'a monad
val whenM : bool -> unit monad -> unit monad
val lift : ('a -> 'b) -> 'a monad -> 'b monad
val (>>|) : 'a monad -> ('a -> 'b) -> 'b monad
val join : 'a monad monad -> 'a monad
val ( >=>) : ('a -> 'b monad) -> ('b -> 'c monad) -> 'a -> 'c monad
end

module MonadOps (M : MONAD) = struct
  open M
  type 'a monad = 'a t
  let run x = x
  let ( let* ) a b = bind a b
  let ( let+ ) a f = bind a (fun x -> return (f x))
  let (>>=) a b = bind a b
  let rec foldM f a = function
    | [] -> return a
    | x::xs ->
      let* a' = f a x in
      foldM f a' xs
  let whenM p s = if p then s else return ()
  let lift f m =
    let* x = m in
    return (f x)
  let (>>|) a b = lift b a
  let join m =
    let* x = m in
    x
  let (>=>) f g = fun x ->
    let* y = f x in
    g y
end

```

We make the monad “safe” by keeping its type abstract, but `run` exposes “what really happened”:

```

module Monad (M : MONAD) : sig
  include MONAD_OPS
  val run : 'a monad -> 'a M.t
end = struct
  include M
  include MonadOps(M)
end

```

Monad-Plus Classes The monad-plus class has many implementations. They need to provide `mzero` and `mplus`:

```

module type MONAD_PLUS = sig
  include MONAD
  val mzero : 'a t
  val mplus : 'a t -> 'a t -> 'a t
end

module type MONAD_PLUS_OPS = sig
  include MONAD_OPS
  val mzero : 'a monad
  val mplus : 'a monad -> 'a monad -> 'a monad
  val fail : 'a monad
  val (++) : 'a monad -> 'a monad -> 'a monad
  val guard : bool -> unit monad
  val msum_map : ('a -> 'b monad) -> 'a list -> 'b monad
end

module MonadPlusOps (M : MONAD_PLUS) = struct
  open M
  include MonadOps(M)
  let fail = mzero
  let (++) a b = mplus a b
  let guard p = if p then return () else fail
  let msum_map f l = List.fold_right
    (fun a acc -> mplus (f a) acc) l mzero
end

module MonadPlus (M : MONAD_PLUS) : sig
  include MONAD_PLUS_OPS
  val run : 'a monad -> 'a M.t
end = struct
  include M
  include MonadPlusOps(M)
end

```

We also need a class for computations with state:

```

module type STATE = sig
  type store
  type 'a t
  val get : store t
  val put : store -> unit t
end

```

8.11 Monad Instances

The Lazy Monad Heavy laziness notation? Try a monad (with access):

```

module LazyM = Monad (struct
  type 'a t = 'a Lazy.t
  let bind a b = lazy (Lazy.force (b (Lazy.force a)))
  let return a = lazy a
end)

let laccess m = Lazy.force (LazyM.run m)

```

The List Monad Our resident list monad (monad-plus):

```

module ListM = MonadPlus (struct
  type 'a t = 'a list
  let bind a b = concat_map b a
  let return a = [a]
  let mzero = []
  let mplus = List.append
end)

```

Backtracking Parameterized by Monad-Plus The Countdown module can be parameterized by any monad-plus:

```

module Countdown (M : MONAD_PLUS_OPS) = struct
  open M (* Open the module to make monad operations visible *)

  let rec insert x = function (* All choice-introducing operations *)
    | [] -> return [x]           (* need to happen in the monad *)
    | y::ys as xs ->
      let* xys = insert x ys in
      return (x::xs) ++ return (y::xys)

  let rec choices = function
    | [] -> return []
    | x::xs ->
      let* cxs = choices xs in
      return cxs ++ insert x cxs          (* Choosing which numbers in what order *)
                                         (* and now whether with or without x *)

  type op = Add | Sub | Mul | Div

  let apply op x y =
    match op with
    | Add -> x + y
    | Sub -> x - y
    | Mul -> x * y
    | Div -> x / y

  let valid op x y =

```

```

match op with
| Add -> x <= y
| Sub -> x > y
| Mul -> x <= y && x <> 1 && y <> 1
| Div -> x mod y = 0 && y <> 1

type expr = Val of int | App of op * expr * expr

let op2str = function
| Add -> "+" | Sub -> "-" | Mul -> "*" | Div -> "/"

let rec expr2str = function (* We will provide solutions as strings *)
| Val n -> string_of_int n
| App (op, l, r) -> "(" ^ expr2str l ^ op2str op ^ expr2str r ^ ")"

let combine (l, x) (r, y) o = (* Try out an operator *)
let* () = guard (valid o x y) in
return (App (o, l, r), apply o x y)

let split l = (* Another choice: which numbers go into which argument *)
let rec aux lhs = function
| [] | [_] -> fail (* Both arguments need numbers *)
| [y; z] -> return (List.rev (y::lhs), [z])
| hd::rhs ->
let lhs = hd::lhs in
return (List.rev lhs, rhs)
++ aux lhs rhs in
aux [] l

let rec results = function (* Build possible expressions once numbers *)
| [] -> fail (* have been picked *)
| [n] ->
let* () = guard (n > 0) in
return (Val n, n)
| ns ->
let* (ls, rs) = split ns in
let* lx = results ls in
let* ly = results rs in (* Collect solutions using each operator *)
msum_map (combine lx ly) [Add; Sub; Mul; Div]

let solutions ns n = (* Solve the problem: *)
let* ns' = choices ns in (* pick numbers and their order, *)
let* (e, m) = results ns' in (* build possible expressions, *)
let* () = guard (m = n) in (* check if the expression gives target value, *)
return (expr2str e) (* "print" the solution *)
end

```

Understanding Laziness Let us measure execution times:

```
let time f =
  let tbeg = Unix.gettimeofday () in
  let res = f () in
  let tend = Unix.gettimeofday () in
  tend -. tbeg, res
```

With the list monad:

```
module ListCountdown = Countdown (ListM)
let test1 () = ListM.run (ListCountdown.solutions [1;3;7;10;25;50] 765)
let t1, sol1 = time test1
(* val t1 : float = 2.28... *)
(* val sol1 : string list = ["/((25-(3+7))*(1+50))"; "/((25-3)-7)*(1+50))"; ...] *)
```

What if we want only one solution? Laziness to the rescue! We define an “odd lazy list”:

```
type 'a llist = LNil | LCons of 'a * 'a llist Lazy.t

let rec ltake n = function
| LCons (a, lazy l) when n > 0 -> a :: (ltake (n-1) l)
| _ -> []

let rec lappend l1 l2 =
  match l1 with
| LNil -> l2
| LCons (hd, tl) ->
  LCons (hd, lazy (lappend (Lazy.force tl) l2))

let rec lconcat_map f = function
| LNil -> LNil
| LCons (a, lazy l) ->
  lappend (f a) (lconcat_map f l)

module LListM = MonadPlus (struct
  type 'a t = 'a llist
  let bind a b = lconcat_map b a
  let return a = LCons (a, lazy LNil)
  let mzero = LNil
  let mplus = lappend
end)
```

Testing shows that the odd lazy list still takes about the same time to even get the lazy list started! The elements are almost already computed once the first one is.

The **option monad** does not help either:

```

module OptionM = MonadPlus (struct
  type 'a t = 'a option
  let bind a b =
    match a with None -> None | Some x -> b x
  let return a = Some a
  let mzero = None
  let mplus a b = match a with None -> b | Some _ -> a
end)

```

This very quickly computes... nothing. The `OptionM` monad (Haskell's `Maybe` monad) is good for computations that might fail, but not for search with multiple solutions.

Our lazy list type is not lazy enough. Whenever we “make” a choice with `a ++ b` or `msum_map`, it computes the first candidate for each choice path immediately.

We need **even lazy lists** (our `llist` above are called “odd lazy lists”):

```

type 'a lazy_list = 'a lazy_list_ Lazy.t
and 'a lazy_list_ = LazNil | LazCons of 'a * 'a lazy_list

let rec laztake n = function
| lazy (LazCons (a, l)) when n > 0 -> a :: (laztake (n-1) l)
| _ -> []

let rec append_aux l1 l2 =
  match l1 with
  | lazy LazNil -> Lazy.force l2
  | lazy (LazCons (hd, tl)) ->
    LazCons (hd, lazy (append_aux tl l2))

let lazappend l1 l2 = lazy (append_aux l1 l2)

let rec concat_map_aux f = function
| lazy LazNil -> LazNil
| lazy (LazCons (a, l)) ->
  append_aux (f a) (lazy (concat_map_aux f l))

let lazconcat_map f l = lazy (concat_map_aux f l)

module LazyListM = MonadPlus (struct
  type 'a t = 'a lazy_list
  let bind a b = lazconcat_map b a
  let return a = lazy (LazCons (a, lazy LazNil))
  let mzero = lazy LazNil
  let mplus = lazappend
end)

```

Now the first solution takes considerably less time than all solutions. The next 9 solutions are almost computed once the first one is. But computing all solutions takes nearly twice as long as without the overhead of lazy computation – the price of laziness.

The Exception Monad Built-in non-functional exceptions in OCaml are more efficient and more flexible. However, monadic exceptions are safer than standard exceptions in situations like multi-threading. The monadic lightweight-thread library Lwt has `throw` (called `fail` there) and `catch` operations in its monad.

```
module ExceptionM (Excn : sig type t end) : sig
  type excn = Excn.t
  type 'a t = OK of 'a | Bad of excn
  include MONAD_OPS
  val run : 'a monad -> 'a t
  val throw : excn -> 'a monad
  val catch : 'a monad -> (excn -> 'a monad) -> 'a monad
end = struct
  type excn = Excn.t
  module M = struct
    type 'a t = OK of 'a | Bad of excn
    let return a = OK a
    let bind m b = match m with
      | OK a -> b a
      | Bad e -> Bad e
    end
    include M
    include MonadOps(M)
    let throw e = Bad e
    let catch m handler = match m with
      | OK _ -> m
      | Bad e -> handler e
  end
```

The State Monad

```
module StateM (Store : sig type t end) : sig
  type store = Store.t
  type 'a t = store -> 'a * store (* Pass the current store value to get the next value *)
  include MONAD_OPS
  include STATE with type 'a t := 'a monad
                and type store := store
  val run : 'a monad -> 'a t
end = struct
  type store = Store.t
```

```

module M = struct
  type 'a t = store -> 'a * store
  let return a = fun s -> a, s      (* Keep the current value unchanged *)
  let bind m b = fun s -> let a, s' = m s in b a s'
end                                (* To bind two steps, pass the value after first step *)
include M
include MonadOps(M)
let get = fun s -> s, s           (* Keep the value unchanged but put it in monad *)
let put s' = fun _ -> (), s'      (* Change the value; a throwaway in monad *)
end

```

The state monad is useful to hide passing-around of a “current” value. Here is an example that renames variables in lambda-terms to eliminate potential name clashes:

```

type term =
| Var of string
| Lam of string * term
| App of term * term

let (!) x = Var x
let (|->) x t = Lam (x, t)
let (@) t1 t2 = App (t1, t2)
let test = "x" |-> ("x" |-> !"y" @ !"x") @ !"x"

module S = StateM (struct type t = int * (string * string) list end)
open S

let rec alpha_conv = function
| Var x as v ->                      (* Function from terms to StateM monad *)
  let* (_, env) = get in                (* Seeing a variable does not change state *)
  let v = try Var (List.assoc x env)    (* but we need its new name *)
  with Not_found -> v in               (* Free variables don't change name *)
  return v
| Lam (x, t) ->                     (* We rename each bound variable *)
  let* (fresh, env) = get in           (* We need a fresh number *)
  let x' = x ^ string_of_int fresh in
  let* () = put (fresh+1, (x, x')::env) in (* Remember new name, update number *)
  let* t' = alpha_conv t in
  let* (fresh', _) = get in           (* We need to restore names, *)
  let* () = put (fresh', env) in      (* but keep the number fresh *)
  return (Lam (x', t'))
| App (t1, t2) ->
  let* t1 = alpha_conv t1 in          (* Passing around of names *)
  let* t2 = alpha_conv t2 in          (* and the currently fresh number *)
  return (App (t1, t2))              (* is done by the monad *)

```

```
(* val test : term = Lam ("x", App (Lam ("x", App (Var "y", Var "x")), Var "x")) *)
(* # StateM.run (alpha_conv test) (5, []);;
 - : term * (int * (string * string) list) =
  (Lam ("x5", App (Lam ("x6", App (Var "y", Var "x6")), Var "x5")), (7, [])) *)
```

Note: This does not make a lambda-term safe for multiple steps of beta-reduction.
Can you find a counter-example?

8.12 Monad Transformers

Sometimes we need merits of multiple monads at the same time, e.g., monads AM and BM . The straightforward idea is to nest one monad within another: either ' $a\ AM.\text{monad } BM.\text{monad}$ ' or ' $a\ BM.\text{monad } AM.\text{monad}$ '. But we want a monad that has operations of both AM and BM .

It turns out that the straightforward approach does not lead to operations with the meaning we want. A **monad transformer** AT takes a monad BM and turns it into a monad $AT(BM)$ which actually wraps around BM on both sides. $AT(BM)$ has operations of both monads.

We will develop a monad transformer StateT which adds state to a monad-plus. The resulting monad has all: return , bind , mzero , mplus , put , get , and their supporting general-purpose functions.

We need monad transformers in OCaml because “monads are contagious”: although we have built-in state and exceptions, we need to use monadic state and exceptions when we are inside a monad. This is the reason Lwt is both a concurrency and an exception monad.

The state monad uses `let x = a in ...` for binding. The transformed monad uses `M.bind` (or `M.let*`) instead:

```
type 'a state = store -> ('a * store)

let return (a : 'a) : 'a state =
  fun s -> (a, s)

let bind (u : 'a state) (f : 'a -> 'b state) : 'b state =
  fun s -> (fun (a, s') -> f a s') (u s)

(* Monad M transformed to add state, in pseudo-code: *)
type 'a stateT(M) = store -> ('a * store) M
(* notice this is not an ('a M) state *)

let return (a : 'a) : 'a stateT(M) =
  fun s -> M.return (a, s)           (* Rather than returning, M.return *)

let bind (u : 'a stateT(M)) (f : 'a -> 'b stateT(M)) : 'b stateT(M) =
  fun s -> M.bind (u s) (fun (a, s') -> f a s') (* Rather than let-binding, M.bind *)
```

State Transformer Implementation

```

module StateT (MP : MONAD_PLUS_OPS) (Store : sig type t end) : sig
  type store = Store.t
  type 'a t = store -> ('a * store) MP.monad
  include MONAD_PLUS_OPS          (* Exporting all monad-plus operations *)
  include STATE with type 'a t := 'a monad
    and type store := store (* and state operations *)
  val run : 'a monad -> 'a t      (* Expose "what happened" -- resulting states *)
  val runT : 'a monad -> store -> 'a MP.monad
end = struct                      (* Run the state transformer -- get resulting values *)
  type store = Store.t
  module M = struct
    type 'a t = store -> ('a * store) MP.monad
    let return a = fun s -> MP.return (a, s)
    let bind m b = fun s ->
      MP.bind (m s) (fun (a, s') -> b a s')
    let mzero = fun _ -> MP.mzero           (* Lift the monad-plus operations *)
    let mplus ma mb = fun s -> MP.mplus (ma s) (mb s)
  end
  include M
  include MonadPlusOps(M)
  let get = fun s -> MP.return (s, s)      (* Instead of just returning, *)
  let put s' = fun _ -> MP.return (((), s')) (* MP.return *)
  let runT m s = MP.lift fst (m s)
end

```

Backtracking with State Using the state transformer with our puzzle solver:

```

module HoneyIslands (M : MONAD_PLUS_OPS) = struct
  type state = {
    been_size : int;
    been_islands : int;
    unvisited : cell list;
    visited : CellSet.t;
    eaten : cell list;
    more_to_eat : int;
  }

  let init_state unvisited more_to_eat = {
    been_size = 0;
    been_islands = 0;
    unvisited;
    visited = CellSet.empty;
    eaten = [];
    more_to_eat;
  }

```

```

}

module BacktrackingM = StateT (M) (struct type t = state end)
open BacktrackingM

let rec visit_cell () = (* State update actions *)
  let* s = get in
  match s.unvisited with
  | [] -> return None
  | c::remaining when CellSet.mem c s.visited ->
    let* () = put {s with unvisited=remaining} in
    visit_cell () (* Throwaway argument because of recursion *)
  | c::remaining ->
    let* () = put {s with
      unvisited=remaining;
      visited = CellSet.add c s.visited} in
    return (Some c) (* This action returns a value *)

let eat_cell c =
  let* s = get in
  let* () = put {s with eaten = c::s.eaten;
    visited = CellSet.add c s.visited;
    more_to_eat = s.more_to_eat - 1} in
  return () (* Remaining state update actions just affect the state *)

let keep_cell c =
  let* s = get in
  let* () = put {s with
    visited = CellSet.add c s.visited;
    been_size = s.been_size + 1} in
  return ()

let fresh_island =
  let* s = get in
  let* () = put {s with been_size = 0;
    been_islands = s.been_islands + 1} in
  return ()

let find_to_eat n island_size num_islands empty_cells =
  let honey = honey_cells n empty_cells in
  let rec find_board () =
    let* cell = visit_cell () in
    match cell with
    | None ->
      let* s = get in
      let* () = guard (s.been_islands = num_islands) in

```

```

        return s.eaten
| Some cell ->
  let* () = fresh_island in
  let* () = find_island cell in
  let* s = get in
  let* () = guard (s.been_size = island_size) in
  find_board ()

and find_island current =
  let* () = keep_cell current in
  neighbors n empty_cells current
|> foldM
  (fun () neighbor ->
    let* s = get in
    whenM (not (CellSet.mem neighbor s.visited))
      (let choose_eat =
        let* () = guard (s.more_to_eat > 0) in
        eat_cell neighbor
      and choose_keep =
        let* () = guard (s.been_size < island_size) in
        find_island neighbor in
        choose_eat ++ choose_keep)) () in

  let cells_to_eat =
    List.length honey - island_size * num_islands in
  init_state honey cells_to_eat
|> runT (find_board ())
end

module HoneyL = HoneyIslands (ListM)
let find_to_eat a b c d =
  ListM.run (HoneyL.find_to_eat a b c d)

```

8.13 Probabilistic Programming

Using a random number generator, we can define procedures that produce various output. This is **not functional** – mathematical functions have a deterministic result for fixed arguments.

Similarly to how we can “simulate” (mutable) variables with state monad and non-determinism with list monad, we can “simulate” random computation with a probability monad.

The probability monad class means much more than having randomized computation. We can ask questions about probabilities of results. Monad instances can make tradeoffs of efficiency vs. accuracy (exact vs. approximate probabilities).

The Probability Monad The essential functions for the probability monad class are `choose` and `distrib`. Remaining functions could be defined in terms of these but are provided by each instance for efficiency.

Inside-monad operations:

- `choose : float -> 'a monad -> 'a monad`: choose p a b represents an event or distribution which is a with probability p and is b with probability $1 - p$.
- `pick : ('a * float) list -> 'a monad`: A result from the provided distribution over values. The argument must be a probability distribution: positive values summing to 1.
- `uniform : 'a list -> 'a monad`: Uniform distribution over given values.
- `flip : float -> bool monad`: Equal to `choose p (return true) (return false)`.
- `coin : bool monad`: Equal to `flip 0.5`.

Outside-monad operations:

- `prob : ('a -> bool) -> 'a monad -> float`: Returns the probability that the predicate holds.
- `distrib : 'a monad -> ('a * float) list`: Returns the distribution of probabilities over the resulting values.
- `access : 'a monad -> 'a`: Samples a random result from the distribution – **non-functional** behavior.

```
module type PROBABILITY = sig
  include MONAD_OPS
  val choose : float -> 'a monad -> 'a monad -> 'a monad
  val pick : ('a * float) list -> 'a monad
  val uniform : 'a list -> 'a monad
  val coin : bool monad
  val flip : float -> bool monad
  val prob : ('a -> bool) -> 'a monad -> float
  val distrib : 'a monad -> ('a * float) list
  val access : 'a monad -> 'a
end
```

Helper functions:

```
let total dist =
  List.fold_left (fun a (_,b) -> a +. b) 0. dist

let merge dist = map_reduce (fun x -> x) (+.) 0. dist (* Merge repeating elements *)

let normalize dist =                                     (* Normalize a measure into a distribution *)
  let tot = total dist in
  if tot = 0. then dist
```

```

else List.map (fun (e,w) -> e, w /. tot) dist

let roulette dist =                                (* Roulette wheel from a distribution/measure *)
  let tot = total dist in
  let rec aux r = function
    | [] -> assert false
    | (e, w)::_ when w <= r -> e
    | (_, w)::tl -> aux (r -. w) tl in
  aux (Random.float tot) dist

```

Exact Distribution Monad

```

module DistribM : PROBABILITY = struct
  module M = struct                                (* Exact probability distribution -- naive implemen
    type 'a t = ('a * float) list
    let bind a b = merge                         (* x w.p. p and then y w.p. q happens = *)
      (List.concat_map (fun (x, p) ->
        List.map (fun (y, q) -> (y, q *. p)) (b x)) a)  (* y results w.p. p*q *)
    let return a = [a, 1.]                         (* Certainly a *)
  end
  include M
  include MonadOps (M)
  let choose p a b =
    List.map (fun (e,w) -> e, p *. w) a @
    List.map (fun (e,w) -> e, (1. -. p) *. w) b
  let pick dist = dist
  let uniform elems = normalize
    (List.map (fun e -> e, 1.) elems)
  let coin = [true, 0.5; false, 0.5]
  let flip p = [true, p; false, 1. -. p]
  let prob p m = m
    |> List.filter (fun (e,_) -> p e)      (* All cases where p holds, *)
    |> List.map snd |> List.fold_left (+.) 0.  (* add up *)
  let distrib m = m
  let access m = roulette m
end

```

Sampling Monad

```

module SamplingM (S : sig val samples : int end) : PROBABILITY = struct
  module M = struct                                (* Parameterized by how many samples *)
    type 'a t = unit -> 'a                      (* used to approximate prob or distrib *)
    let bind a b () = b (a ()) ()                (* Randomized computation -- each call a() *)
    let return a = fun () -> a                  (* is an independent sample. Always a. *)
  end
  include M

```

```

include MonadOps (M)
let choose p a b () =
  if Random.float 1. <= p then a () else b ()
let pick dist = fun () -> roulette dist
let uniform elems =
  let n = List.length elems in
  fun () -> List.nth elems (Random.int n)
let coin = Random.bool
let flip p = choose p (return true) (return false)
let prob p m =
  let count = ref 0 in
  for i = 1 to S.samples do
    if p (m ()) then incr count
  done;
  float_of_int !count /. float_of_int S.samples
let distrib m =
  let dist = ref [] in
  for i = 1 to S.samples do
    dist := (m (), 1.) :: !dist done;
  normalize (merge !dist)
let access m = m ()
end

```

Example: The Monty Hall Problem In search of a new car, the player picks a door, say 1. The game host then opens one of the other doors, say 3, to reveal a goat and offers to let the player pick door 2 instead of door 1.

```

module MontyHall (P : PROBABILITY) = struct
  open P
  type door = A | B | C
  let doors = [A; B; C]

  let monty_win switch =
    let* prize = uniform doors in
    let* chosen = uniform doors in
    let* opened = uniform (list_diff doors [prize; chosen]) in
    let final =
      if switch then List.hd (list_diff doors [opened; chosen])
      else chosen in
    return (final = prize)
end

module MontyExact = MontyHall (DistribM)
module Sampling1000 =
  SamplingM (struct let samples = 1000 end)
module MontySimul = MontyHall (Sampling1000)

```

```

(* DistribM.distrib (MontyExact.monty_win false);;
 - : (bool * float) list = [(true, 0.333...); (false, 0.666...)]

DistribM.distrib (MontyExact.monty_win true);;
 - : (bool * float) list = [(true, 0.666...); (false, 0.333...)] *)

```

The famous result: switching doubles your chances of winning!

Conditional Probabilities Wouldn't it be nice to have a monad-plus rather than just a monad? We could use `guard` for conditional probabilities!

To compute $P(A|B)$: 1. Compute what is needed for both A and B 2. Guard B 3. Return A

For the exact distribution monad, we just need to allow intermediate distributions to be unnormalized (sum to less than 1). For the sampling monad, we use rejection sampling (though `mplus` has no straightforward correct implementation).

```

module type COND_PROBAB = sig
  include PROBABILITY
  include MONAD_PLUS_OPS with type 'a monad := 'a monad
end

module DistribMP : COND_PROBAB = struct
  module MP = struct
    type 'a t = ('a * float) list      (* Measures no longer restricted to *)
    let bind a b = merge              (* probability distributions *)
      (List.concat_map (fun (x, p) ->
        List.map (fun (y, q) -> (y, q *. p)) (b x)) a)
    let return a = [a, 1.]            (* Measure equal 0 everywhere is OK *)
    let mzero = []                   (* Measure equal 0 everywhere is OK *)
    let mplus = List.append
  end
  include MP
  include MonadPlusOps (MP)
  let choose p a b =                (* It isn't a w.p. p & b w.p. (1-p) since a and b *)
    List.map (fun (e,w) -> e, p *. w) a @ (* are not normalized! *)
    List.map (fun (e,w) -> e, (1. -. p) *. w) b
  let pick dist = dist
  let uniform elems = normalize
    (List.map (fun e -> e, 1.) elems)
  let coin = [true, 0.5; false, 0.5]
  let flip p = [true, p; false, 1. -. p]
  let prob p m = normalize m          (* Final normalization step *)
    |> List.filter (fun (e,_) -> p e)
    |> List.map snd |> List.fold_left (+.) 0.

```

```

let distrib m = normalize m
let access m = roulette m
end

module SamplingMP (S : sig val samples : int end) : COND_PROBAB = struct
  exception Rejected           (* For rejecting current sample *)
  module MP = struct           (* Monad operations are exactly as for SamplingM *)
    type 'a t = unit -> 'a
    let bind a b () = b (a ()) ()
    let return a = fun () -> a
    let mzero = fun () -> raise Rejected (* but now we can fail *)
    let mplus a b = fun () ->
      failwith "SamplingMP.mplus not implemented"
  end
  include MP
  include MonadPlusOps (MP)
  let choose p a b () =
    if Random.float 1. <= p then a () else b ()           (* Inside-monad operations don't change *)
  let pick dist = fun () -> roulette dist
  let uniform elems =
    let n = List.length elems in
    fun () -> List.nth elems (Random.int n)
  let coin = Random.bool
  let flip p = choose p (return true) (return false)
  let prob p m =                                         (* Getting out of monad: handle rejected samples *)
    let count = ref 0 and tot = ref 0 in
    while !tot < S.samples do                         (* Count up to the required *)
      try                                                 (* number of samples *)
        if p (m ()) then incr count;                  (* m() can fail *)
        incr tot                                     (* But if we got here it hasn't *)
      with Rejected -> ()                           (* Rejected, keep sampling *)
    done;
    float_of_int !count /. float_of_int S.samples
  let distrib m =
    let dist = ref [] and tot = ref 0 in
    while !tot < S.samples do
      try
        dist := (m (), 1.) :: !dist;
        incr tot
      with Rejected -> ()
    done;
    normalize (merge !dist)
  let rec access m =
    try m () with Rejected -> access m
end

```

Burglary Example: Encoding a Bayes Net Consider a problem with this dependency structure:

- An alarm can be due to either a burglary or an earthquake
- You are on vacation and have asked neighbors John and Mary to call if the alarm rings
- Mary only calls when she is really sure about the alarm, but John has better hearing
- Earthquakes are twice as probable as burglaries
- The alarm has about 30% chance of going off during an earthquake
- You can check on the radio if there was an earthquake, but you might miss the news

Probability tables: - $P(\text{Burglary}) = 0.001$ - $P(\text{Earthquake}) = 0.002$ - $P(\text{Alarm}|\text{B}, \text{E})$ varies (0.001 for FF, 0.29 for FT, 0.94 for TF, 0.95 for TT) - $P(\text{John calls}|\text{Alarm})$ is 0.9 if alarm, 0.05 otherwise - $P(\text{Mary calls}|\text{Alarm})$ is 0.7 if alarm, 0.01 otherwise

```
module Burglary (P : COND_PROBAB) = struct
  open P
  type what_happened =
    | Safe | Burgl | Earthq | Burgl_n_earthq

  let check ~john_called ~mary_called ~radio =
    let* earthquake = flip 0.002 in
    let* () = guard (radio = None || radio = Some earthquake) in
    let* burglary = flip 0.001 in
    let alarm_p =
      match burglary, earthquake with
      | false, false -> 0.001
      | false, true -> 0.29
      | true, false -> 0.94
      | true, true -> 0.95 in
    let* alarm = flip alarm_p in
    let john_p = if alarm then 0.9 else 0.05 in
    let* john_calls = flip john_p in
    let* () = guard (john_calls = john_called) in
    let mary_p = if alarm then 0.7 else 0.01 in
    let* mary_calls = flip mary_p in
    let* () = guard (mary_calls = mary_called) in
    match burglary, earthquake with
    | false, false -> return Safe
    | true, false -> return Burgl
    | false, true -> return Earthq
    | true, true -> return Burgl_n_earthq
  end
```

```

module BurglaryExact = Burglary (DistribMP)
module Sampling2000 =
  SamplingMP (struct let samples = 2000 end)
module BurglarySimul = Burglary (Sampling2000)

(* DistribMP.distrib
   (BurglaryExact.check ~john_called:true ~mary_called:true ~radio:None);;
- : (BurglaryExact.what_happened * float) list =
 [(Burgl_n_earthq, 0.000574...); (Earthq, 0.175...);
 (Burgl, 0.283...); (Safe, 0.540...)])

```

8.14 Lightweight Cooperative Threads

The `bind` operation is inherently sequential: `bind a (fun x -> b)` computes `a`, and resumes computing `b` only once the result `x` is known.

For concurrency, we need to “suppress” this sequentiality. We introduce:

```
parallel : 'a monad -> 'b monad -> ('a -> 'b -> 'c monad) -> 'c monad
```

where `parallel a b (fun x y -> c)` does not wait for `a` to be computed before it can start computing `b`.

If the monad starts computing right away (as in the Lwt library), `parallel ea eb c` is equivalent to:

```

let a = ea in
let b = eb in
let* x = a in
let* y = b in
c x y

```

Fine-Grained vs. Coarse-Grained Concurrency Under **fine-grained** concurrency, every `bind` is suspended and computation moves to other threads. It comes back to complete the `bind` before running threads created since the `bind` was suspended.

Under **coarse-grained** concurrency, computation is only suspended when requested via a `suspend` (often called `yield`) operation. Library operations that need to wait for an event or completion of I/O should call `suspend` internally.

Thread Monad Signatures

```

module type THREADS = sig
  include MONAD
  val parallel :
    'a t -> 'b t -> ('a -> 'b -> 'c t) -> 'c t
end

```

```

module type THREAD_OPS = sig
  include MONAD_OPS
  include THREADS with type 'a t := 'a monad
  val parallel_map :
    'a list -> ('a -> 'b monad) -> 'b list monad
  val (>||=) :
    'a monad -> 'b monad -> ('a -> 'b -> 'c monad) -> 'c monad
  val (>||) :
    'a monad -> 'b monad -> (unit -> 'c monad) -> 'c monad
end

module type THREADSYS = sig
  include THREADS
  val access : 'a t -> 'a
  val kill_threads : unit -> unit
end

module ThreadOps (M : THREADS) = struct
  open M
  include MonadOps (M)
  let parallel_map l f =
    List.fold_right (fun a bs ->
      parallel (f a) bs
      (fun a bs -> return (a::bs))) l (return [])
  let (>||=) = parallel
  let (>||) a b c = parallel a b (fun _ _ -> c ())
end

module Threads (M : THREADSYS) : sig
  include THREAD_OPS
  val access : 'a monad -> 'a
  val kill_threads : unit -> unit
end = struct
  include M
  include ThreadOps(M)
end

```

Cooperative Thread Implementation

```

module Cooperative = Threads(struct
  type 'a state =
    | Return of 'a                                (* The thread has returned *)
    | Sleep of ('a -> unit) list                 (* When thread returns, wake up waiters *)
    | Link of 'a t                                 (* A link to the actual thread *)
  and 'a t = {mutable state : 'a state}          (* State of the thread can change *)
                                                (* -- it can return, or more waiters can be added *)

```

```

let rec find t =                                     (* Union-find style link chasing *)
  match t.state with
  | Link t -> find t
  | _ -> t

let jobs = Queue.create ()                         (* Work queue -- will store unit -> unit procedures *

let wakeup m a =                                 (* Thread m has actually finished -- *)
  let m = find m in                             (* updating its state *)
  match m.state with
  | Return _ -> assert false
  | Sleep waiters ->
    m.state <- Return a;                      (* Set the state, and only then *)
    List.iter ((|>) a) waiters             (* wake up the waiters *)
  | Link _ -> assert false

let return a = {state = Return a}

let connect t t' =                                (* t was a placeholder for t' *)
  let t' = find t' in
  match t'.state with
  | Sleep waiters' ->
    let t = find t in
    (match t.state with
    | Sleep waiters ->                      (* If both sleep, collect their waiters *)
      t.state <- Sleep (waiters' @ waiters);
      t'.state <- Link t                   (* and link one to the other *)
    | _ -> assert false)
  | Return x -> wakeup t x                  (* If t' returned, wake up the placeholder *)
  | Link _ -> assert false

let rec bind a b =
  let a = find a in
  let m = {state = Sleep []} in           (* The resulting monad *)
  (match a.state with
  | Return x ->                         (* If a returned, we suspend further work *)
    let job () = connect m (b x) in   (* In exercise 11, this should *)
    Queue.push job jobs                (* only happen after suspend) *)
  | Sleep waiters ->                    (* If a sleeps, we wait for it to return *)
    let job x = connect m (b x) in
    a.state <- Sleep (job::waiters)
  | Link _ -> assert false);
  m

let parallel a b c =                            (* Since in our implementation *)
  bind a (fun x ->                            (* the threads run as soon as they are created, *)

```

```

bind b (fun y ->
          (* parallel is redundant *)
          c x y))

let rec access m =
  let m = find m in
  match m.state with
  | Return x -> x
    (* No further work *)
  | Sleep _ ->
    (try Queue.pop jobs ()
       with Queue.Empty ->
         failwith "access: result not available");
     access m
  | Link _ -> assert false

let kill_threads () = Queue.clear jobs (* Remove pending work *)
end)

```

Testing the Thread Implementation

```

module TTest (T : THREAD_OPS) = struct
  open T
  let rec loop s n =
    let* () = return (Printf.printf "-- %s(%d)\n%" s n) in
    if n > 0 then loop s (n-1)           (* We cannot use whenM because *)
    else return ()                         (* the thread would be created regardless of condition
end

module TT = TTest (Cooperative)

let test =
  Cooperative.kill_threads ();           (* Clean-up after previous tests *)
  let thread1 = TT.loop "A" 5 in
  let thread2 = TT.loop "B" 4 in
  Cooperative.access thread1;           (* We ensure threads finish computing *)
  Cooperative.access thread2           (* before we proceed *)

(* Output:
   -- A(5)
   -- B(4)
   -- A(4)
   -- B(3)
   -- A(3)
   -- B(2)
   -- A(2)
   -- B(1)
   -- A(1)

```

```
-- B(0)
-- A(0)
val test : unit = () *
```

The output shows that the threads interleave their execution, with each `bind` causing a context switch.

8.15 Exercises

Exercise 1.

(Puzzle via Oleg Kiselyov)

“U2” has a concert that starts in 17 minutes and they must all cross a bridge to get there. All four men begin on the same side of the bridge. It is night. There is one flashlight. A maximum of two people can cross at one time. Any party who crosses, either 1 or 2 people, must have the flashlight with them. The flashlight must be walked back and forth, it cannot be thrown, etc. Each band member walks at a different speed. A pair must walk together at the rate of the slower man’s pace:

- Bono: 1 minute to cross
- Edge: 2 minutes to cross
- Adam: 5 minutes to cross
- Larry: 10 minutes to cross

For example: if Bono and Larry walk across first, 10 minutes have elapsed when they get to the other side of the bridge. If Larry then returns with the flashlight, a total of 20 minutes have passed and you have failed the mission.

Find all answers to the puzzle using a list comprehension. The comprehension will be a bit long but recursion is not needed.

Exercise 2. Assume `concat_map` as defined in lecture 6 and the binding operators defined above. What will the following expressions return? Why?

1. `let* _ = return 5 in return 7`
2. `let guard p = if p then [()] else [] in let* () = guard false in return 7`
3. `let* _ = return 5 in let* () = guard false in return 7`

Exercise 3. Define `bind` in terms of `lift` and `join`.

Exercise 4. Define a monad-plus implementation based on binary trees, with constant-time `mzero` and `mpplus`. Starter code:

```
type 'a tree = Empty | Leaf of 'a | T of 'a tree * 'a tree

module TreeM = MonadPlus (struct
  type 'a t = 'a tree
  let bind a b = (* TODO *)
  let return a = (* TODO *)
  let mzero = (* TODO *)
```

```

let mplus a b = (* TODO *)
end)

```

Exercise 5. Show the monad-plus laws for one of: 1. TreeM from your solution of exercise 4 2. ListM from lecture

Exercise 6. Why is the following monad-plus not lazy enough?

```

let rec badappend l1 l2 =
  match l1 with lazy LazNil -> l2
  | lazy (LazCons (hd, tl)) ->
    lazy (LazCons (hd, badappend tl l2))

```

```

let rec badconcatmap f = function
  | lazy LazNil -> lazy LazNil
  | lazy (LazCons (a, l)) ->
    badappend (f a) (badconcatmap f l)

```

```

module BadyListM = MonadPlus (struct
  type 'a t = 'a lazylist
  let bind a b = badconcatmap b a
  let return a = lazy (LazCons (a, lazy LazNil))
  let mzero = lazy LazNil
  let mplus = badappend
end)

```

Exercise 7. Convert a “rectangular” list of lists of strings, representing a matrix with inner lists being rows, into a string, where elements are column-aligned. (Exercise not related to monads.)

Exercise 8. Recall the enriched monad signature with $('s, 'a) t$ type. Design the signatures for the exception monad operations to provide more flexibility than our exception monad. Does the implementation need to change?

Exercise 9. Implement the following constructs for *all* monads:

1. for...to...
2. for...downto...
3. while...do...
4. do...while...
5. repeat...until...

Explain how, when your implementation is instantiated with the StateM monad, we get the solution to exercise 2 from lecture 4.

Exercise 10. A canonical example of a probabilistic model is that of a lawn whose grass may be wet because it rained, because the sprinkler was on, or for some other reason. The probability tables are:

$$\begin{aligned}
P(\text{cloudy}) &= 0.5 \\
P(\text{rain}|\text{cloudy}) &= 0.8 \\
P(\text{rain}|\neg\text{cloudy}) &= 0.2 \\
P(\text{sprinkler}|\text{cloudy}) &= 0.1 \\
P(\text{sprinkler}|\neg\text{cloudy}) &= 0.5 \\
P(\text{wet_roof}|\neg\text{rain}) &= 0 \\
P(\text{wet_roof}|\text{rain}) &= 0.7 \\
P(\text{wet_grass}|\text{rain} \wedge \neg\text{sprinkler}) &= 0.9 \\
P(\text{wet_grass}|\text{sprinkler} \wedge \neg\text{rain}) &= 0.9
\end{aligned}$$

We observe whether the grass is wet and whether the roof is wet. What is the probability that it rained?

Exercise 11. Implement the coarse-grained concurrency model:

- Modify `bind` to compute the resulting monad straight away if the input monad has returned.
- Introduce `suspend` to do what in the fine-grained model was the effect of `bind (return a) b`, i.e., suspend the work although it could already be started.
- One possibility is to introduce `suspend` of type `unit monad`, introduce a “dummy” monadic value `Suspend` (besides `Return` and `Sleep`), and define `bind suspend b` to do what `bind (return ()) b` would formerly do.

Chapter 9: Compiler

Compilation, Runtime, Optimization, and Parsing

This chapter explores the practical aspects of OCaml development, from compilation and runtime systems to performance optimization and parsing techniques. We cover the OCaml toolchain, garbage collection, stack frames and closures, profiling strategies, and both lexing with `ocamllex` and parsing with Menhir.

References: - Andrew W. Appel, *Modern Compiler Implementation in ML* - E. Chailloux, P. Manoury, B. Pagano, *Developing Applications with OCaml* - Jon D. Harrop, *OCaml for Scientists* - Francois Pottier, Yann Regis-Gianas, *Menhir Reference Manual*

9.1 OCaml Compilers

OCaml provides two primary compilers: the bytecode compiler `ocamlc` and the native code compiler `ocamlopt`. Natively compiled code typically runs about 10 times faster than bytecode, though this varies by program.

OCaml also includes an interactive shell called the *toplevel* (or *REPL* in other languages): `ocaml`, which is based on the bytecode compiler. There is also

a native-code-based toplevel `ocamlnat`, though it is not currently part of the binary distribution.

Third-party compilers extend OCaml's reach to other platforms. Most notably, `js_of_ocaml` translates OCaml bytecode into JavaScript source. On modern JS virtual machines like V8, the resulting code can be 2-3 times faster than code running on the OCaml virtual machine (though it can also be slower in some cases).

Compilation Stages The compilation process proceeds through several stages:

Stage	Input	Output
Preprocessing	Source program	Transformed source or AST
Compiling	Source/AST	Assembly program
Assembling	Assembly program	Machine instructions
Linking	Machine instructions	Executable code

OCaml Programs and Tools

Program	Description
<code>ocaml</code>	Toplevel loop (REPL)
<code>ocamlrun</code>	Bytecode interpreter (VM)
<code>camlp4</code>	Preprocessor (syntax extensions)
<code>ocamlc</code>	Bytecode compiler
<code>ocamlopt</code>	Native code compiler
<code>ocamlmktop</code>	New toplevel constructor
<code>ocamldep</code>	Dependencies between modules
<code>ocamlbuild</code>	Building projects tool
<code>ocamlbrowser</code>	Graphical browsing of sources

File Extensions

Extension	Description
<code>.ml</code>	OCaml source file
<code>.mli</code>	OCaml interface source file
<code>.cmi</code>	Compiled interface
<code>.cmo</code>	Bytecode-compiled file
<code>.cmx</code>	Native-code-compiled file
<code>.cma</code>	Bytecode-compiled library (several source files)
<code>.cmxa</code>	Native-code-compiled library
<code>.cmt/.cmti/.annot</code>	Type information for editors
<code>.c</code>	C source file

Extension	Description
.o	C native-code-compiled file
.a	C native-code-compiled library

Compiler Options Both compilers share common command-line options:

Option	Description
-a	Construct a runtime library
-c	Compile without linking
-o name	Specify the name of the executable
-linkall	Link with all libraries used
-i	Display all compiled global declarations
-pp command	Use command as preprocessor
-unsafe	Turn off index checking for arrays
-v	Display the version of the compiler
-w list	Choose warning levels
-impl file	Indicate that file is a Caml source (.ml)
-intf file	Indicate that file is a Caml interface (.mli)
-I directory	Add directory to the search path; prefix + for relative
-g	Generate debugging information

Warning levels control which messages the compiler produces:

Level	Description
A/a	Enable/disable all messages
F/f	Partial application in a sequence
P/p	Incomplete pattern matching
U/u	Missing cases in pattern matching
X/x	Enable/disable messages for hidden objects
M/m, V/v	Object-oriented related warnings

The native compiler has additional options:

Option	Description
-compact	Optimize the produced code for space
-S	Keep the assembly code in a file
-inline level	Set the aggressiveness of inlining

Runtime Configuration The OCamlrunparam environment variable controls runtime behavior:

Parameter	Description
b	Print detailed stack backtrace of runtime exceptions
s/h/i	Size of the minor heap/major heap/size increment
o/o	Major GC speed setting / heap compaction trigger setting

Typical usage for running a program `prog` with backtraces:

```
export OCAMLRUNPARAM='b'; ./prog
```

To have stack backtraces, compile with the `-g` option.

Toplevel Loop Directives The interactive toplevel supports several directives:

Directive	Description
<code>#quit;;</code>	Exit the toplevel
<code>#directory "dir";;</code>	Add <code>dir</code> to the search path; + for relative
<code>#cd "dir-name";;</code>	Change directory
<code>#load "file-name";;</code>	Load a bytecode .cmo/.cma file
<code>#load_rec "file-name";;</code>	Load the files <code>file-name</code> depends on too
<code>#use "file-name";;</code>	Read, compile and execute source phrases
<code>#install_printer pr_nm;;</code>	Register <code>pr_nm</code> to print values of a type
<code>#print_depth num;;</code>	How many nestings to print
<code>#print_length num;;</code>	How many nodes to print (the rest shows as ...)
<code>#trace func;;/#untrace</code>	Trace calls to <code>func</code> / stop tracing

9.1.1 Compiling Multiple-File Projects Traditionally, the file containing a module has a lowercase name, although the module name is always uppercase. Some developers prefer uppercase file names to reflect module names (e.g., `MyModule.ml` rather than `myModule.ml` for a module `MyModule`).

Consider a project with a main module `main.ml` and helper modules `sub1.ml` and `sub2.ml` with corresponding interfaces.

Native compilation by hand:

```
$ ocamlopt sub1.mli
$ ocamlopt sub2.mli
$ ocamlopt -c sub1.ml
$ ocamlopt -c sub2.ml
$ ocamlopt -c main.ml
```

```
$ ocamlopt unix.cmxa sub1.cmx sub2.cmx main.cmx -o prog
$ ./prog
```

Native compilation using make:

```
PROG := prog
LIBS := unix
SOURCES := sub1.ml sub2.ml main.ml
INTERFACES := $(wildcard *.mli)
OBJS := $(patsubst %.ml,%.cmx,$(SOURCES))
LIBS := $(patsubst %,%.cmxa,$(LIBS))

$(PROG): $(OBJS)
    ocamlopt -o $@ $(LIBS) $(OBJS)

clean:
    rm -rf $(PROG) *.* *.cmx *.cmi *~

%.cmx: %.ml
    ocamlopt -c $*.ml

%.cmi: %.mli
    ocamlopt -c $*.mli

depend: $(SOURCES) $(INTERFACES)
    ocamldep -native $(SOURCES) $(INTERFACES)
```

First use: touch .depend; make depend; make. Later, just make, and run make depend after creating new source files.

Using ocamlbuild: - Files with compiled code are created in the _build directory - Command: ocamlbuild -libs unix main.native - The resulting program is called main.native (in directory _build, with a link in the project directory) - Multiple libraries: ocamlbuild -libs nums,unix,graphics main.native - Passing compiler parameters: ocamlbuild -cflags -I,+lablgtk,-rectypes hello.native - Adding -- at the end compiles and runs the program: ocamlbuild -libs unix main.native --

9.1.2 Editors and IDEs

Several editors provide OCaml support:

Emacs: - ocaml-mode from the standard distribution - Alternative: tuareg-mode (<https://forge.ocamlcore.org/projects/tuareg/>) - camldebug integration with debugger - Type feedback with C-c C-t key shortcut (needs .annot files)

Vim: - OMlet plugin - Type lookup with ocaml-annot or other plugins

Eclipse: - OCaml Development Tools (<http://ocamldt.free.fr/>)

TypeRex (<http://www.typerex.org/>): - Auto-completion of identifiers (experimental)
- Browsing of identifiers: show type and comment, go to definition
- Local and whole-program refactoring: renaming identifiers and compilation units, open elimination

Indentation tool ocp-indent (<https://github.com/OCamlPro/ocp-indent>): - Installation instructions for Emacs and Vim - Can be used with other editors

Dedicated editors: - OCamlEditor - `ocamlbrowser` inspects libraries and programs (browsing contents of modules, search by name and by type, basic editing with syntax highlighting)

9.2 Imperative Features in OCaml

OCaml is **not** a *purely functional* language. It has built-in imperative features:

Mutable arrays:

```
let a = Array.make 5 0 in
a.(3) <- 7; a.(2), a.(3)
```

Hashtables in the standard distribution are based on arrays:

```
let h = Hashtbl.create 11 in (* Takes initial size of the array *)
Hashtbl.add h "Alpha" 5; Hashtbl.find h "Alpha"
```

Mutable strings (for historical reasons):

```
let a = Bytes.make 4 'a' in
Bytes.set a 2 'b'; Bytes.get a 2, Bytes.get a 3
```

Extensible mutable strings are provided by `Buffer.t` in the standard distribution.

Loops: - `for i = a to/downto b do body done` - `while condition do body done`

Mutable record fields:

```
type 'a ref = { mutable contents : 'a } (* Single, mutable field *)
```

A record can have both mutable and immutable fields.

- Modifying the field: `record.field <- new_value`
- The `ref` type has operations:

```
let (:=) r v = r.contents <- v
let (!) r = r.contents
```

Exceptions are defined by `exception`, raised by `raise`, and caught by `try-with` clauses. An exception is a variant of type `exception`, which is the only open algebraic datatype – new variants can be added to it.

Input-output functions have no “type safeguards” (no *IO monad*).

Using **global** state (e.g., reference cells) makes code *non re-entrant*: you must finish one task before starting another – any form of concurrency is excluded.

9.2.1 Parsing Command-Line Arguments To go beyond the `Sys.argv` array, see the `Arg` module (<http://caml.inria.fr/pub/docs/manual-ocaml/libref/Arg.html>).

Example: configuring a Mine Sweeper game:

```
type config = {
    nbcols : int; nbrows : int; nbmines : int
}

let default_config = { nbcols=10; nbrows=10; nbmines=15 }

let set_nbcols cf n = cf := {!cf with nbcols = n}
let set_nbrows cf n = cf := {!cf with nbrows = n}
let set_nbmines cf n = cf := {!cf with nbmines = n}

let read_args() =
    let cf = ref default_config in (* State of configuration *)
    let speclist = (* will be updated by given functions *)
        [("-col", Arg.Int (set_nbcols cf), "number of columns");
         ("-lin", Arg.Int (set_nbrows cf), "number of lines");
         ("-min", Arg.Int (set_nbmines cf), "number of mines")] in
    let usage_msg =
        "usage : minesweep [-col n] [-lin n] [-min n]" in
    Arg.parse speclist (fun s -> () ) usage_msg; !cf
```

9.3 OCaml Garbage Collection

9.3.1 Representation of Values OCaml uses a uniform value representation to support garbage collection:

- **Pointers** always end with 00 in binary (addresses are in number of bytes)
- **Integers** are represented by shifting them 1 bit, setting the last bit to 1
- **Constant constructors** (variants without parameters) like `None`, `[]`, and `()`, and other integer-like types (`char`, `bool`) are represented the same way as integers
- **Pointers** are always to OCaml *blocks*. Variants with parameters, strings, and OCaml arrays are stored as blocks
- A **block** starts with a header, followed by an array of values of size 1 word (either integer-like, or pointers)
- The **header** stores the size of the block, the 2-bit color used for garbage collection, and an 8-bit **tag** (which variant it is)
 - Therefore there can be at most about 240 variants with parameters in a variant type (some tag numbers are reserved)

– *Polymorphic variants* are a different story

9.3.2 Generational Garbage Collection OCaml has two heaps to store blocks: a small, continuous *minor heap* and a growing-as-necessary *major heap*.

Allocation simply moves the minor heap pointer (aka. the *young pointer*) and returns the pointed address. Allocation of very large blocks uses the major heap instead.

When the minor heap runs out of space, it triggers the *minor (garbage) collection*, which uses the *Stop and Copy* algorithm. Together with the minor collection, a slice of *major (garbage) collection* is performed to clean up the major heap a bit.

The major heap is not cleaned all at once because it might stop the main program for too long. Major collection uses the *Mark and Sweep* algorithm.

This generational approach works well when most minor heap blocks are already not needed when collection starts – garbage does **not** slow down collection.

9.3.3 Stop and Copy GC Minor collection starts from a set of *roots* – young blocks that definitely are not garbage.

Besides the root set, OCaml also maintains the *remembered set* of minor heap blocks pointed at from the major heap. Most mutations must check whether they assign a minor heap block to a major heap block field. This is called the *write barrier*.

Immutable blocks cannot contain pointers from major to minor heap (unless they are `lazy` blocks).

Collection follows pointers in the root set and remembered set to find other used blocks. Every found block is copied to the major heap. At the end of collection, the young pointer is reset so that the minor heap is empty again.

9.3.4 Mark and Sweep GC Major collection starts from a separate root set – old blocks that definitely are not garbage.

Major garbage collection consists of a *mark* phase which colors blocks that are still in use and a *sweep* phase that searches for stretches of unused memory. Slices of the mark phase are performed by/after each minor collection. Unused memory is stored in a *free list*.

The “proper” major collection is started when a minor collection consumes the remaining free list. The mark phase is finished and sweep phase performed.

The colors used are: - **gray**: marked cells whose descendants are not yet marked - **black**: marked cells whose descendants are also marked - **hatched**: free list element - **white**: elements previously being in use

Example of garbage collection in action:

```

# let u = let l = ['c'; 'a'; 'm'] in List.tl l;;
val u : char list = ['a'; 'm']

# let v = let r = ( ['z'] , u ) in match r with p -> (fst p) @ (snd p);;
val v : char list = ['z'; 'a'; 'm']

```

During the marking phase, reachable cells are colored gray, then black as their descendants are processed. During the sweep phase, white (unreachable) cells are added to the free list.

9.4 Stack Frames and Closures

The nesting of procedure calls is reflected in the *stack* of procedure data. The stretch of stack dedicated to a single function is a *stack frame* (aka. *activation record*).

The *stack pointer* is where we create new frames, stored in a special register. The *frame pointer* allows referring to function data by offset – data known early in compilation is close to the frame pointer.

Local variables are stored in the stack frame or in registers. Some registers need to be saved prior to function call (*caller-save*) or at entry to a function (*callee-save*). OCaml avoids callee-save registers.

Up to 4-6 arguments can be passed in registers, with remaining ones on the stack. Note that the *x86* architecture has a small number of registers.

Using registers, tail call optimization, and function inlining can eliminate the use of stack entirely. The OCaml compiler can also use stack more efficiently than by creating full stack frames.

A typical stack frame contains (from higher to lower addresses): - Incoming arguments (argument n ... argument 1, static link) - Local variables - Return address - Temporaries - Saved registers - Outgoing arguments (argument m ... argument 1, static link)

Static links point to stack frames of parent functions, so we can access stack-based data (e.g., arguments of a main function from inside an `aux` helper function).

A *closure* represents a function: it is a block that contains the address of the function (either another closure or a machine-code pointer) and a way to access non-local variables of the function. For partially applied functions, it contains the values of arguments and the address of the original function.

Escaping variables are the variables of a function `f` (arguments and local definitions) which are accessed from a nested function that is part of the returned value of `f` (or assigned to a mutable field). Escaping variables must be either part of the closures representing the nested functions, or of a closure representing the function `f` – in the latter case, the nested functions must also be represented by closures that have a link to the closure of `f`.

9.4.1 Tail Recursion A function call `f x` within the body of another function `g` is in *tail position* if, roughly, “calling `f` is the last thing that `g` will do before returning.”

Important: A call inside a `try ... with` clause is **not** in tail position! For efficient exceptions, OCaml stores *traps* for `try-with` on the stack with the topmost trap in a register; after `raise`, it unwinds directly to the trap.

The steps for a tail call are: 1. Move actual parameters into argument registers (if they aren’t already there) 2. Restore callee-save registers (if needed) 3. Pop the stack frame of the calling function (if it has one) 4. Jump to the callee

Bytecode always throws a `Stack_overflow` exception on too deep recursion; native code will sometimes cause a *segmentation fault*!

Note: `List.map` from the standard distribution is **not** tail-recursive.

9.4.2 Generated Assembly For examples of generated assembly code and performance analysis, see http://ocaml.org/tutorials/performance_and_profiling.html

9.5 Profiling and Optimization

Steps for optimizing a program:

1. **Profile** the program to find bottlenecks: where the time is spent
2. If possible, modify the algorithm used by the bottleneck to an algorithm with better asymptotic complexity
3. If possible, modify the bottleneck algorithm to access data less randomly, to increase *cache locality*
 - Additionally, *realtime* systems may require avoiding use of huge arrays, traversed by the garbage collector in one go
4. Experiment with various implementations of data structures used (related to step 3)
5. Avoid *boxing* and polymorphic functions, especially for numerical processing (OCaml specific)
6. *Deforestation*
7. *Defunctionalization*

9.5.1 Profiling We cover native code profiling because it is more useful. It relies on the Unix profiling program `gprof`.

First, compile the sources in profiling mode: `ocamlopt -p ...`

Or using `ocamlbuild` when the program source is in `prog.ml`:

```
ocamlbuild prog.p.native --
```

The execution of program `./prog` produces a file `gmon.out`.

We call `gprof prog > profile.txt` (or `gprof prog.p.native > profile.txt` when using `ocamlbuild`). This redirects profiling analysis to `profile.txt`.

The result `profile.txt` has three parts: 1. List of functions in the program in descending order of time spent within the body of the function, excluding time spent in other functions 2. A hierarchical representation of the time taken by each function, and the total time spent in it, including time spent in functions it called 3. A bibliography of function references

It contains C/assembly function names like `camlList__assoc_1169`: - Prefix `caml` means the function comes from OCaml source - `List__` means it belongs to a `List` module - `assoc` is the name of the function in source - Postfix `_1169` is used to avoid name clashes, as in OCaml different functions often have the same names

Example: computing a words histogram for a large file (`Optim0.ml`):

```
let read_words file = (* Imperative programming example *)
  let input = open_in file in
  let words = ref [] and more = ref true in
  try (* Lecture 6 read_lines function would stack-overflow *)
    while !more do (* because of the try-with clause *)
      Scanf.fscanf input "%[^a-zA-Z0-9]%[a-zA-Z0-9]"
      (fun _ x -> words := x :: !words; more := x <> "")
    done;
    List.rev (List.tl !words)
  with End_of_file -> List.rev !words

let empty () = []

let increment h w = (* Inefficient map update *)
  try
    let c = List.assoc w h in
    (w, c+1) :: List.remove_assoc w h
  with Not_found -> (w, 1)::h

let iterate f h =
  List.iter (fun (k,v)->f k v) h

let histogram words =
  List.fold_left increment (empty ()) words
```

The profiling analysis first part begins with:

%	cumulative	self	self	total		
time	seconds	seconds	calls	s/call	s/call	name
37.88	8.54	8.54	306656698	0.00	0.00	compare_val
19.97	13.04	4.50	273169	0.00	0.00	camlList__assoc_1169
9.17	15.10	2.07	633527269	0.00	0.00	caml_page_table_lookup
8.72	17.07	1.97	260756	0.00	0.00	camlList__remove_assoc_1189
7.10	18.67	1.60	612779467	0.00	0.00	caml_string_length

4.97	19.79	1.12	306656692	0.00	0.00	caml_compare
2.84	20.43	0.64				caml_c_call
1.53	20.77	0.35	14417	0.00	0.00	caml_page_table_modify
1.07	21.01	0.24	1115	0.00	0.00	sweep_slice
0.89	21.21	0.20	484	0.00	0.00	mark_slice

List.assoc and List.remove_assoc high in the ranking suggests that increment could be the bottleneck. They both use comparison, which could explain why compare_val consumes the most time.

The second part shows data about the increment function. Each block describes the function whose line starts with an index in brackets, with callers above and callees below:

index	% time	self	children	called	name
[8]	28.7	0.00	6.47	273169/273169	camlList_fold_left_1078 [7]
		0.00	6.47	273169	camlOptim0_increment_1038 [8]
		4.50	0.00	273169/273169	camlList_assoc_1169 [9]
		1.97	0.00	260756/260756	camlList_remove_assoc_1189 [11]

As expected, increment is only called by List.fold_left. But it seems to account for only 29% of time because compare is not analyzed correctly, thus not included in time for increment.

9.5.2 Algorithmic Optimizations All times measured with profiling turned on:

Optim0.ml: asymptotic time complexity $\mathcal{O}(n^2)$, time: 22.53s - Garbage collection takes 6% of time - So little because data access wastes a lot of time

Optimize the data structure, keep the algorithm:

```
let empty () = Hashtbl.create 511

let increment h w =
  try
    let c = Hashtbl.find h w in
    Hashtbl.replace h w (c+1); h
  with Not_found -> Hashtbl.add h w 1; h
```

Optim1.ml: asymptotic time complexity $\mathcal{O}(n)$, time: 0.63s - Garbage collection takes 17% of time

Optimize the algorithm, keep the data structure:

```
let histogram words =
  let words = List.sort String.compare words in
```

```

let k,c,h = List.fold_left
  (fun (k,c,h) w ->
    if k = w then k, c+1, h else w, 1, ((k,c)::h))
  ("", 0, [])
words in
(k,c)::h

```

Optim2.ml: asymptotic time complexity $\mathcal{O}(n \log n)$, time: 1s - Garbage collection takes 40% of time

Optimizing for cache efficiency is more advanced; we will not attempt it here.

With algorithmic optimizations we should be concerned with **asymptotic complexity** in terms of the $\mathcal{O}(\cdot)$ notation, but we will not pursue complexity analysis in the remainder of the chapter.

9.5.3 Low-Level Optimizations The optimizations below have been made for *educational purposes only*.

Avoid polymorphism in generic comparison function (=):

```

let rec assoc x = function
  [] -> raise Not_found
  | (a,b)::l -> if String.compare a x = 0 then b else assoc x l

let rec remove_assoc x = function
  | [] -> []
  | (a, b as pair) :: l ->
    if String.compare a x = 0 then l else pair :: remove_assoc x l

```

Optim3.ml (based on Optim0.ml): time: 19s

Despite implementation-wise the code being the same (as `String.compare` equals `Pervasives.compare` inside module `String`, and `List.assoc` is like above but uses `Pervasives.compare`), we removed polymorphism, so no longer need the `caml_compare_val` function.

Usually, adding type annotations would be enough (especially useful for numeric types `int`, `float`).

Deforestation means removing intermediate data structures:

```

let read_to_histogram file =
  let input = open_in file in
  let h = empty () and more = ref true in
  try
    while !more do
      Scanf.fscanf input "%[^a-zA-Z0-9]%[a-zA-Z0-9]" 
        (fun _ w ->
          let w = String.lowercase_ascii w in
          ignore (increment h w); more := w <> "")

```

```

done; h
with End_of_file -> h

```

Optim4.ml (based on **Optim1.ml**): time: 0.51s - Garbage collection takes 8% of time - So little because we have eliminated garbage

Defunctionization means computing functor applications by hand. - There was a tool `ocamldefun` but it is out of date - The slight speedup comes from the fact that functor arguments are implemented as records of functions

9.5.4 Comparison of Data Structure Implementations We perform a rough comparison of association lists, tree-based maps, and hashtables. Sets would give the same results. We always create hashtables with initial size 511.

10^7 operations of: adding an association (creation), finding a key that is in the map, finding a key out of a small number of keys not in the map.

First row gives sizes of maps. Time in seconds, to two significant digits.

Create operations:

Structure	2^1	2^2	2^3	2^4	2^5	2^6	2^7	2^8	2^9	2^{10}
assoc	0.25	0.25	0.18	0.19	0.17	0.22	0.19	0.19	0.19	-
list										
tree	0.48	0.81	0.82	1.2	1.6	2.3	2.7	3.6	4.1	5.1
map										
hashtable	2^7	9.1	5.5	4	2.9	2.4	2.1	1.9	1.8	3.7

Structure	2^{11}	2^{12}	2^{13}	2^{14}	2^{15}	2^{16}	2^{17}	2^{18}	2^{19}	2^{20}	2^{21}	2^{22}
tree	6.5	8	9.8	15	19	26	34	41	51	67	80	130
map												
hashtable	2^8	5.6	6.4	8.4	12	15	19	20	22	24	23	33

Found operations:

Structure	2^1	2^2	2^3	2^4	2^5	2^6	2^7	2^8	2^9	2^{10}
assoc	1.1	1.5	2.5	4.2	8.1	17	30	60	120	-
list										
tree	1	1.1	1.3	1.5	1.9	2.1	2.5	2.8	3.1	3.6
map										
hashtable	2^4	1.4	1.5	1.4	1.4	1.5	1.5	1.6	1.6	1.8

Struct	2^{11}	2^{12}	2^{13}	2^{14}	2^{15}	2^{16}	2^{17}	2^{18}	2^{19}	2^{20}	2^{21}	2^{22}
tree	4.3	5.2	6	7.6	9.4	12	15	17	19	24	28	32
map												
hashtable	1.8	2	2.5	3.1	4	5.1	5.9	6.4	6.8	7.6	6.7	7.5

Not found operations:

Structure	2^1	2^2	2^3	2^4	2^5	2^6	2^7	2^8	2^9	2^{10}
assoc	1.8	2.6	4.6	8	16	32	60	120	240	-
list										
tree	1.5	1.5	1.8	2.1	2.4	2.7	3	3.2	3.5	3.8
map										
hashtable	1.4	1.4	1.5	1.5	1.6	1.5	1.7	1.9	2	2.1

Struct	2^{11}	2^{12}	2^{13}	2^{14}	2^{15}	2^{16}	2^{17}	2^{18}	2^{19}	2^{20}	2^{21}	2^{22}
tree	4.2	4.3	4.7	4.9	5.3	5.5	6.1	6.3	6.6	7.2	7.5	7.3
map												
hashtable	1.9	2	1.9	1.9	1.9	2	2	2.2	2	2	2	1.9

Using lists makes sense for up to about 15 elements.

Unfortunately, OCaml and Haskell do not encourage the use of efficient maps the way Scala and Python have built-in syntax for them.

9.6 Parsing: `ocamllex` and `Menhir`

Parsing means transforming text (a string of characters) into a data structure that is well fitted for a given task, or generally makes information in the text more explicit.

Parsing is usually done in stages: 1. *Lexing* or *tokenizing*: dividing the text into smallest meaningful pieces called *lexemes* or *tokens* 2. Composing bigger structures out of lexemes/tokens (and smaller structures) according to a *grammar* - Alternatively to building such hierarchical structure, sometimes we build relational structure over the tokens, e.g., *dependency grammars*

We will use `ocamllex` for lexing, whose rules are like pattern matching functions, but with patterns being *regular expressions*.

We will either consume the results from the lexer directly, or use `Menhir` for parsing, a successor of `ocamlyacc`, belonging to the *yacc/bison* family of parsers.

9.6.1 Lexing with ocamllex The format of lexer definitions is as follows (file with extension .mll):

```
{ header }
let ident1 = regexp ...
rule entrypoint1 [arg1 ... argN] =
  parse regexp { action1 }
  | ...
  | regexp { actionN }
and entrypointN [arg1 ... argN] =
  parse ...
and ...
{ trailer }

• Comments are delimited by (* *), as in OCaml
• The parse keyword can be replaced by the shortest keyword
• “Header”, “trailer”, “action1”, ... “actionN” are arbitrary OCaml code
• There can be multiple let-clauses and rule-clauses
```

Let-clauses are shorthands for regular expressions.

Each rule-clause `entrypoint` defines function(s) that take as the last argument (after `arg1 ... argN` if $N > 0$) an argument `lexbuf` of type `Lexing.lexbuf`. The `lexbuf` is also visible in actions, just as a regular argument. Entry points can be mutually recursive if we need to read more before we can return output.

We can use `lexbuf` in actions:

- `Lexing.lexeme lexbuf` – Return the matched string
- `Lexing.lexeme_char lexbuf n` – Return the nth character in the matched string (first character is $n = 0$)
- `Lexing.lexeme_start/lexeme_end lexbuf` – Return the absolute position in the input text of the beginning/end of the matched string (first character read has offset 0)

The parser will call an `entrypoint` when it needs another lexeme/token.

The syntax of regular expressions:

Pattern	Meaning
'c'	Match the character 'c'
-	Match a single character
eof	Match end of lexer input
"string"	Match the corresponding sequence of characters
[character set]	Match the character set (characters 'c' and ranges 'c'-'d' separated by space)
[^character set]	Match characters outside the character set
[set1] # [set2]	Match the difference (only characters in set1 that are not in set2)

Pattern	Meaning
<code>regexp*</code>	(repetition) Match zero or more strings matching regexp
<code>regexp+</code>	(strict repetition) Match one or more strings matching regexp
<code>regexp?</code>	(option) Match the empty string, or a string matching regexp
<code>regexp1 \ regexp2</code>	(alternative) Match any string that matches regexp1 or regexp2
<code>regexp1 regexp2</code>	(concatenation) Match concatenation of two strings
<code>(regexp)</code>	Match the same strings as regexp
<code>ident</code>	Reference the regular expression bound by <code>let ident = regexp</code>
<code>regexp as ident</code>	Bind the substring matched by regexp to identifier ident

The precedences are: # highest, followed by *, +, ?, concatenation, |, as.

The type of as ident variables can be `string`, `char`, `string option`, or `char option`: - `char` means obviously a single character pattern - `option` means situations like `(regexp as ident)?` or `regexp1|(regexp2 as ident)` - The variables can repeat in the pattern (**unlike** in normal patterns) – meaning both regexps match the same substrings

`ocamllex Lexer.mll` produces the lexer code in `Lexer.ml`. The `ocamlbuild` tool will call `ocamllex` and `ocamlyacc/menhir` if needed.

Unfortunately if the lexer patterns are big we get an error: *transition table overflow, automaton is too big*.

Example: Finding Email Addresses We mine a text file for email addresses that could have been obfuscated.

To compile and run `Emails.mll`, processing a file `email_corpus.xml`:

```
ocamlbuild Emails.native -- email_corpus.xml
```

The lexer header defines types and helper functions:

```
{
    (* The header with OCaml code *)
    open Lexing
    (* Make accessing Lexing easier *)
    let nextline lexbuf =      (* Typical lexer function: move position to next line *)
        let pos = lexbuf.lex_curr_p in
        lexbuf.lex_curr_p <- { pos with
            pos_lnum = pos.pos_lnum + 1;
            pos_bol = pos.pos_cnum;
```

```

        }
type state =          (* Which step of searching for address we're at: *)
| Seek             (* Seek: still seeking, Addr (true...): possibly finished *)
| Addr of bool * string * string list (* Addr (false...): no domain *)

let report state lexbuf =           (* Report the found address, if any *)
  match state with
  | Seek -> ()
  | Addr (false, _, _) -> ()
  | Addr (true, name, addr) ->          (* With line at which it is found *)
    Printf.printf "%d: %s@%s\n" lexbuf.lex_curr_p.pos_lnum
      name (String.concat "." (List.rev addr))
}

```

The lexer defines regular expressions for email patterns and rules for matching:

```

let newline = ('\n' | "\r\n")           (* Regexp for end of line *)
let addr_char = ['a'-'z' 'A'-'Z' '0'-'9' '-' '_']
let at_w_symb = "where" | "WHERE" | "at" | "At" | "AT"
let at_nw_symb = '@' | "&#x40;" | "&#64;"
let open_symb = ' ' * '(' ' ' '*' | ' ' '+'
let close_symb = ' ' * ')' ' ' '*' | ' ' '+'
let at_sep_symb =
  open_symb? at_nw_symb close_symb? |
  open_symb at_w_symb close_symb

let dot_w_symb = "dot" | "DOT" | "dt" | "DT"
let dom_w_symb = dot_w_symb | "dom" | "DOM"           (* Obfuscation for last dot *)
let dot_sep_symb =
  open_symb dot_w_symb close_symb |
  open_symb? '.' close_symb?
let dom_sep_symb =
  open_symb dom_w_symb close_symb |
  open_symb? '.' close_symb?
let addr_dom = addr_char addr_char           (* Restricted form of last *)
  | "edu" | "EDU" | "org" | "ORG" | "com" | "COM"           (* part of address *)

```

The main rule processes the input character by character:

```

rule email state = parse
| newline           (* Check state before moving on *)
  { report state lexbuf; nextline lexbuf;
    email Seek lexbuf }
| (addr_char+ as name) at_sep_symb (addr_char+ as addr) (* Detected possible start *)
  { email (Addr (false, name, [addr])) lexbuf }
| dom_sep_symb (addr_dom as dom)           (* Detected possible finish of address *)
  { let state =
    match state with

```

```

| Seek -> Seek                                (* We weren't looking at an address *)
| Addr (_, name, addrs) ->                   (* Bingo *)
  Addr (true, name, dom::addrs) in
  email state lexbuf }
| dot_sep_symb (addr_char+ as addr)           (* Next part of address -- must be continued *)
{ let state =
  match state with
  | Seek -> Seek
  | Addr (_, name, addrs) ->
    Addr (false, name, addr::addrs) in
    email state lexbuf }
| eof                                         (* End of file -- end loop *)
{ report state lexbuf }
| _                                           (* Some boring character -- not looking at address *)
{ report state lexbuf; email Seek lexbuf }

```

The trailer opens a file and starts mining:

```

{
let _ =                                         (* The trailer with OCaml code *)
  let ch = open_in Sys.argv.(1) in
  email Seek (Lexing.from_channel ch);
  close_in ch                                  (* Close the file at the end *)
}

```

9.6.2 Parsing with Menhir The format of parser definitions is as follows
(file with extension .mly):

```

%{ header %}                                     (* OCaml code put in front *)
%parameter < M : signature >                  (* Parameters make a functor *)
%token < type1 > Token1 Token2                (* Terminal productions, variants returned from lexer *)
%token < type3 > Token3
%token NoArgToken                               (* Without an argument, e.g. keywords or symbols *)
%nonassoc Token1                                (* This token cannot be stacked without parentheses *)
%left Token3                                    (* Associates to left *)
%right Token2                                   (* Associates to right *)
%type < type4 > rule1                         (* Type of the action of the rule *)
%start < type5 > rule2                         (* The entry point of the grammar *)
%%                                         (* Separate out the rules part *)
%inline rule1 (id1, ..., idN) :
| production1 { action1 }                     (* Inlined rules can propagate priorities *)
| production2 | production3 { action2 }        (* If production matches, perform action *)
                                         (* Several productions *)
                                         (* with the same action *)

%public rule2 :                                 (* Visible in other files of the grammar *)
| production4 { action4 }
%public rule3 :                                 (* Override precedence of production5 *)

```

```

| production5 { action5 } %prec Token1 (* to that of Token1 *)
%%                                         (* Separations needed even if sections empty *)
trailer                                (* OCaml code put at the end *)

```

Header, actions, and trailer are OCaml code. Comments are (* ... *) in OCaml code, /* ... */ or // ... outside.

Rules can optionally be separated by ;. The **%parameter** declaration turns the **whole** resulting grammar into a functor; multiple parameters are allowed and visible in %{...%}.

Terminal symbols **Token1** and **Token2** are both variants with argument of type **type1**, called their *semantic value*.

rule1 ... **ruleN** must be lower-case identifiers. Parameters **id1** ... **idN** can be lower- or upper-case.

Priorities (precedence) are declared implicitly: **%nonassoc**, **%left**, **%right** list tokens in increasing priority (**Token2** has highest precedence).

- Higher precedence = a rule is applied even when tokens so far could be part of the other rule
- Precedence of a production comes from its rightmost terminal
- **%left**/**%right** means left/right associativity: the rule will/won't be applied if the "other" rule is the same production

%start symbols become names of functions exported in the **.mli** file to invoke the parser. They are automatically **%public**.

%public rules can even be defined over multiple files, with productions joined by |.

The syntax of productions:

Pattern	Meaning
rule2 Token1 rule3	Match tokens in sequence with Token1 in the middle
a=rule2 t=Token3	Name semantic values produced by rules/tokens
rule2; Token3	Parts of pattern can be separated by semicolon
rule1(arg1,...,argN)	Use a rule that takes arguments
rule2?	Shorthand for option(rule2)
rule2+	Shorthand for nonempty_list(rule2)
rule2*	Shorthand for list(rule2)

Always-visible "standard library" (most of rules shown):

```

%public option(X):
/* nothing */
{ None }
| x = X
{ Some x }

%public %inline pair(X, Y):
x = X; y = Y
{ (x, y) }

%public %inline separated_pair(X, sep, Y):
x = X; sep; y = Y
{ (x, y) }

%public %inline delimited(opening, X, closing):
opening; x = X; closing
{ x }

%public list(X):
/* nothing */
{ [] }
| x = X; xs = list(X)
{ x :: xs }

%public nonempty_list(X):
x = X
{ [ x ] }
| x = X; xs = nonempty_list(X)
{ x :: xs }

%public %inline separated_list(separator, X):
xs = loption(separated_nonempty_list(separator, X))
{ xs }

%public separated_nonempty_list(separator, X):
x = X
{ [ x ] }
| x = X; separator; xs = separated_nonempty_list(separator, X)
{ x :: xs }

```

Only *left-recursive* rules are truly tail-recursive:

```

declarations:
| { [] }
| ds = declarations; option(COMMA);
d = declaration { d :: ds }

```

This is opposite to code expressions (or *recursive descent parsers*) – if both are OK, the first rather than last invocation should be recursive.

Invocations can be nested in arguments:

```
plist(X):
| xs = loption(
    delimited(LPAREN,
              separated_nonempty_list(COMMA, X),
              RPAREN)) { xs }
```

Higher-order parameters are allowed:

```
procedure(list):
| PROCEDURE ID list(formal) SEMICOLON block SEMICOLON {...}
```

Example where inlining is required (besides being an optimization):

```
%token < int > INT
%token PLUS TIMES
%left PLUS
%left TIMES                                (* Multiplication has higher priority *)
%
expression:
| i = INT { i }                            (* Without inlining, would not distinguish priorities *)
| e = expression; o = op; f = expression { o e f }
%inline op:                                  (* Inline operator -- generate corresponding rules *)
| PLUS { ( + ) }
| TIMES { ( * ) }
```

Menhir is an LR(1) parser generator, i.e., it fails for grammars where looking one token ahead, together with precedences, is insufficient to determine whether a rule applies. In particular, only unambiguous grammars are supported.

Although LR(1) grammars are a small subset of *context free grammars*, the semantic actions can depend on context: actions can be functions that take some form of context as input.

Positions are available in actions via keywords `$startpos(x)` and `$endpos(x)` where `x` is a name given to part of pattern.

Note: Do not use the `Parsing` module from OCaml standard library.

Example: Parsing Arithmetic Expressions Example based on a Menhir demo. Due to difficulties with `ocamlbuild`, we use option `--external-tokens` to provide type `token` directly rather than having it generated.

File `lexer.mll`:

```
{  
  type token =
```

```

| TIMES
| RPAREN
| PLUS
| MINUS
| LPAREN
| INT of (int)
| EOL
| DIV
exception Error of string
}

rule line = parse
| (^'\n']* '\n') as line { line }
| eof { exit 0 }

and token = parse
| [' ' '\t'] { token lexbuf }
| '\n' { EOL }
| ['0'-'9']+ as i { INT (int_of_string i) }
| '+' { PLUS }
| '-' { MINUS }
| '*' { TIMES }
| '/' { DIV }
| '(' { LPAREN }
| ')' { RPAREN }
| eof { exit 0 }
| _ { raise (Error (Printf.sprintf "At offset %d: unexpected character.\n"
                                (Lexing.lexeme_start lexbuf))) }

```

File parser.mly:

```

%token <int> INT          (* We still need to define tokens, *)
%token PLUS MINUS TIMES DIV (* Menhir does its own checks *)
%token LPAREN RPAREN
%token EOL
%left PLUS MINUS      /* lowest precedence */
%left TIMES DIV        /* medium precedence */
%nonassoc UMINUS       /* highest precedence */
%parameter<Semantics : sig
  type number
  val inject: int -> number
  val (+): number -> number -> number
  val (-): number -> number -> number
  val (*): number -> number -> number
  val (/): number -> number -> number
  val (~-): number -> number
end>

```

```

%start <Semantics.number> main
%{ open Semantics %}

%%
main:
| e = expr EOL    { e }

expr:
| i = INT      { inject i }
| LPAREN e = expr RPAREN   { e }
| e1 = expr PLUS e2 = expr { e1 + e2 }
| e1 = expr MINUS e2 = expr { e1 - e2 }
| e1 = expr TIMES e2 = expr { e1 * e2 }
| e1 = expr DIV e2 = expr  { e1 / e2 }
| MINUS e = expr %prec UMINUS { - e }

File calc.ml:

module FloatSemantics = struct
  type number = float
  let inject = float_of_int
  let (+) = ( +. )
  let (-) = ( -. )
  let (*) = ( *. )
  let (/) = ( /. )
  let (~-) = ( ~-. )
end

module FloatParser = Parser.Make(FloatSemantics)

let () =
  let stdinbuf = Lexing.from_channel stdin in
  while true do
    let linebuf =
      Lexing.from_string (Lexer.line stdinbuf) in
    try
      Printf.printf "%.\n%!"
      (FloatParser.main Lexer.token linebuf)
    with
    | Lexer.Error msg ->
      Printffprintf stderr "%s%!" msg
    | FloatParser.Error ->
      Printffprintf stderr
        "At offset %d: syntax error.\n%!"
        (Lexing.lexeme_start linebuf)
  done

```

Build and run command:

```
ocamlbuild calc.native -use-menhir -menhir "menhir parser.mly --base  
parser --external-tokens Lexer" --
```

- Other grammar files can be provided besides `parser.mly`
- `--base` gives the file (without extension) which will become the module accessed from OCaml
- `--external-tokens` provides the OCaml module which defines the `token` type

Example: A Toy Sentence Grammar Our lexer is a simple limited *part-of-speech tagger*. Not re-entrant.

For debugging, we log execution in file `log.txt`.

File `EngLexer.mll`:

```
{  
  type sentence = {  
    subject : string; (* Could be in any module visible to EngParser *)  
    action : string; (* The actor/actors, i.e. subject noun *)  
    plural : bool; (* The action, i.e. verb *)  
    adjs : string list; (* Whether one or multiple actors *)  
    advs : string list; (* Characteristics of actor *)  
  }  
  
  type token =  
  | VERB of string  
  | NOUN of string  
  | ADJ of string  
  | ADV of string  
  | PLURAL | SINGULAR  
  | A_DET | THE_DET | SOME_DET | THIS_DET | THAT_DET  
  | THESE_DET | THOSE_DET  
  | COMMA_CNJ | AND_CNJ | DOT_PUNCT  
  
  let adjectives = (* Recognized adjectives *)  
    ["smart"; "extreme"; "green"; "slow"; "old"; "incredible";  
     "quiet"; "diligent"; "mellow"; "new"]  
  let log_file = open_out "log.txt" (* File with debugging information *)  
  let log s = Printffprintf log_file "%s\n%"! s  
  let last_tok = ref DOT_PUNCT (* State for better tagging *)  
  
  let tokbuf = Queue.create () (* Token buffer, since single word *)  
  let push w = (* is sometimes two tokens *)  
    last_tok := w; Queue.push w tokbuf  
  exception LexError of string
```

```

}

let alphanum = ['0'-'9' 'a'-'z' 'A'-'Z' ' ' '-']

rule line = parse
| ([~'\n']* '\n') as l { l }
| eof { exit 0 }

(* For line-based interface *)

and lex_word = parse
| ' ' '\t'
  { lex_word lexbuf } (* Skip whitespace *)
| '.'
  { push DOT_PUNCT } (* End of sentence *)
| "a" { push A_DET } | "the" { push THE_DET } (* Keywords *)
| "some" { push SOME_DET }
| "this" { push THIS_DET } | "that" { push THAT_DET }
| "these" { push THESE_DET } | "those" { push THOSE_DET }
| "A" { push A_DET } | "The" { push THE_DET }
| "Some" { push SOME_DET }
| "This" { push THIS_DET } | "That" { push THAT_DET }
| "These" { push THESE_DET } | "Those" { push THOSE_DET }
| "and" { push AND_CNJ }
| ',' { push COMMA_CNJ }
| (alphanum+ as w) "ly" (* Adverb is adjective that ends in "ly" *)
{
  if List.mem w adjectives
  then push (ADV w)
  else if List.mem (w^"le") adjectives
  then push (ADV (w^"le"))
  else (push (NOUN w); push SINGULAR)
}

| (alphanum+ as w) "s" (* Plural noun or singular verb *)
{
  if List.mem w adjectives then push (ADJ w)
  else match !last_tok with
    | THE_DET | SOME_DET | THESE_DET | THOSE_DET
    | DOT_PUNCT | ADJ _ ->
      push (NOUN w); push PLURAL
    | _ -> push (VERB w); push SINGULAR
}
| alphanum+ as w (* Noun contexts vs. verb contexts *)
{
  if List.mem w adjectives then push (ADJ w)
  else match !last_tok with
    | A_DET | THE_DET | SOME_DET | THIS_DET | THAT_DET
    | DOT_PUNCT | ADJ _ ->

```

```

        push (NOUN w); push SINGULAR
    | _ -> push (VERB w); push PLURAL
}

| _ as w
{ raise (LexError ("Unrecognized character " ^
                    Char.escaped w)) }
{
let lexeme lexbuf =      (* The proper interface reads from the token buffer *)
  if Queue.is_empty tokbuf then lex_word lexbuf;
  Queue.pop tokbuf
}
File EngParser.mly:
%{
open EngLexer                      (* Source of the token type and sentence type *)
%}
%token <string> VERB NOUN ADJ ADV      (* Open word classes *)
%token PLURAL SINGULAR                (* Number marker *)
%token A_DET THE_DET SOME_DET THIS_DET THAT_DET   (* "Keywords" *)
%token THESE_DET THOSE_DET
%token COMMA_CNJ AND_CNJ DOT_PUNCT
%start <EngLexer.sentence> sentence    (* Grammar entry *)
%%

%public %inline sep2_list(sep1, sep2, X):          (* General purpose *)
| xs = separated_nonempty_list(sep1, X) sep2 x=X
  { xs @ [x] }                                (* We use it for "comma-and" lists: *)
| x=option(X)                                 (* smart, quiet and diligent *)
  { match x with None->[] | Some x->[x] }

sing_only_det:                                (* How determiners relate to number *)
| A_DET | THIS_DET | THAT_DET { log "prs: sing_only_det" }

plu_only_det:
| THESE_DET | THOSE_DET { log "prs: plu_only_det" }

other_det:
| THE_DET | SOME_DET { log "prs: other_det" }

np(det):
| det adjs=list(ADJ) subject=NOUN
  { log "prs: np"; adjs, subject }

vp(NUM):
| advs=separated_list(AND_CNJ,ADV) action=VERB NUM

```

```

| action=VERB NUM advs=sep2_list(COMMA_CNJ,AND_CNJ,ADV)
  { log "prs: vp"; action, advs }

sent(det,NUM):                                     (* Sentence parameterized by number *)
| adjsub=np(det) NUM vbadv=vp(NUM)
  { log "prs: sent";
    {subject=snd adjsub; action=fst vbadv; plural=false;
     adjs=fst adjsub; advs=snd vbadv} }

vbsent(NUM):                                     (* Unfortunately, it doesn't always work... *)
| NUM vbadv=vp(NUM)   { log "prs: vbsent"; vbadv }

sentence:                                         (* Sentence, either singular or plural *)
| s=sent(sing_only_det,SINGULAR) DOT_PUNCT
  { log "prs: sentence1";
    {s with plural = false} }
| s=sent(plu_only_det,PLURAL) DOT_PUNCT
  { log "prs: sentence2";
    {s with plural = true} }

| adjsub=np(other_det) vbadv=vbsent(SINGULAR) DOT_PUNCT
  { log "prs: sentence3";      (* Because parser allows only one token look-ahead *)
    {subject=snd adjsub; action=fst vbadv; plural=false;
     adjs=fst adjsub; advs=snd vbadv} }
| adjsub=np(other_det) vbadv=vbsent(PLURAL) DOT_PUNCT
  { log "prs: sentence4";      (* we need to factor-out the "common subset" *)
    {subject=snd adjsub; action=fst vbadv; plural=true;
     adjs=fst adjsub; advs=snd vbadv} }

```

File Eng.ml:

```

open EngLexer

let () =
  let stdinbuf = Lexing.from_channel stdin in
  while true do
    (* Read line by line. *)
    let linebuf = Lexing.from_string (line stdinbuf) in

    try
      (* Run the parser on a single line of input. *)
      let s = EngParser.sentence lexeme linebuf in
      Printf.printf
        "subject=%s\nplural=%b\nadjs=%s\naction=%s\nadvs=%s\n\n!""
          s.subject s.plural (String.concat ", " s.adjs)
          s.action (String.concat ", " s.advs)
    with

```

```

| LexError msg ->
  Printffprintf stderr "%s\n%" msg
| EngParser.Error ->
  Printffprintf stderr "At offset %d: syntax error.\n%""
    (Lexing.lexeme_start linebuf)
done

```

Build and run command:

```
ocamlbuild Eng.native -use-menhir -menhir "menhir EngParser.mly
--base EngParser --external-tokens EngLexer" --
```

9.7 Example: Phrase Search

In lecture 6 we performed keyword search; now we turn to *phrase search*, i.e., require that given words be consecutive in the document.

We start with some English-specific transformations used in the lexer:

```

let wh_or_pronoun w =
  w = "where" || w = "what" || w = "who" ||
  w = "he" || w = "she" || w = "it" ||
  w = "I" || w = "you" || w = "we" || w = "they"

let abridged w1 w2 =                                     (* Remove shortened forms like I'll or press'd
  if w2 = "ll" then [w1; "will"]
  else if w2 = "s" then
    if wh_or_pronoun w1 then [w1; "is"]
    else ["of"; w1]
  else if w2 = "d" then [w1^"ed"]
  else if w1 = "o" || w1 = "O"
  then
    if w2.[0] = 'e' && w2.[1] = 'r' then [w1^"v"^w2]
    else ["of"; w2]
  else if w2 = "t" then [w1; "it"]
  else [w1^"'"^w2]

```

For now we normalize words just by lowercasing (but see exercise 8 for Porter stemming).

In the lexer we *tokenize* text: separate words and normalize them. We also handle simple aspects of *XML* syntax. We store the number of each word occurrence, excluding XML tags.

Parsing: the inverted index and the query:

```

type token =
| WORDS of (string * int) list
| OPEN of string | CLOSE of string | COMMENT of string
| SENTENCE of string | PUNCT of string

```

```

| EOF

let inv_index update ii lexer lexbuf =
  let rec aux ii =
    match lexer lexbuf with
    | WORDS ws ->
      let ws = List.map (fun (w,p)->EngMorph.normalize w, p) ws in
      aux (List.fold_left update ii ws)
    | OPEN _ | CLOSE _ | SENTENCE _ | PUNCT _ | COMMENT _ ->
      aux ii
    | EOF -> ii in
  aux ii

let phrase lexer lexbuf =
  let rec aux words =
    match lexer lexbuf with
    | WORDS ws ->
      let ws = List.map (fun (w,p)->EngMorph.normalize w) ws in
      aux (List.rev_append ws words)
    | OPEN _ | CLOSE _ | SENTENCE _ | PUNCT _ | COMMENT _ ->
      aux words
    | EOF -> List.rev words in
  aux []

```

9.7.1 Naive Implementation of Phrase Search We need *postings lists* with positions of words rather than just the document or line of document they belong to.

First approach: association lists and merge postings lists word-by-word.

```

let update ii (w, p) =
  try
    let ps = List.assoc w ii in (* Add position to the postings list of w *)
    (w, p::ps) :: List.remove_assoc w ii
  with Not_found -> (w, [p]):::ii

let empty = []
let find w ii = List.assoc w ii
let mapv f ii = List.map (fun (k,v)->k, f v) ii

let index file =
  let ch = open_in file in
  let lexbuf = Lexing.from_channel ch in
  EngLexer.reset_as_file lexbuf file;
  let ii =
    IndexParser.inv_index update empty EngLexer.token lexbuf in

```

```

close_in ch;
mapv List.rev ii, List.rev !EngLexer.linebreaks (* Keep postings lists in increasing order *)

let find_line linebreaks p = (* Recover the line in document of a position *)
  let rec aux line =
    | [] -> line
    | bp::_ when p < bp -> line
    | _::breaks -> aux (line+1) breaks in
  aux 1 linebreaks

let search (ii, linebreaks) phrase =
  let lexbuf = Lexing.from_string phrase in
  EngLexer.reset_as_file lexbuf ("search phrase: ^phrase");
  let phrase = IndexParser.phrase EngLexer.token lexbuf in
  let rec aux wpos = function (* Merge postings lists for words in query: *)
    | [] -> wpos (* no more words in query *)
    | w::ws -> (* for positions of w, keep those that are next to previous *)
      let nwpos = find w ii in (* filtered positions of previous word *)
      aux (List.filter (fun p->List.mem (p-1) wpos) nwpos) ws in
  let wpos =
    match phrase with
    | [] -> [] (* No results for an empty query *)
    | w::ws -> aux (find w ii) ws in
  List.map (find_line linebreaks) wpos (* Answer in terms of document lines *)

let shakespeare = index "./shakespeare.xml"

let query q =
  let lines = search shakespeare q in
  Printf.printf "%s: lines %s\n!" q
  (String.concat ", " (List.map string_of_int lines))

```

Test: 200 searches of the queries:

```

["first witch"; "wherefore art thou";
 "captain's captain"; "flatter'd"; "of Fulvia";
 "that which we call a rose"; "the undiscovered country"]

```

Invocation: ocamlbuild InvIndex.native -libs unix --

Time: 7.3s

9.7.2 Replace Association List with Hash Table I recommend using either *OCaml Batteries* or *OCaml Core* – replacement for the standard library. *Batteries* has efficient Hashtbl.map (our mapv).

Invocation: ocamlbuild InvIndex1.native -libs unix --

Time: 6.3s

9.7.3 Replace Naive Merging with Ordered Merging Postings lists are already ordered.

Invocation: `ocamlbuild InvIndex2.native -libs unix --`

Time: 2.5s

9.7.4 Bruteforce Optimization: Biword Indexes Pairs of words are much less frequent than single words, so storing them means less work for postings lists merging.

Can result in much bigger index size: $\min(W^2, N)$ where W is the number of distinct words and N the total number of words in documents.

Invocation that gives us stack backtraces:

```
ocamlbuild InvIndex3.native -cflag -g -libs unix; export OCAMLRUNPARAM="b"; ./InvIndex3.native
```

Time: 2.4s – disappointing.

9.7.5 Smart Way: Information Retrieval (G.V. Cormack et al.) You should classify your problem and search literature for state-of-the-art algorithms to solve it.

The algorithm needs a data structure for inverted index that supports:

- `first(w)` – first position in documents at which `w` appears
- `last(w)` – last position of `w`
- `next(w, cp)` – first position of `w` after position `cp`
- `prev(w, cp)` – last position of `w` before position `cp`

We develop `next` and `prev` operations in stages: 1. First, a naive (but FP) approach using the `Set` module of OCaml - We could use our balanced binary search tree implementation to avoid the overhead due to limitations of `Set` API 2. Then, *binary search* based on arrays 3. Imperative linear search 4. Imperative *galloping search* optimization of binary search

The Phrase Search Algorithm During search we maintain *current position* `cp` of the last found word or phrase.

The algorithm is almost purely functional; we use the `Not_found` exception instead of option type for convenience.

```
let rec next_phrase ii phrase cp =      (* Return the beginning and end position *)
  let rec aux cp = function
    | [] -> raise Not_found           (* of occurrence of phrase after position cp *)
    | [w] ->                          (* Empty phrase counts as not occurring *)
    | w::ws ->                      (* Single or last word of phrase has the same *)
      let np = next ii w cp in np, np   (* beg. and end position *)
      | _ -> aux (next ii w cp) ws in (* After locating the endp. move back *)
```

```

    prev ii w np, fp in
let np, fp = aux cp phrase in          (* If distance is this small, *)
if fp - np = List.length phrase - 1 then np, fp (* words are consecutive *)
else next_phrase ii phrase fp

let search (ii, linebreaks) phrase =
  let lexbuf = Lexing.from_string phrase in
  EngLexer.reset_as_file lexbuf ("search phrase: " ^ phrase);
  let phrase = IndexParser.phrase EngLexer.token lexbuf in
  let rec aux cp =
    try                                (* Find all occurrences of the phrase *)
      let np, fp = next_phrase ii phrase cp in
      np :: aux fp
    with Not_found -> [] in           (* Moved past last occurrence *)
    List.map (find_line linebreaks) (aux (-1))

```

Naive but Purely Functional Inverted Index

```

module S = Set.Make(struct type t=int let compare i j = i-j end)

let update ii (w, p) =
  (try
    let ps = Hashtbl.find ii w in
    Hashtbl.replace ii w (S.add p ps)
  with Not_found -> Hashtbl.add ii w (S.singleton p));
  ii

let first ii w = S.min_elt (find w ii)          (* The functions raise Not_found *)
let last ii w = S.max_elt (find w ii)           (* whenever such position would not exist *)

let prev ii w cp =
  let ps = find w ii in                         (* Split the set into elements *)
  let smaller, _, _ = S.split cp ps in           (* smaller and bigger than cp *)
  S.max_elt smaller

let next ii w cp =
  let ps = find w ii in
  let _, _, bigger = S.split cp ps in
  S.min_elt bigger

Invocation: ocamlbuild InvIndex4.native -libs unix --
Time: 3.3s – would be better without the overhead of S.split.

```

Binary Search Based Inverted Index

```
let prev ii w cp =
```

```

let ps = find w ii in
let rec aux b e =
  if e-b <= 1 then ps.(b) (* We implement binary search separately for prev *)
  else let m = (b+e)/2 in
    if ps.(m) < cp then aux m e
    else aux b m in
let l = Array.length ps in
if l = 0 || ps.(0) >= cp then raise Not_found
else aux 0 (l-1)

let next ii w cp =
  let ps = find w ii in
  let rec aux b e =
    if e-b <= 1 then ps.(e) (* and here more than cp *)
    else let m = (b+e)/2 in
      if ps.(m) <= cp then aux m e
      else aux b m in
  let l = Array.length ps in
  if l = 0 || ps.(l-1) <= cp then raise Not_found
  else aux 0 (l-1)

```

File: InvIndex5.ml. Time: 2.4s

Imperative, Linear Scan

```

let prev ii w cp =
  let cw,ps = find w ii in (* For each word we add a cell with last visited occurrence
  let l = Array.length ps in
  if l = 0 || ps.(0) >= cp then raise Not_found
  else if ps.(l-1) < cp then cw := l-1
  else (
    (* Reset pointer if current position is not "ahead" of it *)
    if !cw < l-1 && ps.(!cw+1) < cp then cw := l-1; (* Otherwise scan *)
    while ps.(!cw) >= cp do decr cw done (* starting from last visited *)
  );
  ps.(!cw)

let next ii w cp =
  let cw,ps = find w ii in
  let l = Array.length ps in
  if l = 0 || ps.(l-1) <= cp then raise Not_found
  else if ps.(0) > cp then cw := 0
  else (
    (* Reset pointer if current position is not ahead of it *)
    if !cw > 0 && ps.(!cw-1) > cp then cw := 0;
    while ps.(!cw) <= cp do incr cw done
  );
  ps.(!cw)

```

End of index-building function:

```
mapv (fun ps->ref 0, Array.of_list (List.rev ps)) ii,...
```

File: InvIndex6.ml. Time: 2.8s

Imperative, Galloping Search

```
let next ii w cp =
  let cw,ps = find w ii in
  let l = Array.length ps in
  if l = 0 || ps.(l-1) <= cp then raise Not_found;
  let rec jump (b,e as bounds) j = (* Locate the interval with cp inside *)
    if e < l-1 && ps.(e) <= cp then jump (e,e+j) (2*j)
    else bounds in
  let rec binse b e = (* Binary search over that interval *)
    if e-b <= 1 then e
    else let m = (b+e)/2 in
      if ps.(m) <= cp then binse m e
      else binse b m in
  if ps.(0) > cp then cw := 0
  else (
    let b =
      if !cw > 0 && ps.(!cw-1) <= cp then !cw-1 else 0 in
    let b,e = jump (b,b+1) 2 in (* Locate interval starting near !cw *)
    let e = if e > l-1 then l-1 else e in
    cw := binse b e
  );
  ps.(!cw)
```

prev is symmetric to next.

File: InvIndex7.ml. Time: 2.4s – minimal speedup in our simple test case.

9.8 Exercises

1. (Exercise 6.1 from “*Modern Compiler Implementation in ML*” by Andrew W. Appel.) Using the `ocamlopt` compiler with parameter `-S` and other parameters turning on all possible compiler optimizations, evaluate the compiled programs by these criteria:

- Are local variables kept in registers? Show on an example.
- If local variable `b` is live across more than one procedure call, is it kept in a callee-save register? Explain how it would speed up the program: `let f a = let b = a+1 in let c = g () in let d = h c in b+c`
- If local variable `x` is never live across a procedure call, is it properly kept in a caller-save register? Explain how doing this would speed up the program: `let h y = let x = y+1 in let z = f y in f z`

2. As above, verify whether escaping variables of a function are kept in a closure corresponding to the function, or in closures corresponding to the local (nested) functions that are returned from the function (or assigned to a mutable field).
3. As above, verify that OCaml compiler performs *inline expansion* of small functions. Check whether the compiler can inline, or specialize (produce a local function to help inlining), recursive functions.
4. Write a `.mll` program that anonymizes (or masks) text. That is, it replaces identified probable full names (of persons, companies etc.) with fresh shorthands *Mr. A*, *Ms. B*, or *Mr./Ms. C* when the gender cannot be easily determined. The same (full) name should be replaced with the same letter.
 - Do only a very rough job of course, starting with recognizing two or more capitalized words in a row.
5. In the lexer `EngLexer` we call function `abridged` from the module `EngMorph`. Inline the operation of `abridged` into the lexer by adding a new regular expression pattern for each if clause. Assess the speedup on the *Shakespeare* corpus and the readability, and either keep the change or revert it.
6. Make the lexer re-entrant for the second Menhir example (toy English grammar parser).
7. Make the determiner optional in the toy English grammar.
 - (*) Can you come up with a factorization that would avoid having two more productions in total?
8. Integrate into the *Phrase search* example the *Porter Stemmer* whose source is in the `stemmer.ml` file.
9. Revisit the search engine example from lecture 6.
 - Perform optimization of data structure, i.e. replace association lists with hash tables.
 - Optimize the algorithm: perform *query optimization*. Measure time gains for selected queries.
 - For bonus points, as time and interest permits, extend the query language with *OR* and *NOT* connectives, in addition to *AND*.
 - () Extend query optimization to the query language with AND, OR* and NOT connectives.
10. Write an XML parser tailored to the `shakespeare.xml` corpus provided with the phrase search example. Modify the phrase search engine to provide detailed information for each found location, e.g. which play and who speaks the phrase.

Chapter 10: Functional Reactive Programming

This chapter explores techniques for dealing with change and interaction in functional programming. We begin with zippers, a data structure for navigating and modifying positions within larger structures, then progress to adaptive (incremental) programming and Functional Reactive Programming (FRP). We conclude with practical examples including graphical user interfaces.

Recommended Reading: - “Zipper” in Haskell Wikibook and “*The Zipper*” by Gerard Huet - “*How froc works*” by Jacob Donham - “*The Haskell School of Expression*” by Paul Hudak - “*Deprecating the Observer Pattern with Scala.React*” by Ingo Maier, Martin Odersky

10.1 Zippers

Often we need to keep track of a position within a data structure: easily access and modify it at that location, and easily move the location around. Recall from earlier chapters how we defined *context types* for datatypes – types that represent a data structure with one of its elements missing.

Consider binary trees:

```
type btree = Tip | Node of int * btree * btree
```

Using our algebraic datatype calculus, where T represents the tree type:

$$\begin{aligned} T &= 1 + xT^2 \\ \frac{\partial T}{\partial x} &= 0 + T^2 + 2xT\frac{\partial T}{\partial x} = TT + 2xT\frac{\partial T}{\partial x} \end{aligned}$$

This derivative gives us the context type:

```
type btree_dir = LeftBranch | RightBranch
type btree_deriv =
  | Here of btree * btree
  | Below of btree_dir * int * btree * btree_deriv
```

The key insight is that **Location = context + subtree!** However, there is a problem with the representation above: we cannot easily move the location if `Here` is at the bottom. The part closest to the location should be on top.

Revisiting the Equations Let us revisit the equations for trees and lists:

$$\begin{aligned} T &= 1 + xT^2 \\ \frac{\partial T}{\partial x} &= 0 + T^2 + 2xT\frac{\partial T}{\partial x} \\ \frac{\partial T}{\partial x} &= \frac{T^2}{1-2xT} \\ L(y) &= 1 + yL(y) \\ L(y) &= \frac{1}{1-y} \\ \frac{\partial T}{\partial x} &= T^2L(2xT) \end{aligned}$$

This tells us that the context can be stored as a list with the root as the last node. It does not matter whether we use built-in lists or a type with `Above` and `Root` variants.

Contexts of subtrees are more useful than contexts of single elements:

```
type 'a tree = Tip | Node of 'a tree * 'a * 'a tree
type tree_dir = Left_br | Right_br
type 'a context = (tree_dir * 'a * 'a tree) list
type 'a location = {sub: 'a tree; ctx: 'a context}

let access {sub} = sub
let change {ctx} sub = {sub; ctx}
let modify f {sub; ctx} = {sub = f sub; ctx}
```

We can imagine a location as a rooted tree which is hanging pinned at one of its nodes. For visualizations, see <http://en.wikibooks.org/wiki/Haskell/Zippers>.

Moving Around Navigation functions allow us to traverse the structure:

```
let ascend loc =
  match loc.ctx with
  | [] -> loc (* Or raise exception *)
  | (Left_br, n, l) :: up_ctx ->
    {sub = Node (l, n, loc.sub); ctx = up_ctx}
  | (Right_br, n, r) :: up_ctx ->
    {sub = Node (loc.sub, n, r); ctx = up_ctx}

let desc_left loc =
  match loc.sub with
  | Tip -> loc (* Or raise exception *)
  | Node (l, n, r) ->
    {sub = l; ctx = (Right_br, n, r) :: loc.ctx}

let desc_right loc =
  match loc.sub with
  | Tip -> loc (* Or raise exception *)
  | Node (l, n, r) ->
    {sub = r; ctx = (Left_br, n, l) :: loc.ctx}
```

Trees with Arbitrary Branching Following *The Zipper* by Gerard Huet, let us look at a tree with an arbitrary number of branches, representing a document structure:

```
type doc = Text of string | Line | Group of doc list
type context = (doc list * doc list) list
type location = {sub: doc; ctx: context}
```

The navigation functions for this more complex structure:

```

let go_up loc =
  match loc.ctx with
  | [] -> invalid_arg "go_up: at top"
  | (left, right) :: up_ctx -> (* Previous subdocument and its siblings *)
    {sub = Group (List.rev left @ loc.sub :: right); ctx = up_ctx}

let go_left loc =
  match loc.ctx with
  | [] -> invalid_arg "go_left: at top"
  | (l :: left, right) :: up_ctx -> (* Left sibling of previous subdocument *)
    {sub = l; ctx = (left, loc.sub :: right) :: up_ctx}
  | [], _ :: _ -> invalid_arg "go_left: at first"

let go_right loc =
  match loc.ctx with
  | [] -> invalid_arg "go_right: at top"
  | (left, r :: right) :: up_ctx ->
    {sub = r; ctx = (loc.sub :: left, right) :: up_ctx}
  | _, [] :: _ -> invalid_arg "go_right: at last"

let go_down loc = (* Go to the first (i.e. leftmost) subdocument *)
  match loc.sub with
  | Text _ -> invalid_arg "go_down: at text"
  | Line -> invalid_arg "go_down: at line"
  | Group [] -> invalid_arg "go_down: at empty"
  | Group (doc :: docs) -> {sub = doc; ctx = ([], docs) :: loc.ctx}

```

10.2 Example: Context Rewriting

Imagine a friend working on string theory asks us for help simplifying equations. The task is to pull out particular subexpressions as far to the left as possible, while changing the whole expression as little as possible.

We can illustrate our algorithm using mathematical notation. Let: $-x$ be the thing we pull out - $C[e]$ and $D[e]$ be big expressions with subexpression e - operator \circ stand for one of: $*$, $+$

The rewriting rules are:

$$\begin{array}{ll}
 D[(C[x] \circ e_1) \circ e_2] & \Rightarrow D[C[x] \circ (e_1 \circ e_2)] \\
 D[e_2 \circ (C[x] \circ e_1)] & \Rightarrow D[C[x] \circ (e_1 \circ e_2)] \\
 D[(C[x] + e_1)e_2] & \Rightarrow D[C[x]e_2 + e_1e_2] \\
 D[e_2(C[x] + e_1)] & \Rightarrow D[C[x]e_2 + e_1e_2] \\
 D[e \circ C[x]] & \Rightarrow D[C[x] \circ e]
 \end{array}$$

First, the groundwork:

```
type op = Add | Mul
type expr = Val of int | Var of string | App of expr * op * expr
type expr_dir = Left_arg | Right_arg
type context = (expr_dir * op * expr) list
type location = {sub: expr; ctx: context}
```

To locate the subexpression described by predicate p:

```
let rec find_aux p e =
  if p e then Some (e, [])
  else match e with
    | Val _ | Var _ -> None
    | App (l, op, r) ->
      match find_aux p l with
      | Some (sub, up_ctx) ->
        Some (sub, (Right_arg, op, r) :: up_ctx)
      | None ->
        match find_aux p r with
        | Some (sub, up_ctx) ->
          Some (sub, (Left_arg, op, l) :: up_ctx)
        | None -> None

let find p e =
  match find_aux p e with
  | None -> None
  | Some (sub, ctx) -> Some {sub; ctx = List.rev ctx}
```

Now we can implement the pull-out transformation:

```
let rec pull_out loc =
  match loc.ctx with
  | [] -> loc (* Done *)
  | (Left_arg, op, l) :: up_ctx ->
    (* D[e . C[x]] => D[C[x] . e] *)
    pull_out {loc with ctx = (Right_arg, op, l) :: up_ctx}
  | (Right_arg, op1, e1) :: (_, op2, e2) :: up_ctx
    when op1 = op2 ->
    (* D[(C[x] . e1) . e2] / D[e2 . (C[x] . e1)] => D[C[x] . (e1 . e2)] *)
    pull_out {loc with ctx = (Right_arg, op1, App(e1, op1, e2)) :: up_ctx}
  | (Right_arg, Add, e1) :: (_, Mul, e2) :: up_ctx ->
    (* D[(C[x] + e1) * e2] / D[e2 * (C[x] + e1)] => D[C[x] * e2 + e1 * e2] *)
    pull_out {loc with ctx =
      (Right_arg, Mul, e2) :: (Right_arg, Add, App(e1, Mul, e2)) :: up_ctx}
  | (Right_arg, op, r) :: up_ctx -> (* Move up the context *)
    pull_out {sub = App(loc.sub, op, r); ctx = up_ctx}
```

Since operators are commutative, we ignore the direction for the second piece of context above.

Testing the implementation:

```
let (+) a b = App (a, Add, b)
let (*) a b = App (a, Mul, b)
let (!) a = Val a
let x = Var "x"
let y = Var "y"
let ex = !5 + y * (!7 + x) * (!3 + y)
let loc = find (fun e -> e = x) ex
let sol =
  match loc with
  | None -> raise Not_found
  | Some loc -> pull_out loc
(* Result: "((x*y)*(3+y))+((7*y)*(3+y))+5)" *)
```

For best results, we can iterate the `pull_out` function until a fixpoint is reached.

10.3 Adaptive Programming (Incremental Computing)

While zippers are useful, they are somewhat unnatural for general-purpose programming. Once we change the data structure, it is difficult to propagate the changes – we would need to rewrite all algorithms to work on context changes.

In *Adaptive Programming*, also known as *incremental computation* or *self-adjusting computation*, we write programs in a straightforward functional manner, but can later modify any data causing only the minimal amount of work required to update results.

The functional description of computation is within a monad. We can change monadic values – for example, parts of input – from outside and propagate the changes. In the *Froc* library, the monadic *changeables* are '`a Froc_sa.t`', and the ability to modify them is exposed by type '`'a Froc_sa.u`' – the *writeables*.

Dependency Graphs The monadic value '`'a changeable`' will be the *dependency graph* of the computation of the represented value '`'a`'. Consider the computation:

```
let u = v / w + x * y + z
```

The dependency graph shows how the result depends on the inputs. When we modify inputs `v` and `z` simultaneously, we need to update intermediate nodes in the correct order. For example, we need to update `n2` before `u`.

We use the order of computation (represented as gray numbers in the dependency visualization) for the order of updates. Similarly to `parallel` in the concurrency monad, we provide `bind2`, `bind3`, etc., and corresponding `lift2`, `lift3`, etc., to introduce nodes with several children:

```

let n0 = bind2 v w (fun v w -> return (v / w))
let n1 = bind2 x y (fun x y -> return (x * y))
let n2 = bind2 n0 n1 (fun n0 n1 -> return (n0 + n1))
let u = bind2 n2 z (fun n2 z -> return (n2 + z))

```

Do-notation is not necessary to have readable expressions:

```

let (/) = lift2 (/)
let (* ) = lift2 ( * )
let (+) = lift2 (+)
let u = v / w + x * y + z

```

As in other monads, we can decrease overhead by using bigger chunks:

```

let n0 = blift2 v w (fun v w -> v / w)
let n2 = blift3 n0 x y (fun n0 x y -> n0 + x * y)
let u = blift2 n2 z (fun n2 z -> n2 + z)

```

Handling Conditional Dependencies We have a problem if we recompute all nodes simply by order of computation:

```

let b = x >>= fun x -> return (x = 0)
let n0 = x >>= fun x -> return (100 / x)
let y = bind2 b n0 (fun b n0 -> if b then return 0 else n0)

```

Rather than a single “time” stamp, we store intervals: begin and end of computation. When updating the `y` node, we first detach nodes in the range 4-9 from the graph. Computing the expression will re-attach the nodes as needed.

When the value of `b` does not change, we skip updating `y` and proceed with updating `n0`. The value of `y` is a link to the value of `n0` so it will change anyway.

We need memoization to re-attach the same nodes in case they do not need updating. Are they up-to-date? Run updating past the node’s timestamp range.

Example Using Froc The `Froc_sa` (for *self-adjusting*) module exports the monadic type `t` for changeable computation, and a handle type `u` for updating the computation:

```
open Froc_sa
```

```

type tree = (* Binary tree with nodes storing their screen location *)
| Leaf of int * int (* We will grow the tree *)
| Node of int * int * tree t * tree t (* by modifying subtrees *)

```

Displaying the tree is a changeable effect. Whenever the tree changes, displaying will be updated. Only new nodes will be drawn after an update:

```

let rec display px py t =
  match t with
  | Leaf (x, y) ->

```

```

    return
      (Graphics.draw_poly_line [|px, py; x, y|]; (* We return *)
       Graphics.draw_circle x y 3) (* a throwaway value *)
  | Node (x, y, l, r) ->
    return (Graphics.draw_poly_line [|px, py; x, y|])
  >>= fun _ -> l >>= display x y
  >>= fun _ -> r >>= display x y

```

Growing the tree:

```

let grow_at (x, depth, upd) =
  let x_l = x - f2i (width *. (2.0 ** (~-. (i2f (depth + 1))))) in
  let l, upd_l = changeable (Leaf (x_l, (depth + 1) * 20)) in
  let x_r = x + f2i (width *. (2.0 ** (~-. (i2f (depth + 1))))) in
  let r, upd_r = changeable (Leaf (x_r, (depth + 1) * 20)) in
  write upd (Node (x, depth * 20, l, r)); (* Update the old leaf *)
  propagate (); (* and keep handles to make future updates *)
  [x_l, depth + 1, upd_l; x_r, depth + 1, upd_r]

```

The main loop:

```

let rec loop t subts steps =
  if steps <= 0 then ()
  else loop t (concat_map grow_at subts) (steps - 1)

let incremental steps () =
  Graphics.open_graph " 1024x600";
  let t, u = changeable (Leaf (512, 20)) in
  let d = t >>= display (f2i (width /. 2.)) 0 in (* Display once *)
  loop t [512, 1, u] steps; (* new nodes will be drawn automatically *)
  Graphics.close_graph ()

```

Unfortunately, the overhead of incremental computation is quite large. Comparing byte code execution times:

depth	12	13	14	15	16	17	18	19	20
incremental	0.166s	1s	2.2s	4.4s	9.3s	21s	50s	140s	255s
rebuilding	0.5s	0.63s	1.3s	3s	5.3s	13s	39s	190s	—

10.4 Functional Reactive Programming

FRP is an attempt to declaratively deal with time. *Behaviors* are functions of time – a behavior has a specific value in each instant. *Events* are sets of (time, value) pairs, organized into streams of actions.

Two problems arise in FRP: 1. Behaviors and events are well-defined when they do not depend on future 2. Efficiency: minimize overhead

FRP is *synchronous*: it is possible to set up for events to happen at the same time. It is also *continuous*: behaviors can have details at arbitrary time resolution. Although the results are *sampled*, there is no fixed (minimal) time step for specifying behavior. (Note: “Asynchrony” refers to various ideas, so always ask what people mean when they use the term.)

Idealized Definitions Ideally we would define:

```
type time = float
type 'a behavior = time -> 'a (* Arbitrary function *)
type 'a event = ('a, time) stream (* Increasing time instants *)
```

Forcing a lazy list (stream) of events would wait until an event arrives. But behaviors need to react to external events:

```
type user_action =
| Key of char * bool
| Button of int * int * bool * bool
| MouseMove of int * int
| Resize of int * int

type 'a behavior = user_action event -> time -> 'a
```

Scanning through an event list since the beginning of time until current time, each time we evaluate a behavior, is very wasteful with respect to time and space. Producing a stream of behaviors for the stream of time allows us to forget about events already in the past:

```
type 'a behavior =
  user_action event -> time stream -> 'a stream
```

The next optimization is to pair user actions with sampling times:

```
type 'a behavior =
  (user_action option * time) stream -> 'a stream
```

The `None` action corresponds to sampling time when nothing happens.

Turning behaviors and events from functions of time into input-output streams is similar to optimizing intersection of ordered lists from $O(mn)$ to $O(m + n)$ time.

Now we can in turn define events in terms of behaviors:

```
type 'a event = 'a option behavior
```

although it betrays the discrete character of events (happening at points in time rather than varying over intervals of time).

We have gotten very close to *stream processing* as discussed in Chapter 7. Recall the incremental pretty-printing example that can “react” to more input. Stream

combinators, *fork* from the exercises for Chapter 7, and a corresponding *merge*, turn stream processing into *synchronous discrete reactive programming*.

Behaviors as Monads Behaviors are monadic (but see the next point) – in the original specification:

```
type 'a behavior = time -> 'a

val return : 'a -> 'a behavior
let return a = fun _ -> a

val bind : 'a behavior -> ('a -> 'b behavior) -> 'b behavior
let bind a f = fun t -> f (a t) t
```

As we have seen with changeables, we mostly use lifting. In the Haskell world we would call behaviors *applicative*. To build our own lifters in any monad:

```
val ap : ('a -> 'b) monad -> 'a monad -> 'b monad
let ap fm am = perform
  f <-> fm;
  a <-> am;
  return (f a)
```

Note that for changeables, the naive implementation above will introduce unnecessary dependencies. Monadic libraries for *incremental computing* or FRP should provide optimized variants if needed. Compare with `parallel` for concurrent computing.

Converting Between Events and Behaviors Going from events to behaviors, `until` and `switch` have type:

```
'a behavior -> 'a behavior event -> 'a behavior
```

while `step` has type:

```
'a -> 'a event -> 'a behavior
```

- `until b es` behaves as `b` until the first event in `es`, then behaves as the behavior in that event
- `switch b es` behaves as the behavior from the last event in `es` prior to current time, if any, otherwise as `b`
- `step a b` starts with behavior returning `a` and then switches to returning the value of the last event in `b` (prior to current time) – a *step function*

We will use “*signal*” to refer to a behavior or an event. Note that often “*signal*” is used to mean what we call behavior (check terminology when looking at a new FRP library).

10.5 Reactivity by Stream Processing

The stream processing infrastructure should be familiar from earlier chapters:

```
type 'a stream = 'a stream_Lazy.t
and 'a stream_ = Cons of 'a * 'a stream

let rec lmap f l = lazy (
  let Cons (x, xs) = Lazy.force l in
  Cons (f x, lmap f xs))

let rec liter (f : 'a -> unit) (l : 'a stream) : unit =
  let Cons (x, xs) = Lazy.force l in
  f x; liter f xs

let rec lmap2 f xs ys = lazy (
  let Cons (x, xs) = Lazy.force xs in
  let Cons (y, ys) = Lazy.force ys in
  Cons (f x y, lmap2 f xs ys))

let rec lmap3 f xs ys zs = lazy (
  let Cons (x, xs) = Lazy.force xs in
  let Cons (y, ys) = Lazy.force ys in
  let Cons (z, zs) = Lazy.force zs in
  Cons (f x y z, lmap3 f xs ys zs))

let rec lfold acc f (l : 'a stream) = lazy (
  let Cons (x, xs) = Lazy.force l in (* Fold a function over the stream *)
  let acc = f acc x in (* producing a stream of partial results *)
  Cons (acc, lfold acc f xs))
```

Since a behavior is a function of user actions and sample times, we need to ensure that only one stream is created for the actual input stream:

```
type ('a, 'b) memo1 =
  {memo_f : 'a -> 'b; mutable memo_r : ('a * 'b) option}

let memo1 f = {memo_f = f; memo_r = None}

let memo1_app f x =
  match f.memo_r with
  | Some (y, res) when x == y -> res (* Physical equality is OK -- *)
  | _ -> (* external input is "physically" unique *)
    let res = f.memo_f x in (* While debugging, we can monitor *)
    f.memo_r <- Some (x, res); (* whether f.memo_r = None before *)
    res
```

```

let ($) = memo1_app

type 'a behavior =
  ((user_action option * time) stream, 'a stream) memo1

Building Complex Behaviors The monadic/applicative functions to build complex behaviors. If you do not provide type annotations in .ml files, work together with an .mli file to catch problems early. You can later add more type annotations as needed to find out what is wrong.

let returnB x : 'a behavior =
  let rec xs = lazy (Cons (x, xs)) in
    memo1 (fun _ -> xs)

let (* !*) = returnB

let liftB f fb = memo1 (fun uts -> lmap f (fb $ uts))

let liftB2 f fb1 fb2 = memo1
  (fun uts -> lmap2 f (fb1 $ uts) (fb2 $ uts))

let liftB3 f fb1 fb2 fb3 = memo1
  (fun uts -> lmap3 f (fb1 $ uts) (fb2 $ uts) (fb3 $ uts))

let liftE f (fe : 'a event) : 'b event = memo1
  (fun uts -> lmap
    (function Some e -> Some (f e) | None -> None)
    (fe $ uts))

let (=>>) fe f = liftE f fe
let (->>) e v = e =>> fun _ -> v

Creating events out of behaviors:

let whileB (fb : bool behavior) : unit event =
  memo1 (fun uts ->
    lmap (function true -> Some () | false -> None)
    (fb $ uts))

let unique fe : 'a event =
  memo1 (fun uts ->
    let xs = fe $ uts in
    lmap2 (fun x y -> if x = y then None else y)
    (lazy (Cons (None, xs))) xs)

let whenB fb =
  memo1 (fun uts -> unique (whileB fb) $ uts)

```

```

let snapshot fe fb : ('a * 'b) event =
  memo1 (fun uts -> lmap2
    (fun x -> function Some y -> Some (y, x) | None -> None)
    (fb $ uts) (fe $ uts))

```

Creating behaviors out of events:

```

let step acc fe = (* The step function: value of last event *)
  memo1 (fun uts -> lfold acc
    (fun acc -> function None -> acc | Some v -> v)
    (fe $ uts))

```

```

let step_accum acc ff = (* Transform a value by a series of functions *)
  memo1 (fun uts ->
    lfold acc (fun acc -> function
      | None -> acc | Some f -> f acc)
    (ff $ uts))

```

To numerically integrate a behavior, we need to access the sampling times:

```

let integral fb =
  let rec loop t0 acc uts bs =
    let Cons ((_, t1), uts) = Lazy.force uts in
    let Cons (b, bs) = Lazy.force bs in
    let acc = acc +. (t1 -. t0) *. b in (* b = fb(t1), acc approx integral up to t0 *)
    Cons (acc, lazy (loop t1 acc uts bs)) in
  memo1 (fun uts -> lazy (
    let Cons ((_, t), uts') = Lazy.force uts in
    Cons (0., lazy (loop t 0. uts' (fb $ uts)))))

```

In our *paddle game* example, we paradoxically express position and velocity in mutually recursive manner. The trick is the same as in Chapter 7 – integration introduces one step of delay.

User actions:

```

let lbp : unit event =
  memo1 (fun uts -> lmap
    (function Some(Button(_, _)), _ -> Some() | _ -> None)
    uts)

let mm : (int * int) event =
  memo1 (fun uts -> lmap
    (function Some(MouseMove(x, y)), _ -> Some(x, y) | _ -> None)
    uts)

let screen : (int * int) event =
  memo1 (fun uts -> lmap

```

```

(function Some(Resize(x, y)), _ -> Some(x, y) | _ -> None)
uts)

let mouse_x : int behavior = step 0 (liftE fst mm)
let mouse_y : int behavior = step 0 (liftE snd mm)
let width : int behavior = step 640 (liftE fst screen)
let height : int behavior = step 512 (liftE snd screen)

```

The Paddle Game Example A *scene graph* is a data structure that represents a “world” which can be drawn on screen:

```

type scene =
| Rect of int * int * int * int (* position, width, height *)
| Circle of int * int * int (* position, radius *)
| Group of scene list
| Color of Graphics.color * scene (* color of subscene objects *)
| Translate of float * float * scene (* additional offset of origin *)

```

Drawing a scene explains what we mean above:

```

let draw sc =
  let f2i = int_of_float in
  let open Graphics in
  let rec aux t_x t_y = function (* Accumulate translations *)
    | Rect (x, y, w, h) ->
      fill_rect (f2i t_x + x) (f2i t_y + y) w h
    | Circle (x, y, r) ->
      fill_circle (f2i t_x + x) (f2i t_y + y) r
    | Group scs ->
      List.iter (aux t_x t_y) scs
    | Color (c, sc) -> set_color c; aux t_x t_y sc (* Set color for sc objects *)
    | Translate (x, y, sc) -> aux (t_x +. x) (t_y +. y) sc in
  clear_graph (); (* "Fast and clean" removing of previous picture *)
  aux 0. 0. sc;
  synchronize () (* Synchronize the double buffer -- avoiding flickering *)

```

An animation is a scene behavior. To animate it we need to create the input stream: the user actions and sampling times stream. We could abstract away drawing from time sampling in `reactimate`, asking for (i.e. passing as argument) a producer of user actions and a consumer of scene graphs (like `draw`).

General-purpose behavior operators:

```

let (+*) = liftB2 (+)
let (-*) = liftB2 (-)
let (***) = liftB2 (*)
let (/*) = liftB2 (/)
let (&&*) = liftB2 (&&)

```

```

let (||*) = liftB2 (||)
let (<*) = liftB2 (<)
let (>*) = liftB2 (>)

```

The walls are drawn on left, top and right borders of the window:

```

let walls =
  liftB2 (fun w h -> Color (Graphics.blue, Group
    [Rect (0, 0, 20, h-1); Rect (0, h-21, w-1, 20);
     Rect (w-21, 0, 20, h-1)]))
  width height

```

The paddle is tied to the mouse at the bottom border of the window:

```

let paddle = liftB (fun mx ->
  Color (Graphics.black, Rect (mx, 0, 50, 10))) mouse_x

```

The ball has a velocity in pixels per second. It bounces from the walls, which is hard-coded in terms of distance from window borders. Unfortunately OCaml, being an eager language, does not let us encode recursive behaviors in an elegant way. We need to unpack behaviors and events as functions of the input stream:

- `xbounce ->> (~-.)` event is just the negation function happening at each horizontal bounce
- `step_accum vel (xbounce ->> (~-.))` behavior is `vel` value changing sign at each horizontal bounce
- `liftB int_of_float (integral xlabel) ** width /* !*2` – first integrate velocity, then truncate it to integers and offset to the middle of the window
- `whenB ((xpos >* width -* !*27) ||* (xpos <* !*27))` – issue an event the first time the position exceeds the bounds. This ensures there are no further bouncings until the ball moves out of the walls

10.6 Reactivity by Incremental Computing

In *Froc*, behaviors and events are both implemented as changeables but only behaviors persist; events are “instantaneous.” Behaviors are composed out of constants and prior events, capturing the “changeable” aspect. Events capture the “writeable” aspect – after their values are propagated, the values are removed. Events and behaviors are collectively called *signals*.

Froc does not represent time, and provides the function `changes : 'a behavior -> 'a event`, which violates the continuous semantics we introduced before. It breaks the illusion that behaviors vary continuously rather than at discrete points in time. But it avoids the need to synchronize global time samples with events in the system. It is “less continuous but more dense.”

Sending an event – `send` – starts an *update cycle*. Signals cannot call `send`, but can `send_deferred` which will send an event in the next cycle. Things

that happen in the same update cycle are *simultaneous*. Events are removed (detached from dependency graph) after an update cycle.

Froc provides the `fix_b`, `fix_e` functions to define signals recursively. Current value refers to value from previous update cycle, and defers next recursive step to next cycle, until convergence.

Update cycles can happen “back-to-back” via `send_deferred` and `fix_b`, `fix_e`, or can be invoked from outside *Froc* by sending events at arbitrary times. With a `time` behavior that holds a `clock` event value, events from “back-to-back” update cycles can be at the same clock time although not simultaneous in this sense. Update cycles prevent *glitches*, where outdated signal is used e.g. to issue an event.

Pure vs. Impure Style A behavior is written in *pure style* when its definition does not use `send`, `send_deferred`, `notify_e`, `notify_b` and `sample`:

- `sample`, `notify_e`, `notify_b` are used from outside the behavior (from its “environment”) analogously to observing the result of a function
- `send`, `send_deferred` are used from outside analogously to providing input to a function

When writing in impure style we need to remember to refer from somewhere to all the pieces of our behavior, otherwise the unrefereed parts will be **garbage collected** breaking the behavior. A value is referred to when it has a name in the global environment or is part of a bigger value that is referred to (for example it is stored somewhere). Signals can be referred to by being part of the dependency graph, but also by any of the more general ways.

Reimplementing the Paddle Game Example Rather than following our incremental computing example (a scene with changeable parts), we follow our FRP example: a scene behavior.

First we introduce time:

```
open Froc
let clock, tick = make_event ()
let time = hold (Unix.gettimeofday ()) clock
```

Next we define integration:

```
let integral fb =
  let aux (sum, t0) t1 =
    sum +. (t1 -. t0) *. sample fb, t1 in
  collect_b aux (0., sample time) clock
```

For convenience, the integral remembers the current upper limit of integration. It will be useful to get the integer part:

```
let integ_res fb =
```

```
lift (fun (v, _) -> int_of_float v) (integral fb)
```

We can also define integration in pure style:

```
let pair fa fb = lift2 (fun x y -> x, y) fa fb

let integral_nice fb =
  let samples = changes (pair fb time) in
  let aux (sum, t0) (fv, t1) =
    sum +. (t1 -. t0) *. fv, t1 in
  collect_b aux (0., sample time) samples
```

The initial value (0., sample time) is not “inside” the behavior so `sample` here does not spoil the pure style.

10.7 Direct Control

Real-world behaviors often are *state machines*, going through several stages. We do not have declarative means for it yet. For example, consider baking recipes: 1. *Preheat the oven.* 2. *Put flour, sugar, eggs into a bowl.* 3. *Spoon the mixture.* etc.

We want a *flow* to be able to proceed through events: when the first event arrives we remember its result and wait for the next event, disregarding any further arrivals of the first event! Therefore *Froc* constructs like mapping an event (`map`) or attaching a notification to a behavior change (`bind b1 (fun v1 -> notify_b ~now:false b2 (fun v2 -> ...))`) will not work.

We also want to be able to repeat or loop a flow, but starting from the notification of the first event that happens after the notification of the last event.

`next e` is an event propagating only the first occurrence of `e`. This will be the basis of our `await` function.

The whole flow should be cancellable from outside at any time.

A flow is a kind of a *lightweight thread* as in the end of Chapter 8; we will make it a monad. It only “stores” a non-unit value when it `awaits` an event. But it has a primitive to `emit` values. We actually implement *coarse-grained* threads (Chapter 8 exercise 11), with `await` in the role of `suspend`.

We build a module `Flow` with monadic type `('a, 'b) flow` “storing” `'b` and emitting `'a`:

```
type ('a, 'b) flow
type cancellable (* A handle to cancel a flow (stop further computation) *)
val noop_flow : ('a, unit) flow (* Same as return () *)
val return : 'b -> ('a, 'b) flow (* Completed flow *)
val await : 'b Froc.event -> ('a, 'b) flow (* Wait and store event: *)
val bind : (* the principled way to input *)
  ('a, 'b) flow -> ('b -> ('a, 'c) flow) -> ('a, 'c) flow
```

```

val emit : 'a -> ('a, unit) flow (* The principled way to output *)
val cancel : cancellable -> unit
val repeat : (* Loop the given flow and store the stop event *)
    ?until:'a Froc.event -> ('b, unit) flow -> ('b, 'a) flow
val event_flow :
    ('a, unit) flow -> 'a Froc.event * cancellable
val behavior_flow : (* The initial value of a behavior and a flow to update it *)
    'a -> ('a, unit) flow -> 'a Froc.behavior * cancellable
val is_cancelled : cancellable -> bool

```

Implementation Details We follow our (or *Lwt*) implementation of lightweight threads, adapting it to the need of cancelling flows:

```

module F = Froc
type 'a result =
  | Return of 'a (* Notifications to cancel when cancelled *)
  | Sleep of ('a -> unit) list * F.cancel ref list
  | Cancelled
  | Link of 'a state
and 'a state = {mutable state : 'a result}
type cancellable = unit state

```

Functions `find`, `wakeup`, `connect` are as in Chapter 8 (but connecting to cancelled thread cancels the other thread).

Our monad is actually a reader monad over the result state. The reader supplies the `emit` function:

```
type ('a, 'b) flow = ('a -> unit) -> 'b state
```

The `return` and `bind` functions are as in our lightweight threads, but we need to handle cancelled flows: for `m = bind a b`, if `a` is cancelled then `m` is cancelled, and if `m` is cancelled then we do not wake up `b`:

```

let waiter x =
  if not (is_cancelled m)
  then connect m (b x emit) in
  ...

```

`await` is implemented like `next`, but it wakes up a flow:

```

let await t = fun emit -
  let c = ref F.no_cancel in
  let m = {state = Sleep ([] , [c])} in
  c :=
  F.notify_e_cancel t begin fun r ->
    F.cancel !c;
    c := F.no_cancel;
    wakeup m r

```

```
    end;
```

```
m
```

`repeat` attaches the whole loop as a waiter for the loop body.

Example: Drawing Shapes The scene is a list of shapes, the first shape is open:

```
type scene = (int * int) list list

let draw sc =
  let open Graphics in
  clear_graph ();
  (match sc with
  | [] -> ()
  | opn :: cld ->
    draw_poly_line (Array.of_list opn);
    List.iter (fill_poly -| Array.of_list) cld);
  synchronize ()
```

We build a flow and turn it into a behavior to animate:

```
let painter =
  let cld = ref [] in (* Global state of painter *)
  repeat (perform
    await mbutton_pressed; (* Start when button down *)
    let opn = ref [] in
    repeat (perform
      mpos <- await mouse_move; (* Add next position to line *)
      emit (opn := mpos :: !opn; !opn :: !cld))
      ~until:mbutton_released; (* Start new shape *)
      emit (cld := !opn :: !cld; opn := []; [] :: !cld))

  let painter, cancel_painter = behavior_flow [] painter
  let () = reactivate painter
```

Flows and State Global state and thread-local state can be used with lightweight threads, but pay attention to semantics – which computations are inside the monad and which while building the initial monadic value.

Side effects hidden in `return` and `emit` arguments are not inside the monad. For example, if in the “first line” of a loop effects are executed only at the start of the loop – but if after bind (“below first line” of a loop), at each step of the loop:

```
let f =
  repeat (perform
    emit (Printf.printf "[0]\n%!";
    () <- await aas;
```

```

emit (Printf.printf "[1]\n%!" ; '1');
() <-- await bs;
emit (Printf.printf "[2]\n%!" ; '2');
() <-- await cs;
emit (Printf.printf "[3]\n%!" ; '3');
() <-- await ds;
emit (Printf.printf "[4]\n%!" ; '4'))

let e, cancel_e = event_flow f
let () =
  F.notify_e e (fun c -> Printf.printf "flow: %c\n%" c);
  Printf.printf "notification installed\n%!"

let () =
  F.send a (); F.send b (); F.send c (); F.send d ();
  F.send a (); F.send b (); F.send c (); F.send d ()

```

The output demonstrates the execution order: - [0] – Only printed once, when building the loop - notification installed – Only installed after the first flow event sent - event: a – Event notification - [1] – Second emit computed after first await returns - flow: 1 – Emitted signal - ... and so on through the loop iterations

10.8 Graphical User Interfaces

In-depth discussion of GUIs is beyond the scope of this course. We only cover what is needed for an example reactive program with direct control.

We demonstrate two libraries: *LablTk* based on optional labelled arguments (discussed in Chapter 2 exercise 2) and polymorphic variants, and *LablGTk* additionally based on objects. We will learn more about objects and polymorphic variants in the next chapter.

Calculator Flow We represent the mechanics of the calculator directly as a flow:

```

let digits, digit = F.make_event ()
let ops, op = F.make_event ()
let dots, dot = F.make_event ()

let calc =
  (* We need two state variables for two arguments of calculation *)
  let f = ref (fun x -> x) and now = ref 0.0 in  (* but we *)
  repeat (perform  (* remember the older argument in partial application *)
    op <-- repeat
    (perform  (* Enter the digits of a number (on later turns *)
      d <-- await digits;  (* starting from the second digit *) *)

```

```

        emit (now := 10. *. !now +. d; !now))
~until:ops; (* until operator button is pressed *)
emit (now := !f !now; f := op !now; !now);
(* Compute the result and "store away" the operator *)
d <- repeat
(perform op <- await ops; return (f := op !now))
~until:digits; (* The user can pick a different operator *)
emit (now := d; !now)) (* Reset the state to a new number *)

let calc_e, cancel_calc = event_flow calc (* Notifies display update *)

```

Tk: LablTk Widget toolkit *Tk* is known from the *Tcl* language.

Layout of the calculator – common across GUIs:

```

let layout = []
[|"7", `Di 7.; "8", `Di 8.; "9", `Di 9.; "+", `O (+.)|];
[|"4", `Di 4.; "5", `Di 5.; "6", `Di 6.; "-", `O (-.)|];
[|"1", `Di 1.; "2", `Di 2.; "3", `Di 3.; "*", `O (*.)|];
[|"0", `Di 0.; ".", `Dot; "=", `O sk; "/", `O (/.)|]
[]
```

Key concepts: - Every *widget* (window gadget) has a parent in which it is located - *Buttons* have action associated with pressing them, *labels* just provide information, *entries* (aka. *edit* fields) are for entering info from keyboard - Actions are *callback* functions passed as the *~command* argument - *Frames* in *Tk* group widgets - The parent is sent as last argument, after optional labelled arguments

```

let top = Tk.openTk ()

let btn_frame =
  Frame.create ~relief:`Groove ~borderwidth:2 top

let buttons =
  Array.map (Array.map (function
    | text, `Dot ->
      Button.create ~text
      ~command:(fun () -> F.send dot ()) btn_frame
    | text, `Di d ->
      Button.create ~text
      ~command:(fun () -> F.send digit d) btn_frame
    | text, `O f ->
      Button.create ~text
      ~command:(fun () -> F.send op f) btn_frame)) layout

let result = Label.create ~text:"0" ~relief:`Sunken top
```

GUI toolkits have layout algorithms, so we only need to tell which widgets hang together and whether they should fill all available space etc. – via `pack`, or `grid` for “rectangular” organization:

- `~fill`: the allocated space in ``X`, ``Y`, ``Both` or ``None` axes
- `~expand`: maximally how much space is allocated or only as needed
- `~anchor`: allows to glue a widget in particular direction (``Center`, ``E`, ``Ne` etc.)
- The `grid` packing flexibility: `~columnspan` and `~rowspan`
- `configure` functions accept the same arguments as `create` but change existing widgets

```
let () =
  Wm.title_set top "Calculator";
  Tk.pack [result] ~side:`Top ~fill:`X;
  Tk.pack [btn_frame] ~side:`Bottom ~expand:true;
  Array.iteri (fun column -> Array.iteri (fun row button ->
    Tk.grid ~column ~row [button])) buttons;
  Wm.geometry_set top "200x200";
  F.notify_e calc_e
  (fun now ->
    Label.configure ~text:(string_of_float now) result);
  Tk.mainloop ()
```

GTk+: *LablGtk* *LablGtk* is built as an object-oriented layer over a low-level layer of functions interfacing with the *Gtk+* library, which is written in *C*.

In OCaml, object fields are only visible to object methods, and methods are called with `#` syntax, e.g. `window#show ()`.

The interaction with the application is reactive: - Our events are called signals in *Gtk+* - Registering a notification is called connecting a signal handler, e.g. `button#connect#clicked ~callback:hello` which takes `~callback:(unit -> unit)` and returns `GtkSignal.id` - As with *Froc* notifications, multiple handlers can be attached - *Gtk+* events are a subclass of signals related to more specific window events, e.g. `window#event#connect#delete ~callback:delete_event` - *Gtk+* event callbacks take more info: `~callback:(event -> unit)` for some type `event`

Automatic layout (aka. packing) seems less sophisticated than in *Tk*: - only horizontal and vertical boxes - therefore `~fill` is binary and `~anchor` is replaced by `~from `START` or ``END`

Automatic grid layout is called `table`: - `~fill` and `~expand` take ``X`, ``Y`, ``BOTH`, ``NONE`

The `coerce` method casts the type of the object (in *Tk* there is `coe` function). Labels do not have a dedicated module. Widgets have setter methods `widget#set_X` (instead of a single `configure` function in *Tk*).

Setup:

```
let _ = GtkMain.Main.init ()
let window =
    GWindow.window ~width:200 ~height:200 ~title:"Calculator" ()
let top = GPack.vbox ~packing:window#add ()
let result = GMisc.label ~text:"0" ~packing:top#add ()
let btn_frame =
    GPack.table ~rows:(Array.length layout)
    ~columns:(Array.length layout.(0)) ~packing:top#add ()
```

Button actions:

```
let buttons =
    Array.map (Array.map (function
        | label, `Dot ->
            let b = GButton.button ~label () in
            let _ = b#connect#clicked
                ~callback:(fun () -> F.send dot ()) in b
        | label, `Di d ->
            let b = GButton.button ~label () in
            let _ = b#connect#clicked
                ~callback:(fun () -> F.send digit d) in b
        | label, `O f ->
            let b = GButton.button ~label () in
            let _ = b#connect#clicked
                ~callback:(fun () -> F.send op f) in b)) layout
```

Button layout, result notification, start application:

```
let delete_event _ = GMain.Main.quit (); false

let () =
    let _ = window#event#connect#delete ~callback:delete_event in
    Array.iteri (fun column -> Array.iteri (fun row button ->
        btn_frame#attach ~left:column ~top:row
        ~fill:`BOTH ~expand:`BOTH (button#coerce))
    ) buttons;
    F.notify_e calc_e
    (fun now -> result#set_label (string_of_float now));
    window#show ();
    GMain.Main.main ()
```

10.9 Exercises

Exercise 1: Introduce operators $-$, $/$ into the context rewriting “pull out subexpression” example. Remember that they are not commutative.

Exercise 2: Add to the *paddle game* example: 1. game restart, 2. score keeping, 3. game quitting (in more-or-less elegant way).

Exercise 3: Our numerical integration function roughly corresponds to the rectangle rule. Modify the rule and write a test for the accuracy of: 1. the trapezoidal rule; 2. the Simpson's rule. See http://en.wikipedia.org/wiki/Simpson%27s_rule

Exercise 4: Explain the recursive behavior of integration: 1. In *paddle game* implemented by stream processing (`Lec10b.ml`), do we look at past velocity to determine current position, at past position to determine current velocity, both, or neither? 2. What is the difference between `integral` and `integral_nice` in `Lec10c.ml`, what happens when we replace the former with the latter in the `pbal` function? How about after rewriting `pbal` into pure style as in the following exercise?

Exercise 5: Reimplement the *Froc* based paddle ball example in a pure style: rewrite the `pbal` function to not use `notify_e`.

Exercise 6: Our implementation of flows is a bit heavy. One alternative approach is to use continuations, as in `Scala.React`. OCaml has a continuations library `Delimcc`; for how it can cooperate with *Froc*, see <http://ambassadortothecomputers.blogspot.com/2010/08/mixing-monadic-and-direct-style-code.html>

Exercise 7: Implement `parallel` for flows, retaining coarse-grained implementation and using the event queue from *Froc* somehow (instead of introducing a new job queue).

Exercise 8: Add quitting, e.g. via a 'q' key press, to the *painter* example. Use the `is_cancelled` function.

Exercise 9: Our calculator example is not finished. Implement entering decimal fractions: add handling of the `dots` event.

Exercise 10: The Flow module has reader monad functions that have not been discussed in this chapter:

```
let local f m = fun emit -> m (fun x -> emit (f x))
let local_opt f m = fun emit ->
  m (fun x -> match f x with None -> () | Some y -> emit y)

val local : ('a -> 'b) -> ('a, 'c) flow -> ('b, 'c) flow
val local_opt : ('a -> 'b option) -> ('a, 'c) flow -> ('b, 'c) flow
```

Implement an example that uses this compositionality-increasing capability.

Chapter 11: The Expression Problem

This chapter explores **the expression problem**, a classic challenge in software engineering that addresses how to design systems that can be extended with both new data variants and new operations without modifying existing code,

while maintaining static type safety. We examine multiple approaches in OCaml, ranging from algebraic data types through object-oriented programming to polymorphic variants with recursive modules. The chapter concludes with an application to parser combinators and dynamic code loading.

11.1 The Expression Problem: Definition

The **Expression Problem** concerns the design of an implementation for expressions where:

- **Datatype extensibility:** New variants of expressions can be added
- **Functional extensibility:** New operations on expressions can be added

By *extensibility* we mean three conditions:

1. **Code-level modularization:** The new datatype variants and new operations are in separate files
2. **Separate compilation:** The files can be compiled and distributed separately
3. **Static type safety:** We do not lose type checking help and guarantees

The name comes from a classic example: extending a language of expressions with new constructs. Consider two sub-languages:

- **Lambda calculus:** variables `Var`, λ -abstractions `Abs`, function applications `App`
- **Arithmetic:** variables `Var`, constants `Num`, addition `Add`, multiplication `Mult`

And operations we want to support:

- Evaluation `eval`
- Pretty-printing to strings `string_of`
- Free variables computation `free_vars`

The challenge is to combine these sub-languages and add new operations without breaking existing code or sacrificing type safety.

References

- Ralf Lammel lectures on MSDN's Channel 9: The Expression Problem, Haskell's Type Classes
- The book *Developing Applications with Objective Caml*: Comparison of Modules and Objects, Extending Components
- *Real World OCaml*: Chapter 11: Objects, Chapter 12: Classes
- Jacques Garrigue's Code reuse through polymorphic variants, and Recursive Modules for Programming with Keiko Nakata
- Extensible variant types
- Graham Hutton's and Erik Meijer's Monadic Parser Combinators

11.2 Functional Programming Non-Solution: Ordinary Algebraic Datatypes

Pattern matching makes **functional extensibility** easy in functional programming. However, ensuring **datatype extensibility** is complicated when using standard variant types.

For brevity, we place examples in a single file, but the component type and function definitions are not mutually recursive, so they can be put in separate modules.

Non-solution penalty points:

- Functions implemented for a broader language (e.g., `lexpr_t`) cannot be used with a value from a narrower language (e.g., `expr_t`)
- Significant memory (and some time) overhead due to *tagging*: the work of `wrap` and `unwrap` functions, adding tags such as `Lambda` and `Expr`
- Some code bloat due to tagging. For example, deep pattern matching needs to be manually unrolled and interspersed with calls to `unwrap`

Verdict: Non-solution, but better than extensible variant types-based approach and direct OOP approach.

Here is the implementation:

```
type var = string (* Variables constitute a sub-language of its own *)
(* We treat this sub-language slightly differently -- no need for a dedicated type*)

let eval_var wrap sub (s : var) =
  try List.assoc s sub with Not_found -> wrap s

type 'a lambda = (* Here we define the sub-language of lambda-expressions *)
  VarL of var | Abs of string * 'a | App of 'a * 'a

(* During evaluation, we need to freshen variables to avoid capture *)
(* (mistaking distinct variables with the same name) *)
let gensym = let n = ref 0 in fun () -> incr n; "_" ^ string_of_int !n

let eval_lambda eval_rec wrap unwrap subst e =
  match unwrap e with (* Alternatively, unwrapping could use an exception *)
  | Some (VarL v) -> eval_var (fun v -> wrap (VarL v)) subst v
  | Some (App (l1, l2)) -> (* but we use the option type as it is safer *)
    let l1' = eval_rec subst l1 (* and more flexible in this context *)
    and l2' = eval_rec subst l2 in (* Recursive processing function returns expression *)
    (match unwrap l1' with (* of the completed language, we need *)
     | Some (Abs (s, body)) -> (* to unwrap it into the current sub-language *)
       eval_rec [s, l2'] body (* The recursive call is already wrapped *)
     | _ -> wrap (App (l1', l2'))) (* Wrap into the completed language *)
  | Some (Abs (s, l1)) ->
```

```

let s' = gensym () in (* Rename variable to avoid capture (alpha-equivalence) *)
wrap (Abs (s', eval_rec ((s, wrap (VarL s'))::subst) 11))
| None -> e (* Falling-through when not in the current sub-language *)

type lambda_t = Lambda_t of lambda_t lambda (* Defining lambda-expressions as the completed*)

let rec eval1 subst = (* and the corresponding eval function *)
  eval_lambda eval1
  (fun e -> Lambda_t e) (fun (Lambda_t e) -> Some e) subst

Now we define the arithmetic sub-language:

type 'a expr = (* The sub-language of arithmetic expressions *)
  VarE of var | Num of int | Add of 'a * 'a | Mult of 'a * 'a

let eval_expr eval_rec wrap unwrap subst e =
  match unwrap e with
  | Some (Num _) -> e
  | Some (VarE v) ->
    eval_var (fun x -> wrap (VarE x)) subst v
  | Some (Add (m, n)) ->
    let m' = eval_rec subst m
    and n' = eval_rec subst n in
    (match unwrap m', unwrap n' with (* Unwrapping to check if the subexpressions *)
     | Some (Num m'), Some (Num n') -> (* got computed to values *)
       wrap (Num (m' + n'))
     | _ -> wrap (Add (m', n'))) (* Here m' and n' are wrapped *)
  | Some (Mult (m, n)) ->
    let m' = eval_rec subst m
    and n' = eval_rec subst n in
    (match unwrap m', unwrap n' with
     | Some (Num m'), Some (Num n') ->
       wrap (Num (m' * n'))
     | _ -> wrap (Mult (m', n'))))
  | None -> e

type expr_t = Expr_t of expr_t expr (* Defining arithmetic expressions as the completed*)

let rec eval2 subst = (* aka "tying the recursive knot" *)
  eval_expr eval2
  (fun e -> Expr_t e) (fun (Expr_t e) -> Some e) subst

Finally, we merge the two sub-languages:

type 'a lexpr = (* The language merging lambda-expressions and arithmetic expressions *)
  Lambda of 'a lambda | Expr of 'a expr (* can also be used as a sub-language for further*)

let eval_lexpr eval_rec wrap unwrap subst e =

```

```

eval_lambda eval_rec
  (fun e -> wrap (Lambda e))
  (fun e ->
    match unwrap e with
    | Some (Lambda e) -> Some e
    | _ -> None)
  subst
  (eval_expr eval_rec (* We use the "fall-through" property of eval_expr *)
    (fun e -> wrap (Expr e)) (* to combine the evaluators *)
    (fun e ->
      match unwrap e with
      | Some (Expr e) -> Some e
      | _ -> None)
    subst e)

type lexpr_t = LExpr_t of lexpr_t lexpr (* Tying the recursive knot one last time *)

let rec eval3 subst =
  eval_lexer eval3
  (fun e -> LExpr_t e)
  (fun (LEExpr_t e) -> Some e) subst

```

11.3 Lightweight FP Non-Solution: Extensible Variant Types

Exceptions have always formed an extensible variant type in OCaml, whose pattern matching is done using the `try...with` syntax. Since recently, new extensible variant types can be defined. This augments the normal function extensibility of FP with straightforward data extensibility.

Non-solution penalty points:

- Giving up exhaustivity checking, which is an important aspect of static type safety
- More natural with “single inheritance” extension chains, although merging is possible and demonstrated in our example
- Requires “tying the recursive knot” for functions

Verdict: Pleasant-looking, but the worst approach because of possible bugginess. Unless bug-proneness is not a concern, then the best approach.

```

type expr = .. (* This is how extensible variant types are defined *)

type var_name = string
type expr += Var of string (* We add a variant case *)

let eval_var sub = function
| Var s as v -> (try List.assoc s sub with Not_found -> v)
| e -> e

```

```

let gensym = let n = ref 0 in fun () -> incr n; "_" ^ string_of_int !n

type expr += Abs of string * expr | App of expr * expr
(* The sub-languages are not differentiated by types, a shortcoming of this non-solution *)

let eval_lambda eval_rec subst = function
| Var _ as v -> eval_var subst v
| App (l1, l2) ->
  let l2' = eval_rec subst l2 in
  (match eval_rec subst l1 with
  | Abs (s, body) ->
    eval_rec [s, l2'] body
  | l1' -> App (l1', l2'))
| Abs (s, l1) ->
  let s' = gensym () in
  Abs (s', eval_rec ((s, Var s')::subst) l1)
| e -> e

let freevars_lambda freevars_rec = function
| Var v -> [v]
| App (l1, l2) -> freevars_rec l1 @ freevars_rec l2
| Abs (s, l1) ->
  List.filter (fun v -> v <> s) (freevars_rec l1)
| _ -> []

let rec eval1 subst e = eval_lambda eval1 subst e
let rec freevars1 e = freevars_lambda freevars1 e

let test1 = App (Abs ("x", Var "x"), Var "y")
let e_test = eval1 [] test1
let fv_test = freevars1 test1

```

Now we extend with arithmetic:

```

type expr += Num of int | Add of expr * expr | Mult of expr * expr

let map_expr f = function
| Add (e1, e2) -> Add (f e1, f e2)
| Mult (e1, e2) -> Mult (f e1, f e2)
| e -> e

let eval_expr eval_rec subst e =
  match map_expr (eval_rec subst) e with
  | Add (Num m, Num n) -> Num (m + n)
  | Mult (Num m, Num n) -> Num (m * n)
  | (Num _ | Add _ | Mult _) as e -> e

```

```

| e -> e

let freevars_expr freevars_rec = function
| Num _ -> []
| Add (e1, e2) | Mult (e1, e2) -> freevars_rec e1 @ freevars_rec e2
| _ -> []

let rec eval2 subst e = eval_expr eval2 subst e
let rec freevars2 e = freevars_expr freevars2 e

let test2 = Add (Mult (Num 3, Var "x"), Num 1)
let e_test2 = eval2 [] test2
let fv_test2 = freevars2 test2

```

Merging the sub-languages:

```

let eval_expr eval_rec subst e =
  eval_expr eval_rec subst (eval_lambda eval_rec subst e)

let freevars_expr freevars_rec e =
  freevars_lambda freevars_rec e @ freevars_expr freevars_rec e

let rec eval3 subst e = eval_expr eval3 subst e
let rec freevars3 e = freevars_expr freevars3 e

let test3 =
  App (Abs ("x", Add (Mult (Num 3, Var "x"), Num 1)),
    Num 2)
let e_test3 = eval3 [] test3
let fv_test3 = freevars3 test3

```

11.4 Object-Oriented Programming: Subtyping

OCaml's **objects** are values, somewhat similar to records. Viewed from the outside, an OCaml object has only **methods**, identifying the code with which to respond to messages (method invocations). All methods are **late-bound**; the object determines what code is run (i.e., *virtual* in C++ parlance).

Subtyping determines if an object can be used in some context. OCaml has **structural subtyping**: the content of the types concerned decides if an object can be used. Parametric polymorphism can be used to infer if an object has the required methods.

```

let f x = x#m (* Method invocation: object#method *)
(* val f : < m : 'a; .. > -> 'a *)
(* Type polymorphic in two ways: 'a is the method type, *)
(* .. means that objects with more methods will be accepted *)

```

Methods are computed when they are invoked, even if they do not take arguments. We define objects inside `object...end` (compare: records `{...}`) using keywords:

- `method` for methods
- `val` for constant fields
- `val mutable` for mutable fields

Constructor arguments can often be used instead of constant fields:

```
let square w = object
  method area = float_of_int (w * w)
  method width = w
end
```

Subtyping often needs to be explicit: we write `(object :> supertype)` or in more complex cases `(object : type :> supertype)`.

Technically speaking, subtyping in OCaml always is explicit, and *open types*, containing ..., use **row polymorphism** rather than subtyping.

```
let a = object method m = 7  method x = "a" end (* Toy example: object types *)
let b = object method m = 42 method y = "b" end (* share some but not all methods *)

(* let l = [a; b] -- Error: the exact types of the objects do not agree *)
(* Error: This expression has type <m : int; y : string>
   but an expression was expected of type <m : int; x : string>
   The second object type has no method y *)

let l = [(a :> <m : 'a>); (b :> <m : 'a>)] (* But the types share a supertype *)
(* val l : <m : int> list *)
```

11.5 Direct Object-Oriented Non-Solution

We can try to solve the expression problem using objects directly. However, adding new functionality still requires modifying old code, so this approach does not fully solve the expression problem.

Non-solution penalty points:

- No way to add functionality without modifying old code (in particular, the abstract class and all concrete classes)
- No deep pattern matching

Verdict: Non-solution, and probably the worst approach.

Here is an implementation using objects:

```
type var_name = string
```

```
let gensym = let n = ref 0 in fun () -> incr n; "_" ^ string_of_int !n
```

```

class virtual ['lang] evaluable =
object
  method virtual eval : (var_name * 'lang) list -> 'lang
  method virtual rename : var_name -> var_name -> 'lang
  method apply (_arg : 'lang)
    (fallback : unit -> 'lang) (_subst : (var_name * 'lang) list) =
    fallback ()
end

class ['lang] var (v : var_name) =
object (self)
  inherit ['lang] evaluable
  val v = v
  method eval subst =
    try List.assoc v subst with Not_found -> self
  method rename v1 v2 =
    if v = v1 then {< v = v2 >} else self
end

class ['lang] abs (v : var_name) (body : 'lang) =
object (self)
  inherit ['lang] evaluable
  val v = v
  val body = body
  method eval subst =
    let v' = gensym () in
    {< v = v'; body = (body#rename v v')#eval subst >}
  method rename v1 v2 =
    if v = v1 then self
    else {< body = body#rename v1 v2 >}
  method apply arg _ subst =
    body#eval ((v, arg)::subst)
end

class ['lang] app (f : 'lang) (arg : 'lang) =
object (self)
  inherit ['lang] evaluable
  val f = f
  val arg = arg
  method eval subst =
    let arg' = arg#eval subst in
    f#apply arg' (fun () -> {< f = f#eval subst; arg = arg' >}) subst
  method rename v1 v2 =
    {< f = f#rename v1 v2; arg = arg#rename v1 v2 >}
end

```

```

type evaluable_t = evaluable_t evaluable
let new_var1 v : evaluable_t = new var v
let new_abs1 v (body : evaluable_t) : evaluable_t = new abs v body
let new_app1 (arg1 : evaluable_t) (arg2 : evaluable_t) : evaluable_t =
  new app arg1 arg2

let test1 = new_app1 (new_abs1 "x" (new_var1 "x")) (new_var1 "y")
let e_test1 = test1#eval []

```

Extending with arithmetic requires additional mixins:

```

class virtual compute_mixin = object
  method compute : int option = None
end

class ['lang] var_c v = object
  inherit ['lang] var v
  inherit compute_mixin
end

class ['lang] abs_c v body = object
  inherit ['lang] abs v body
  inherit compute_mixin
end

class ['lang] app_c f arg = object
  inherit ['lang] app f arg
  inherit compute_mixin
end

class ['lang] num (i : int) =
object (self)
  inherit ['lang] evaluable
  val i = i
  method eval _subst = self
  method rename _ _ = self
  method compute = Some i
end

class virtual ['lang] operation
  (num_inst : int -> 'lang) (n1 : 'lang) (n2 : 'lang) =
object (self)
  inherit ['lang] evaluable
  val n1 = n1
  val n2 = n2
  method eval subst =
    let self' = {< n1 = n1#eval subst; n2 = n2#eval subst >} in

```

```

    match self'#compute with
    | Some i -> num_inst i
    | _ -> self'
  end

  class ['lang] add num_inst n1 n2 =
object (self)
  inherit ['lang] operation num_inst n1 n2
  method compute =
    match n1#compute, n2#compute with
    | Some i1, Some i2 -> Some (i1 + i2)
    | _ -> None
end

  class ['lang] mult num_inst n1 n2 =
object (self)
  inherit ['lang] operation num_inst n1 n2
  method compute =
    match n1#compute, n2#compute with
    | Some i1, Some i2 -> Some (i1 * i2)
    | _ -> None
end

  class virtual ['lang] computable =
object
  inherit ['lang] evaluable
  inherit compute_mixin
end

  type computable_t = computable_t computable
  let new_var2 v : computable_t = new var_c v
  let new_abs2 v (body : computable_t) : computable_t = new abs_c v body
  let new_app2 v (body : computable_t) : computable_t = new app_c v body
  let new_num2 i : computable_t = new num i
  let new_add2 (n1 : computable_t) (n2 : computable_t) : computable_t =
    new add new_num2 n1 n2
  let new_mult2 (n1 : computable_t) (n2 : computable_t) : computable_t =
    new mult new_num2 n1 n2

  let test2 =
    new_app2 (new_abs2 "x" (new_add2 (new_mult2 (new_num2 3) (new_var2 "x"))
                                (new_num2 1))))
               (new_num2 2)
  let e_test2 = test2#eval []

```

11.6 OOP Non-Solution: The Visitor Pattern

The **visitor pattern** is a design pattern that separates an algorithm from the object structure on which it operates. This allows adding new operations to existing object structures without modifying those structures.

Non-solution penalty points:

- Adding new functionality requires modifying old code (the abstract visitor class)
- No deep pattern matching
- Uses mutable state for returning results

Verdict: Poor solution, better than approaches we considered so far, and worse than approaches we consider next.

```
type 'visitor visitable = < accept : 'visitor -> unit >
(* The variants need be visitable *)
(* We store the computation as side effect because of the difficulty *)
(* to keep the visitor polymorphic but have the result type depend on the visitor *)

type var_name = string

class ['visitor] var (v : var_name) =
object (self)
    (* The 'visitor will determine the (sub)language *)
    (* to which a given var variant belongs *)
    method v = v
    method accept : 'visitor -> unit = (* The visitor pattern inverts the way *)
        fun visitor -> visitor#visitVar self (* pattern matching proceeds: the variant *)
end (* selects the computation *)
let new_var v = (new var v :> 'a visitable)

class ['visitor] abs (v : var_name) (body : 'visitor visitable) =
object (self)
    method v = v
    method body = body
    method accept : 'visitor -> unit =
        fun visitor -> visitor#visitAbs self
end
let new_abs v body = (new abs v body :> 'a visitable)

class ['visitor] app (f : 'visitor visitable) (arg : 'visitor visitable) =
object (self)
    method f = f
    method arg = arg
    method accept : 'visitor -> unit =
        fun visitor -> visitor#visitApp self
end
```

```

let new_app f arg = (new app f arg :> 'a visitable)

class virtual ['visitor] lambda_visit =
object
  method virtual visitVar : 'visitor var -> unit
  method virtual visitAbs : 'visitor abs -> unit
  method virtual visitApp : 'visitor app -> unit
end

let gensym = let n = ref 0 in fun () -> incr n; "_" ^ string_of_int !n

class ['visitor] eval_lambda
  (subst : (var_name * 'visitor visitable) list)
  (result : 'visitor visitable ref) =
object (self)
  inherit ['visitor] lambda_visit
  val mutable subst = subst
  val mutable beta_redex : (var_name * 'visitor visitable) option = None
  method visitVar var =
    beta_redex <- None;
    try result := List.assoc var#v subst
    with Not_found -> result := (var :> 'visitor visitable)
  method visitAbs abs =
    let v' = gensym () in
    let orig_subst = subst in
    subst <- (abs#v, new_var v')::subst;
    (abs#body)#accept self;
    let body' = !result in
    subst <- orig_subst;
    beta_redex <- Some (v', body');
    result := new_abs v' body'
  method visitApp app =
    app#arg#accept self;
    let arg' = !result in
    app#f#accept self;
    let f' = !result in
    match beta_redex with
    | Some (v', body') ->
        beta_redex <- None;
        let orig_subst = subst in
        subst <- (v', arg')::subst;
        body'#accept self;
        subst <- orig_subst
    | None -> result := new_app f' arg'
end

```

```

class ['visitor] freevars_lambda (result : var_name list ref) =
object (self)
  inherit ['visitor] lambda_visit
  method visitVar var =
    result := var#v :: !result
  method visitAbs abs =
    (abs#body)#accept self;
    result := List.filter (fun v' -> v' <> abs#v) !result
  method visitApp app =
    app#arg#accept self; app#f#accept self
end

type lambda_visit_t = lambda_visit_t lambda_visit
type lambda_t = lambda_visit_t visitable

let eval1 (e : lambda_t) subst : lambda_t =
  let result = ref (new_var "") in
  e#accept (new eval_lambda subst result :> lambda_visit_t);
  !result

let freevars1 (e : lambda_t) =
  let result = ref [] in
  e#accept (new freevars_lambda result);
  !result

let test1 =
  (new_app (new_abs "x" (new_var "x")) (new_var "y") :> lambda_t)
let e_test = eval1 test1 []
let fv_test = freevars1 test1

```

Extending with arithmetic expressions follows a similar pattern, and the merged language visitor inherits from both `lambda_visit` and `expr_visit`.

11.7 Polymorphic Variants

Polymorphic variants provide a flexible alternative to standard variants. They allow combining types from different sources without explicitly defining a common parent type.

Penalty points:

- Requires explicit type annotations more often
- Requires “tying the recursive knots” for types, e.g., `type lambda_t = lambda_t lambda`
- Some loss of type-level distinction between sub-languages

Verdict: A flexible solution, better than the previous approaches but still not perfect.

```

type var = [`Var of string]

let eval_var sub (`Var s as v : var) =
  try List.assoc s sub with Not_found -> v

type 'a lambda =
  [`Var of string | `Abs of string * 'a | `App of 'a * 'a]

let gensym = let n = ref 0 in fun () -> incr n; "_" ^ string_of_int !n

let eval_lambda eval_rec subst : 'a lambda -> 'a = function
  | #var as v -> eval_var subst v
  | `App (l1, l2) ->
    let l2' = eval_rec subst l2 in
    (match eval_rec subst l1 with
     | `Abs (s, body) ->
       eval_rec [s, l2'] body
     | l1' -> `App (l1', l2'))
  | `Abs (s, l1) ->
    let s' = gensym () in
    `Abs (s', eval_rec ((s, `Var s')::subst) l1)

let freevars_lambda freevars_rec : 'a lambda -> 'b = function
  | `Var v -> [v]
  | `App (l1, l2) -> freevars_rec l1 @ freevars_rec l2
  | `Abs (s, l1) ->
    List.filter (fun v -> v <> s) (freevars_rec l1)

type lambda_t = lambda_t lambda

let rec eval1 subst e : lambda_t = eval_lambda eval1 subst e
let rec freevars1 (e : lambda_t) = freevars_lambda freevars1 e

let test1 = (`App (`Abs ("x", `Var "x"), `Var "y") :> lambda_t)
let e_test = eval1 [] test1
let fv_test = freevars1 test1

```

The arithmetic expression sub-language:

```

type 'a expr =
  [`Var of string | `Num of int | `Add of 'a * 'a | `Mult of 'a * 'a]

let map_expr (f : _ -> 'a) : 'a expr -> 'a = function
  | #var as v -> v
  | `Num _ as n -> n
  | `Add (e1, e2) -> `Add (f e1, f e2)
  | `Mult (e1, e2) -> `Mult (f e1, f e2)

```

```

let eval_expr eval_rec subst (e : 'a expr) : 'a =
  match map_expr (eval_rec subst) e with
  | #var as v -> eval_var subst v
  | `Add (`Num m, `Num n) -> `Num (m + n)
  | `Mult (`Num m, `Num n) -> `Num (m * n)
  | e -> e

let freevars_expr freevars_rec : 'a expr -> 'b = function
  | `Var v -> [v]
  | `Num _ -> []
  | `Add (e1, e2) | `Mult (e1, e2) -> freevars_rec e1 @ freevars_rec e2

type expr_t = expr_t expr

let rec eval2 subst e : expr_t = eval_expr eval2 subst e
let rec freevars2 (e : expr_t) = freevars_expr freevars2 e

let test2 = (`Add (`Mult (`Num 3, `Var "x"), `Num 1) : expr_t)
let e_test2 = eval2 ["x", `Num 2] test2
let fv_test2 = freevars2 test2

```

Merging the sub-languages:

```

type 'a lexpr = ['a lambda | 'a expr]

let eval_lexpr eval_rec subst : 'a lexpr -> 'a = function
  | #lambda as x -> eval_lambda eval_rec subst x
  | #expr as x -> eval_expr eval_rec subst x

let freevars_lexpr freevars_rec : 'a lexpr -> 'b = function
  | #lambda as x -> freevars_lambda freevars_rec x
  | #expr as x -> freevars_expr freevars_rec x

type lexpr_t = lexpr_t lexpr

let rec eval3 subst e : lexpr_t = eval_lexpr eval3 subst e
let rec freevars3 (e : lexpr_t) = freevars_lexpr freevars3 e

let test3 =
  (`App (`Abs ("x", `Add (`Mult (`Num 3, `Var "x"), `Num 1)),
    `Num 2) : lexpr_t)
let e_test3 = eval3 [] test3
let fv_test3 = freevars3 test3
let e_old_test = eval3 [] (test2 :> lexpr_t)
let fv_old_test = freevars3 (test2 :> lexpr_t)

```

11.8 Polymorphic Variants with Recursive Modules

Using recursive modules, we can clean up the confusing or cluttering aspects of tying the recursive knots: type variables and recursive call arguments.

We need **private types**, which for objects and polymorphic variants means *private rows*. We can conceive of open row types, e.g., $[> \text{Int of int} | \text{'String of string}]$ as using a *row variable*, e.g., 'a':

$[\text{Int of int} | \text{'String of string} | 'a]$

and then of private row types as abstracting the row variable:

```
type 'row t = [`Int of int | `String of string | 'row]
```

But the actual formalization of private row types is more complex.

Penalty points:

- We still need to tie the recursive knots for types, for example `private [> 'a lambda] as 'a`
- There can be slight time costs due to the use of functors and dispatch on merging of sub-languages

Verdict: A clean solution, best place.

```
type var = [`Var of string]

let eval_var subst (`Var s as v : var) =
  try List.assoc s subst with Not_found -> v

type 'a lambda =
  [`Var of string | `Abs of string * 'a | `App of 'a * 'a]

module type Eval =
sig type exp val eval : (string * exp) list -> exp -> exp end

module LF(X : Eval with type exp = private [> 'a lambda] as 'a) =
struct
  type exp = X.exp lambda

  let gensym = let n = ref 0 in fun () -> incr n; "_" ^ string_of_int !n

  let eval subst : exp -> X.exp = function
    | #var as v -> eval_var subst v
    | `App (l1, l2) ->
        let l2' = X.eval subst l2 in
        (match X.eval subst l1 with
        | `Abs (s, body) ->
            X.eval [s, l2'] body
```

```

| `App (l1', l2') )
| `Abs (s, l1) -
let s' = gensym () in
`Abs (s', X.eval ((s, `Var s'))::subst) l1)
end
module rec Lambda : (Eval with type exp = Lambda.exp lambda) =
LF(Lambda)

module type FreeVars =
sig type exp val freevars : exp -> string list end

module LFVF(X : FreeVars with type exp = private [> 'a lambda] as 'a) =
struct
type exp = X.exp lambda

let freevars : exp -> 'b = function
| `Var v -> [v]
| `App (l1, l2) -> X.freevars l1 @ X.freevars l2
| `Abs (s, l1) ->
List.filter (fun v -> v <> s) (X.freevars l1)
end
module rec LambdaFV : (FreeVars with type exp = LambdaFV.exp lambda) =
LFVF(LambdaFV)

let test1 = (`App (`Abs ("x", `Var "x"), `Var "y") : Lambda.exp)
let e_test = Lambda.eval [] test1
let fv_test = LambdaFV.freevars test1

The arithmetic expression sub-language:

type 'a expr =
[ `Var of string | `Num of int | `Add of 'a * 'a | `Mult of 'a * 'a]

module type Operations =
sig include Eval include FreeVars with type exp := exp end

module EF(X : Operations with type exp = private [> 'a expr] as 'a) =
struct
type exp = X.exp expr

let map_expr f = function
| #var as v -> v
| `Num _ as n -> n
| `Add (e1, e2) -> `Add (f e1, f e2)
| `Mult (e1, e2) -> `Mult (f e1, f e2)

let eval subst (e : exp) : X.exp =

```

```

match map_expr (X.eval subst) e with
| #var as v -> eval_var subst v
| `Add (`Num m, `Num n) -> `Num (m + n)
| `Mult (`Num m, `Num n) -> `Num (m * n)
| e -> e

let freevars : exp -> 'b = function
| `Var v -> [v]
| `Num _ -> []
| `Add (e1, e2) | `Mult (e1, e2) -> X.freevars e1 @ X.freevars e2
end
module rec Expr : (Operations with type exp = Expr.exp expr) =
  EF(Expr)

let test2 = (`Add (`Mult (`Num 3, `Var "x"), `Num 1) : Expr.exp)
let e_test2 = Expr.eval ["x", `Num 2] test2
let fvs_test2 = Expr.freevars test2

```

Merging the sub-languages:

```

type 'a lexpr = ['a lambda | 'a expr]

module LEF(X : Operations with type exp = private [> 'a lexpr] as 'a) =
  struct
    type exp = X.exp lexpr
    module LambdaX = LF(X)
    module LambdaFVX = LFVF(X)
    module ExprX = EF(X)

    let eval subst : exp -> X.exp = function
      | #LambdaX.exp as x -> LambdaX.eval subst x
      | #ExprX.exp as x -> ExprX.eval subst x

    let freevars : exp -> 'b = function
      | #lambda as x -> LambdaFVX.freevars x (* Either of #lambda or #LambdaX.exp is fine *)
      | #expr as x -> ExprX.freevars x (* Either of #expr or #ExprX.exp is fine *)
  end
  module rec LExpr : (Operations with type exp = LExpr.exp lexpr) =
    LEF(LExpr)

let test3 =
  (`App (`Abs ("x", `Add (`Mult (`Num 3, `Var "x"), `Num 1)),
         `Num 2) : LExpr.exp)
let e_test3 = LExpr.eval [] test3
let fv_test3 = LExpr.freevars test3
let e_old_test = LExpr.eval [] (test2 :> LExpr.exp)
let fv_old_test = LExpr.freevars (test2 :> LExpr.exp)

```

11.9 Parser Combinators

We have done parsing using external languages OCamlLex and Menhir. Now we look at parsers written directly in OCaml.

Language **combinators** are ways of defining languages by composing definitions of smaller languages. For example, the combinators of the **Extended Backus-Naur Form** notation are:

- **Concatenation:** $S = A, B$ stands for $S = \{ab \mid a \in A, b \in B\}$
- **Alternation:** $S = A \mid B$ stands for $S = \{a \mid a \in A \vee a \in B\}$
- **Option:** $S = [A]$ stands for $S = \{\epsilon\} \cup A$, where ϵ is an empty string
- **Repetition:** $S = \{A\}$ stands for $S = \{\epsilon\} \cup \{as \mid a \in A, s \in S\}$
- **Terminal string:** $S = "a"$ stands for $S = \{a\}$

Parsers implemented directly in a functional programming paradigm are functions from character streams to the parsed values. Algorithmically they are **recursive descent parsers**.

Parser combinators approach builds parsers as **monad plus** values:

- **Bind:** `val (>>=) : 'a parser -> ('a -> 'b parser) -> 'b parser`
 - $p >>= f$ is a parser that first parses p , and makes the result available for parsing f
- **Return:** `val return : 'a -> 'a parser`
 - $\text{return } x$ parses an empty string, symbolically $S = \{\epsilon\}$, and returns x
- **MZero:** `val fail : 'a parser`
 - fail fails to parse anything, symbolically $S = \emptyset = \{\}$
- **MPlus:** `val (<|>) : 'a parser -> 'a parser -> 'a parser`
 - $p <|> q$ tries p , and if p succeeds, its result is returned, otherwise the parser q is used

The only non-monad-plus operation that has to be built into the monad is some way to consume a single character from the input stream, for example:

- `val satisfy : (char -> bool) -> char parser`
 - $\text{satisfy } (\text{fun } c \rightarrow c = 'a')$ consumes the character “ a ” from the input stream and returns it; if the input stream starts with a different character, this parser fails

Ordinary monadic recursive descent parsers **do not allow left-recursion**: if a cycle of calls not consuming any character can be entered when a parse failure should occur, the cycle will keep repeating.

For example, if we define numbers $N := D \mid ND$, where D stands for digits, then a stack of uses of the rule $N \rightarrow ND$ will build up when the next character is not a digit.

On the other hand, rules can share common prefixes.

11.10 Parser Combinators: Implementation

The parser monad is actually a composition of two monads:

- The **state monad** for storing the stream of characters that remain to be parsed
- The **backtracking monad** for handling parse failures and ambiguities

Alternatively, one can split the state monad into a reader monad with the parsed string, and a state monad with the parsing position.

We experiment with a different approach to monad-plus: **lazy-monad-plus**:

```
val mplus : 'a monad -> 'a monad Lazy.t -> 'a monad
```

Implementation of lazy-monad-plus First an operation from `MonadPlusOps`:

```
let msum_map f l =
  List.fold_left (* Folding left reverses the apparent order of composition *)
    (fun acc a -> mplus acc (lazy (f a))) mzero l (* order from l is preserved *)
```

The implementation of the lazy-monad-plus using lazy lists:

```
type 'a llist = LNil | LCons of 'a * 'a llist Lazy.t

let rec ltake n = function
  | LCons (a, l) when n > 1 -> a :: (ltake (n-1) (Lazy.force l))
  | LCons (a, l) when n = 1 -> [a] (* Avoid forcing the tail if not needed *)
  | _ -> []

let rec lappend l1 l2 =
  match l1 with LNil -> Lazy.force l2
  | LCons (hd, tl) -> LCons (hd, lazy (lappend (Lazy.force tl) l2))

let rec lconcat_map f = function
  | LNil -> LNil
  | LCons (a, l) -> lappend (f a) (lazy (lconcat_map f (Lazy.force l)))

module LListM = MonadPlus (struct
  type 'a t = 'a llist
  let bind a b = lconcat_map b a
  let return a = LCons (a, lazy LNil)
  let mzero = LNil
  let mplus = lappend
end)
```

The Parsec Monad File `Parsec.ml`:

```
open Monad
```

```

module type PARSE = sig
  type 'a backtracking_monad (* Name for the underlying monad-plus *)
  type 'a parsing_state = int -> ('a * int) backtracking_monad (* Processing state -- position *)
  type 'a t = string -> 'a parsing_state (* Reader for the parsed text *)
  include MONAD_PLUS_OPS
  val (<|>) : 'a monad -> 'a monad Lazy.t -> 'a monad (* A synonym for mplus *)
  val run : 'a monad -> 'a t
  val runT : 'a monad -> string -> int -> 'a backtracking_monad
  val satisfy : (char -> bool) -> char monad (* Consume a character of the specified class *)
  val end_of_text : unit monad (* Check for end of the processed text *)
end

module ParseT (MP : MONAD_PLUS_OPS) :
  PARSE with type 'a backtracking_monad := 'a MP.monad =
struct
  type 'a backtracking_monad = 'a MP.monad
  type 'a parsing_state = int -> ('a * int) MP.monad
  module M = struct
    type 'a t = string -> 'a parsing_state
    let return a = fun s p -> MP.return (a, p)
    let bind m b = fun s p ->
      MP.bind (m s p) (fun (a, p') -> b a s p')
    let mzero = fun _ p -> MP.mzero
    let mplus ma mb = fun s p ->
      MP.mplus (ma s p) (lazy (Lazy.force mb s p))
  end
  include M
  include MonadPlusOps(M)
  let (<|>) ma mb = mplus ma mb
  let runT m s p = MP.lift fst (m s p)
  let satisfy f s p =
    if p < String.length s && f s.[p] (* Consuming a character means accessing it *)
    then MP.return (s.[p], p + 1) else MP.mzero (* and advancing the parsing position *)
  let end_of_text s p =
    if p >= String.length s then MP.return ((), p) else MP.mzero
end

```

Additional Parser Operations

```

module type PARSE_OPS = sig
  include PARSE
  val many : 'a monad -> 'a list monad
  val opt : 'a monad -> 'a option monad
  val (?! : 'a monad -> 'a option monad
  val seq : 'a monad -> 'b monad Lazy.t -> ('a * 'b) monad (* Exercise: why laziness here?*)

```

```

val (<*>) : 'a monad -> 'b monad Lazy.t -> ('a * 'b) monad (* Synonym for seq *)
val lowercase : char monad
val uppercase : char monad
val digit : char monad
val alpha : char monad
val alphanum : char monad
val literal : string -> unit monad (* Consume characters of the given string *)
val (<<>) : string -> 'a monad -> 'a monad (* Prefix and postfix keywords *)
val (<>>) : 'a monad -> string -> 'a monad
end

module ParseOps (R : MONAD_PLUS_OPS)
  (P : PARSE with type 'a backtracking_monad := 'a R.monad) :
  PARSE_OPS with type 'a backtracking_monad := 'a R.monad =
struct
  include P
  let rec many p =
    (perform
      r <- p; rs <- many p; return (r::rs))
    ++ lazy (return [])
  let opt p = (p >>= (fun x -> return (Some x))) ++ lazy (return None)
  let (?!) p = opt p
  let seq p q = perform
    x <- p; y <- Lazy.force q; return (x, y)
  let (<*>) p q = seq p q
  let lowercase = satisfy (fun c -> c >= 'a' && c <= 'z')
  let uppercase = satisfy (fun c -> c >= 'A' && c <= 'Z')
  let digit = satisfy (fun c -> c >= '0' && c <= '9')
  let alpha = lowercase ++ lazy uppercase
  let alphanum = alpha ++ lazy digit
  let literal l =
    let rec loop pos =
      if pos = String.length l then return ()
      else satisfy (fun c -> c = l.[pos]) >>- loop (pos + 1) in
    loop 0
  let (<<>) bra p = literal bra >>- p
  let (<>>) p ket = p >>= (fun x -> literal ket >>- return x)
end

```

11.11 Parser Combinators: Tying the Recursive Knot

File PluginBase.ml:

```

module ParseM =
  Parsec.ParseOps (Monad.LListM) (Parsec.ParseT (Monad.LListM))
open ParseM

```

```

let grammar_rules : (int monad -> int monad) list ref = ref []

let get_language () : int monad =
  let rec result =
    lazy
      (List.fold_left
        (fun acc lang -> acc <|> lazy (lang (Lazy.force result)))
        mzero !grammar_rules) in
  perform r <-- Lazy.force result; end_of_text; return r (* Ensure we parse the whole text *)

```

11.12 Parser Combinators: Dynamic Code Loading

File PluginRun.ml:

```

let load_plug fname : unit =
  let fname = Dynlink.adapt_filename fname in
  if Sys.file_exists fname then
    try Dynlink.loadfile fname
    with
    | (Dynlink.Error err) as e ->
        Printf.printf "\nERROR loading plugin: %s\n%"!
          (Dynlink.error_message err);
        raise e
    | e -> Printf.printf "\nUnknow error while loading plugin\n%"!
  else (
    Printf.printf "\nPlugin file %s does not exist\n%"! fname;
    exit (-1))

let () =
  for i = 2 to Array.length Sys.argv - 1 do
    load_plug Sys.argv.(i) done;
  let lang = PluginBase.get_language () in
  let result =
    Monad.LListM.run
      (PluginBase.ParseM.runT lang Sys.argv.(1) 0) in
  match Monad.ltake 1 result with
  | [] -> Printf.printf "\nParse error\n%"!
  | r::_ -> Printf.printf "\nResult: %d\n%"! r

```

11.13 Parser Combinators: Toy Example

File Plugin1.ml:

```

open PluginBase.ParseM
let digit_of_char d = int_of_char d - int_of_char '0'

```

```

let number _ = (* Numbers: N := D N | D where D is digits *)
  let rec num =
    lazy ( (perform
      d <-- digit;
      (n, b) <-- Lazy.force num;
      return (digit_of_char d * b + n, b * 10))
      <|> lazy (digit >>= (fun d -> return (digit_of_char d, 10)))) in
    Lazy.force num >>| fst

let addition lang = (* Addition rule: S -> (S + S) *)
  perform (* Requiring a parenthesis ( turns the rule into non-left-recursive *)
    literal "("; n1 <-- lang; literal "+"; n2 <-- lang; literal ")";
    return (n1 + n2)

let () =
  PluginBase.(grammar_rules := number :: addition :: !grammar_rules)

File Plugin2.ml:
open PluginBase.ParseM

let multiplication lang = (* Multiplication rule: S -> (S * S) *)
  perform
    literal "("; n1 <-- lang; literal "*"; n2 <-- lang; literal ")";
    return (n1 * n2)

let () =
  PluginBase.(grammar_rules := multiplication :: !grammar_rules)

```

11.14 Exercises

Exercise 1: Implement the `string_of_` functions or methods, covering all data cases, corresponding to the `eval_` functions in at least two examples from the lecture, including both an object-based example and a variant-based example (either standard, or polymorphic, or extensible variants).

Exercise 2: Split at least one of the examples from the previous exercise into multiple files and demonstrate separate compilation.

Exercise 3: Can we drop the tags `Lambda_t`, `Expr_t` and `LExpr_t` used in the examples based on standard variants (file `FP_ADT.ml`)? When using polymorphic variants, such tags are not needed.

Exercise 4: Factor-out the sub-language consisting only of variables, thus eliminating the duplication of tags `VarL`, `VarE` in the examples based on standard variants (file `FP_ADT.ml`).

Exercise 5: Come up with a scenario where the extensible variant types-based solution leads to a non-obvious or hard to locate bug.

Exercise 6: Re-implement the direct object-based solution to the expression problem (file `Objects.ml`) to make it more satisfying. For example, eliminate the need for some of the `rename`, `apply`, `compute` methods.

Exercise 7: Re-implement the visitor pattern-based solution to the expression problem (file `Visitor.ml`) in a functional way, i.e., replace the mutable fields `subst` and `beta_redex` in the `eval_lambda` class with a different solution to the problem of treating `abs` and non-`abs` expressions differently.

Exercise 8: Extend the sub-language `expr_visit` with variables, and add to arguments of the evaluation constructor `eval_expr` the substitution. Handle the problem of potentially duplicate fields `subst`. (One approach might be to use ideas from exercise 6.)

Exercise 9: Implement the following modifications to the example from the file `PolyV.ml`:

1. Factor-out the sub-language of variables, around the already present `var` type.
2. Open the types of functions `eval3`, `freevars3` and other functions as required, so that explicit subtyping, e.g., in `eval3 [] (test2 :> lexpr_t)`, is not necessary.
3. Remove the double-dispatch currently in `eval_lexpr` and `freevars_lexpr`, by implementing a cascading design rather than a “divide-and-conquer” design.

Exercise 10: Streamline the solution `PolyRecM.ml` by extending the language of λ -expressions with arithmetic expressions, rather than defining the sub-languages separately and then merging them. See slide on page 15 of Jacques Garrigue *Structural Types, Recursive Modules, and the Expression Problem*.

Exercise 11: Transform a parser monad, or rewrite the parser monad transformer, by adding state for the line and column numbers.

Exercise 12: Implement `_of_string` functions as parser combinators on top of the example `PolyRecM.ml`. Sections 4.3 and 6.2 of *Monadic Parser Combinators* by Graham Hutton and Erik Meijer might be helpful. Split the result into multiple files as in Exercise 2 and demonstrate dynamic loading of code.

Exercise 13: What are the benefits and drawbacks of our lazy-monad-plus (built on top of *odd lazy lists*) approach, as compared to regular monad-plus built on top of *even lazy lists*? To additionally illustrate your answer:

1. Rewrite the parser combinators example to use regular monad-plus and even lazy lists.
2. Select one example from Lecture 8 and rewrite it using lazy-monad-plus and odd lazy lists.